

token-based \Rightarrow 1000 tokens



1000 embeddings

a single word \Rightarrow multiple embeddings

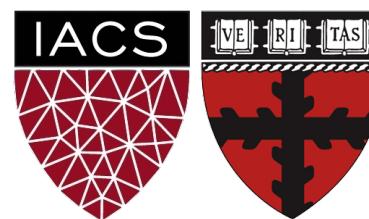
Lecture 25: Transformers

NLP Lectures: Part 4 of 4

Harvard IACS

CS109B

Pavlos Protopapas, Mark Glickman, and Chris Tanner



Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

First, we learned about **language models (LMs)**

$P(\text{"What is the weather today?")}$

$P(\text{"What is the whether two day?")}$

$P(\text{"What is the whether too day?")}$

Being able to correctly estimate the likelihood
of sentences is useful for many other tasks

First, we learned about **language models (LMs)**

Auto-complete

$P(\text{What is the weather tomorrow?})$

Machine Translation

$P(\text{What is the weather two days?})$

Text Generation

$P(\text{Text Classification})$

Speech Recognition

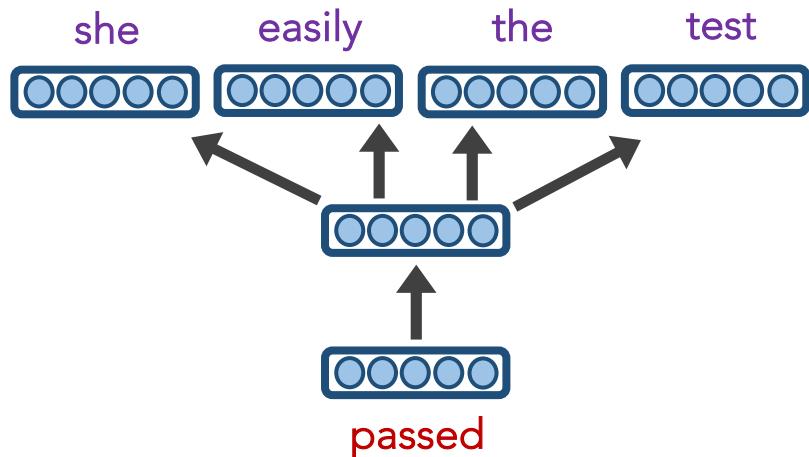
$P(\text{What is the weather today?})$

Being able to correctly estimate the likelihood
of sentences is useful for many other tasks

Next, we learned about **word embeddings**

TYPE-BASED

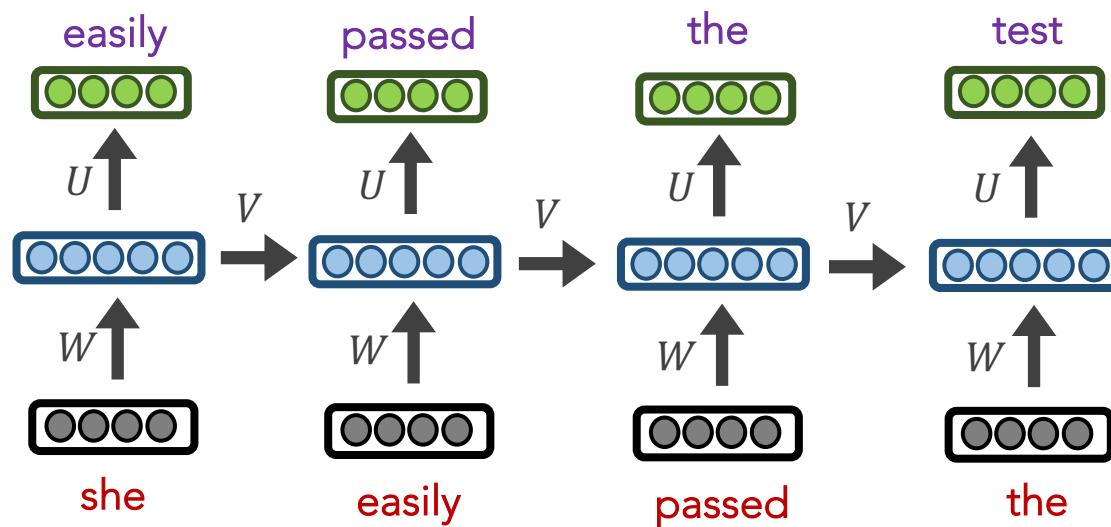
a single, **global word embedding** for each word, independent of its context.



word2vec (skip-gram)

TOKEN-BASED

contextualized embeddings are distinct for every occurrence of a word, completely dependent on its context

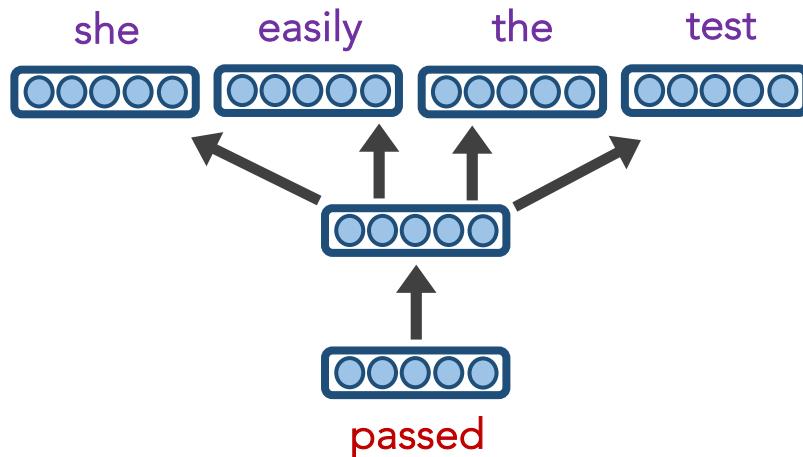


Bidirectional LSTM

Next, we learned about **word embeddings**

TYPE-BASED

a single, **global word embedding** for each word, independent of its context.



word2vec (skip-gram)

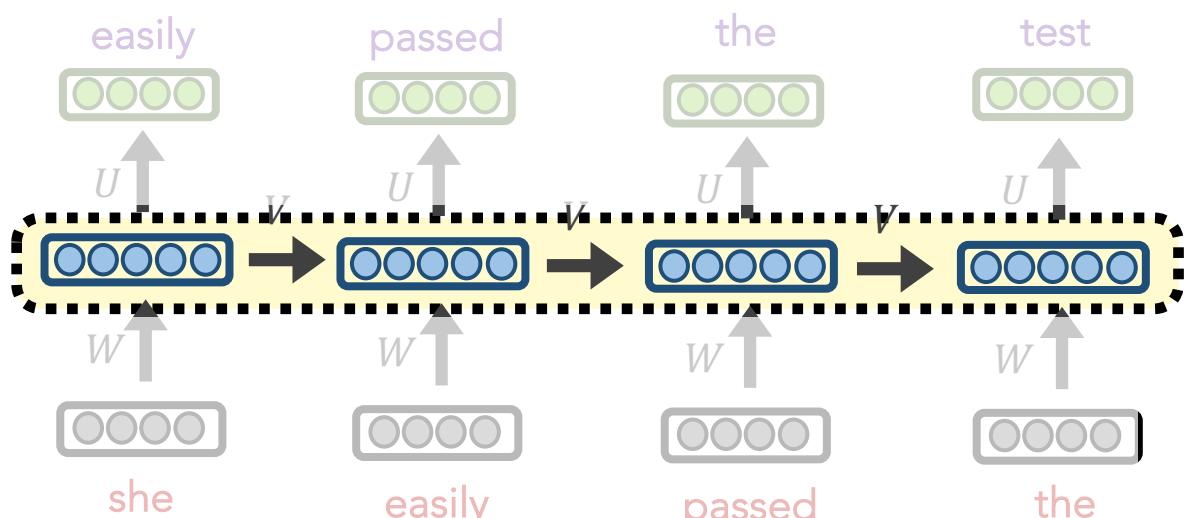
- These models **output** embeddings we can save to a file and use however we wish
- We then create a **separate model** that uses these embeddings
- Kind of limiting
- Often inferior, as of **2015**

Next, we learned about **word embeddings**

- These models are trained on a **specific task** (e.g., LM, text classification, etc)
 - The hidden layer(s) contains the “*meaning*” and are **very useful**
 - We can extract those embeddings if we wish, or grab the learned weights and re-use for another task
 - Dominating NLP from **2015 - present**

TOKEN-BASED

contextualized embeddings are
distinct for every occurrence of a word,
completely dependent on its context



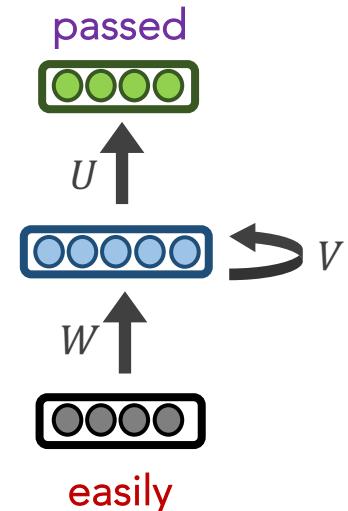
Bidirectional LSTM

Next, we learned about **word embeddings**

- LSTMs are amazing but ultimately only look at **1 word at a time, sequentially**
- Sure, they maintain long-term memory, but **they are short-sighted** in terms of knowing what to hold onto and how to weight each input

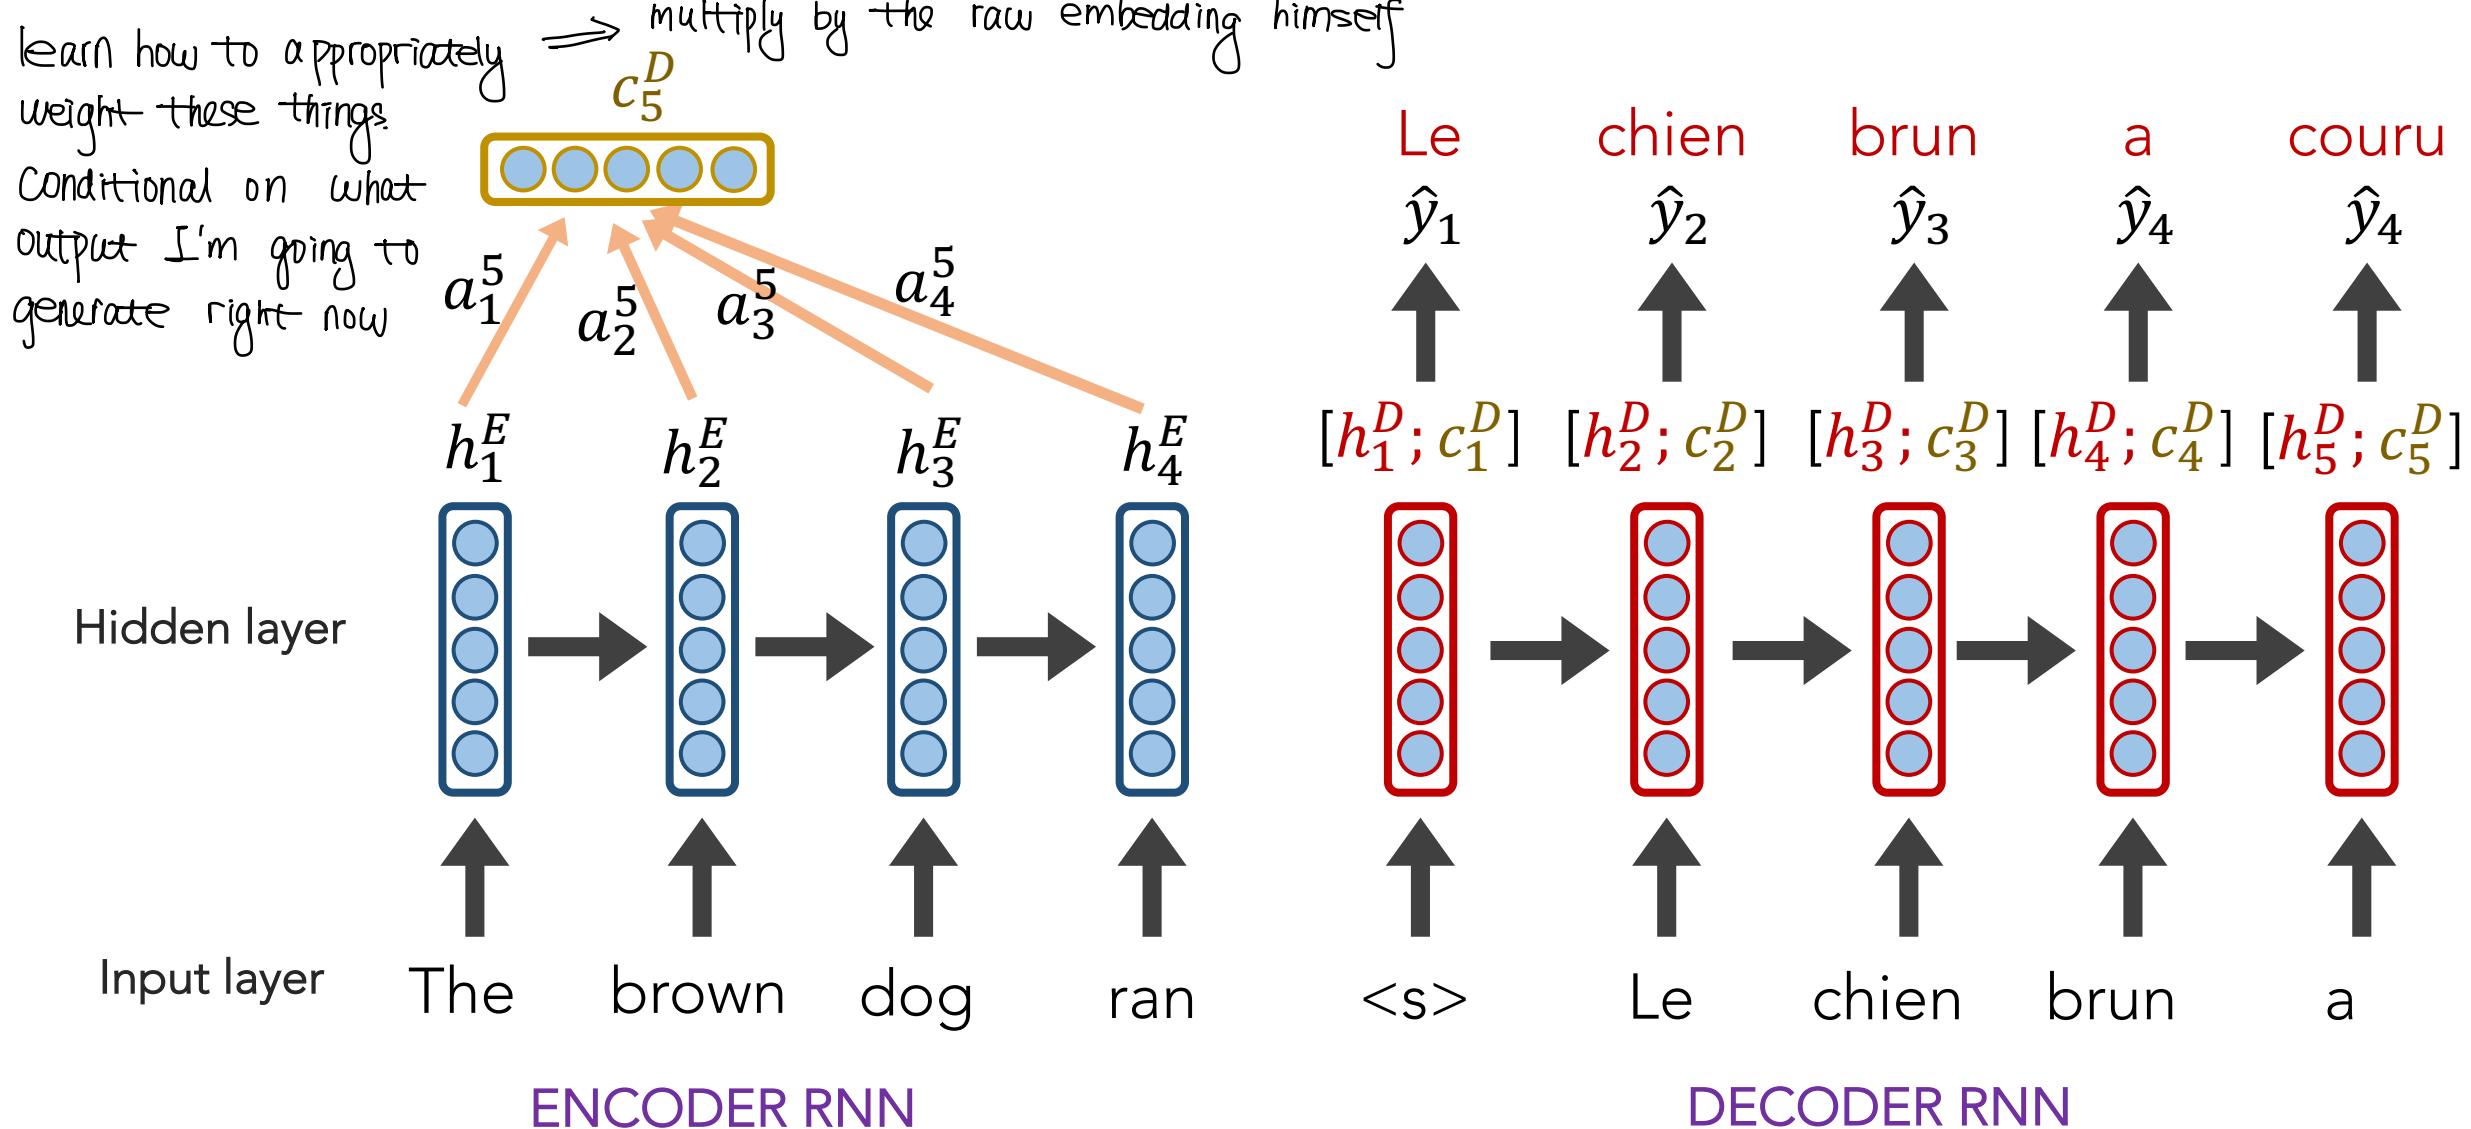
TOKEN-BASED

contextualized embeddings are distinct for every occurrence of a word, completely dependent on its context

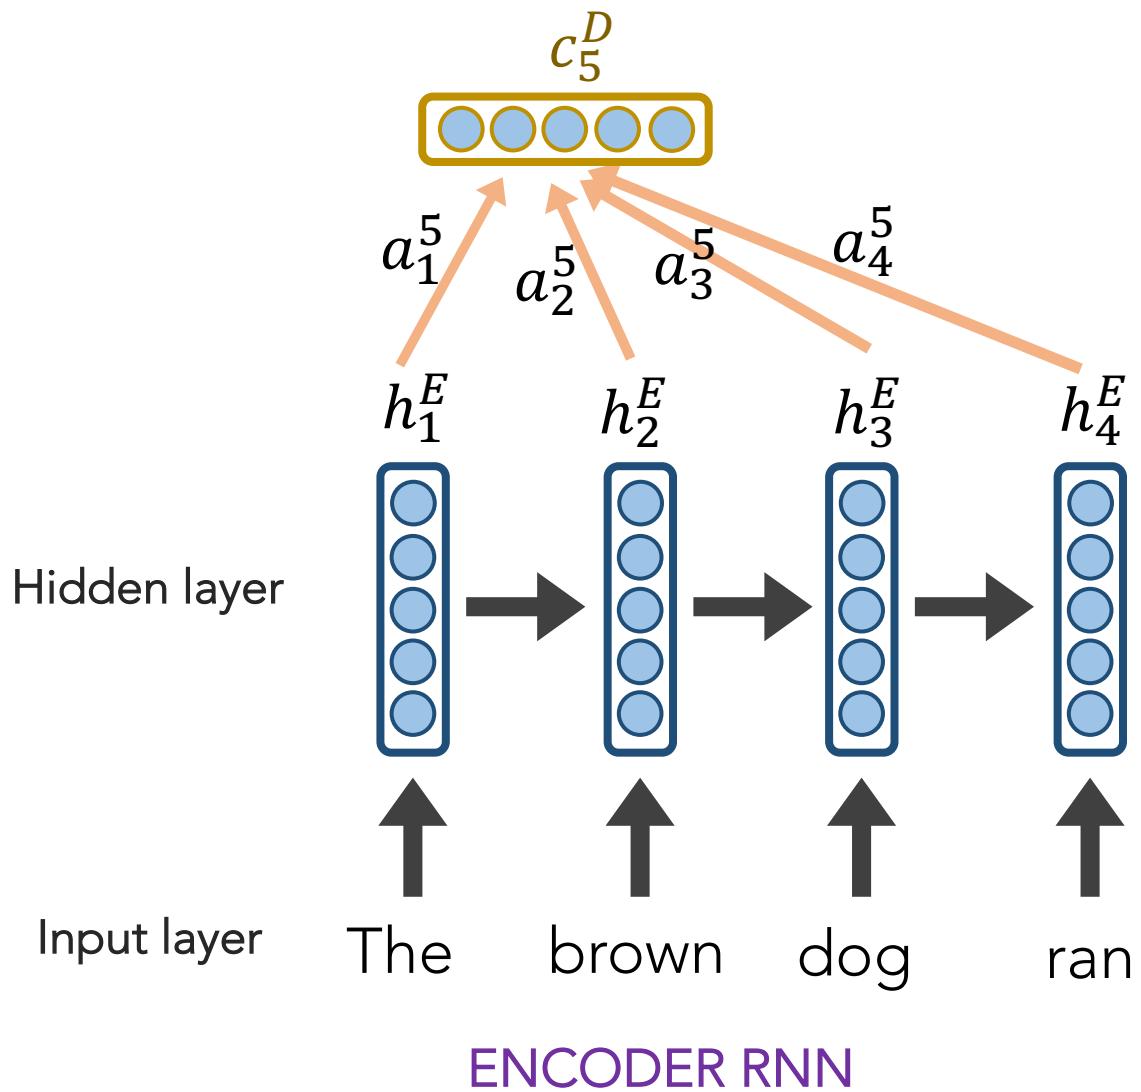


Bidirectional LSTM

Next, we learned about **seq2seq** and **Attention**



Next, we learned about **seq2seq** and **Attention**



- Revolutionary idea
- Decoder has access to **all input words** and appropriately **focuses** on select parts
- It's conditioned on the current word we're **decoding**

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Outline



Recap



Transformers



BERT



GPT-2



Concerns



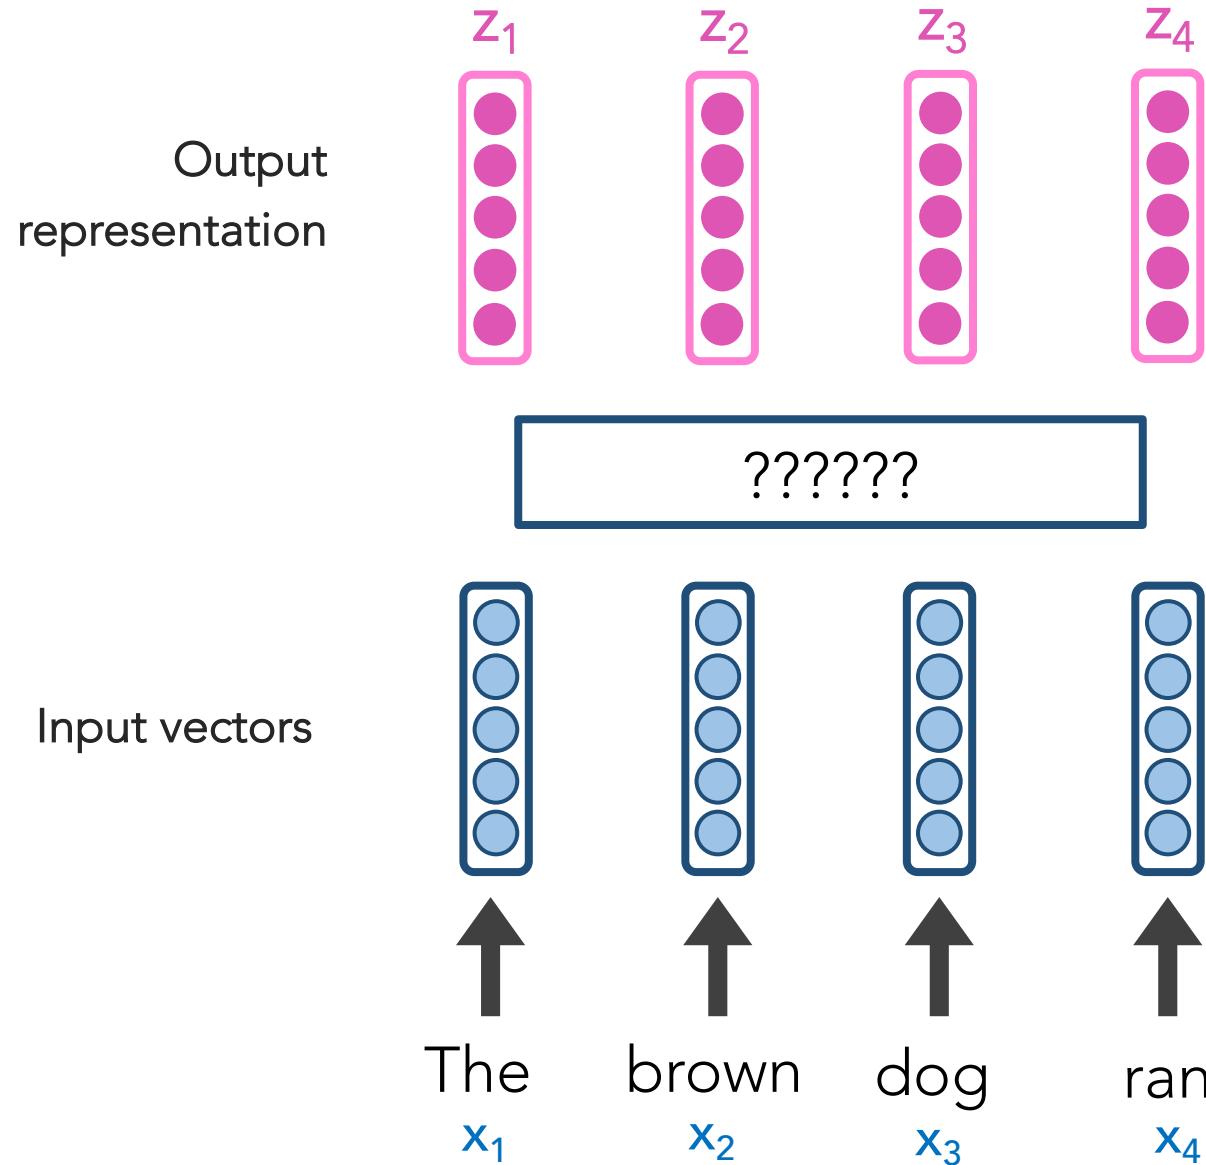
Summary

Goals

- Each word in a sequence to be transformed into a rich, abstract representation (**context embedding**) based on the weighted sums of the other words in the same sequence (akin to deep CNN layers)
- Inspired by **Attention**, we want each word to determine, “**how much should I be influenced by each of my neighbors**”
- Want positionality

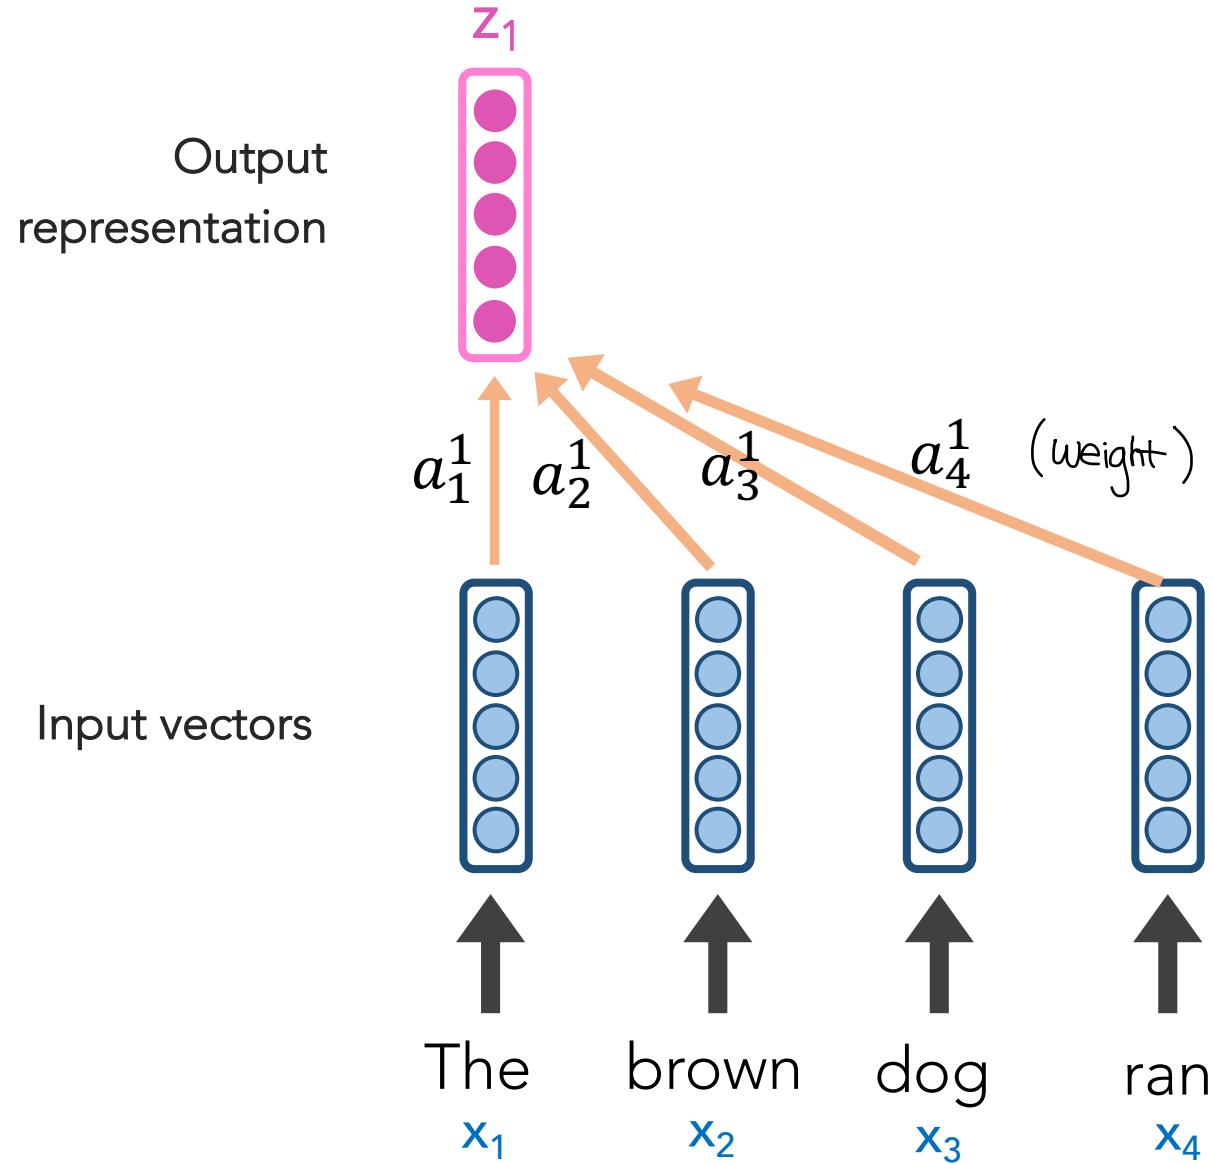
Self-Attention

same concept \Rightarrow how much would be influenced by the near neighbor



Self-Attention's goal is to create great representations, z_i , of the input

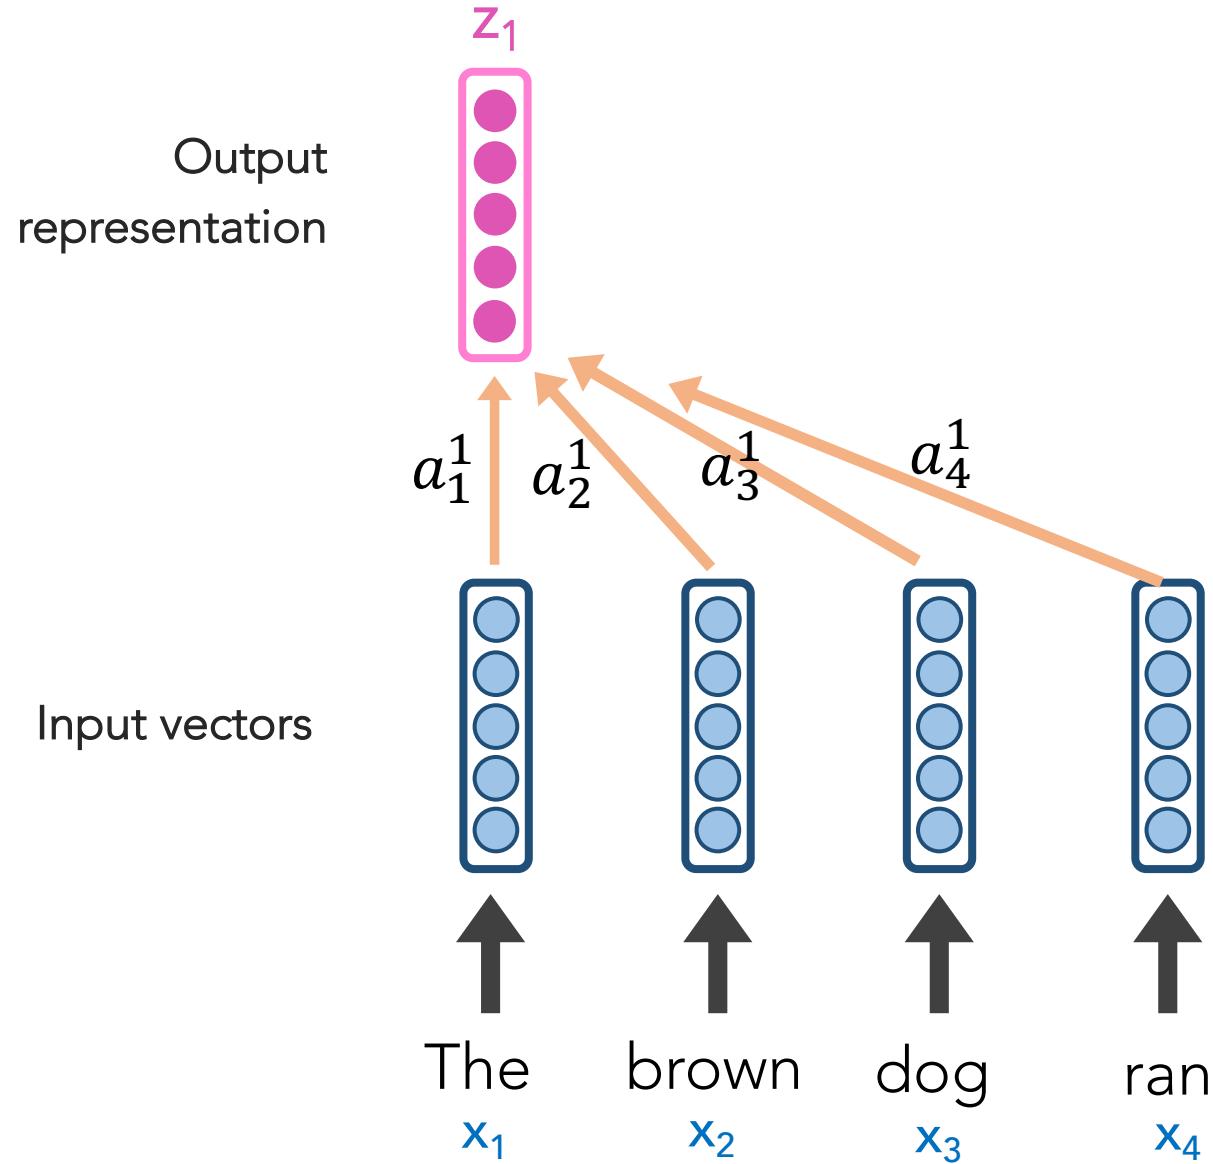
Self-Attention



Self-Attention's goal is to create great representations, z_i , of the input

z_1 will be based on a weighted contribution of x_1, x_2, x_3, x_4

Self-Attention

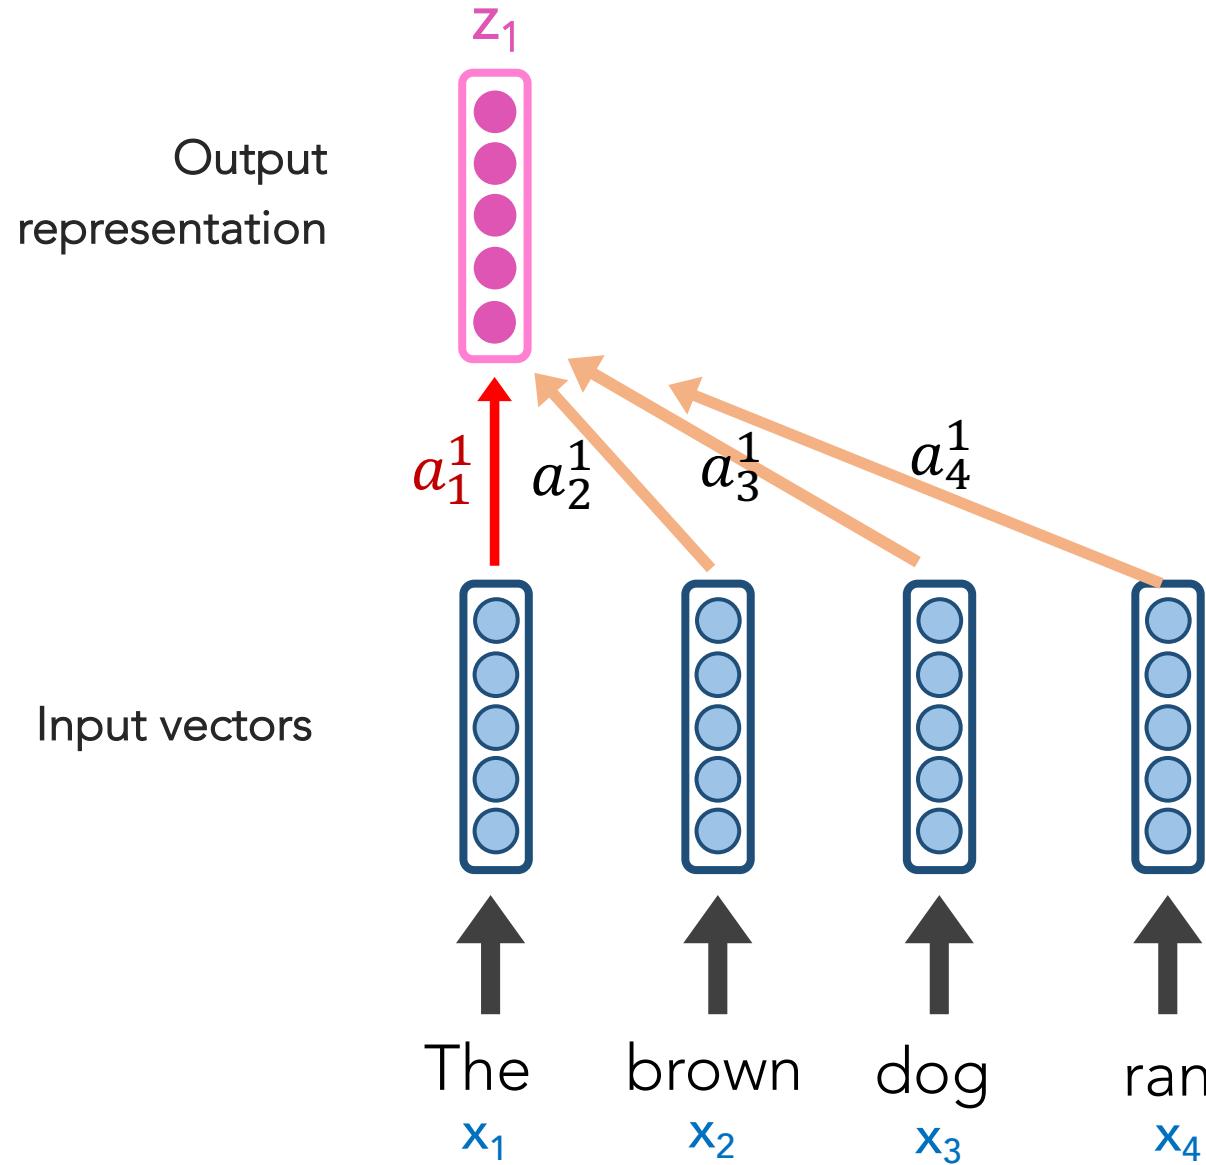


Self-Attention's goal is to create great representations, z_i , of the input

z_1 will be based on a weighted contribution of x_1, x_2, x_3, x_4

a_i^1 is “just” a weight. More is happening under the hood, but it’s effectively weighting versions of x_1, x_2, x_3, x_4

Self-Attention



Under the hood, each x_i has 3 small, associated vectors.

For example, x_1 has:

- Query q_i
- Key k_i
- Value v_i

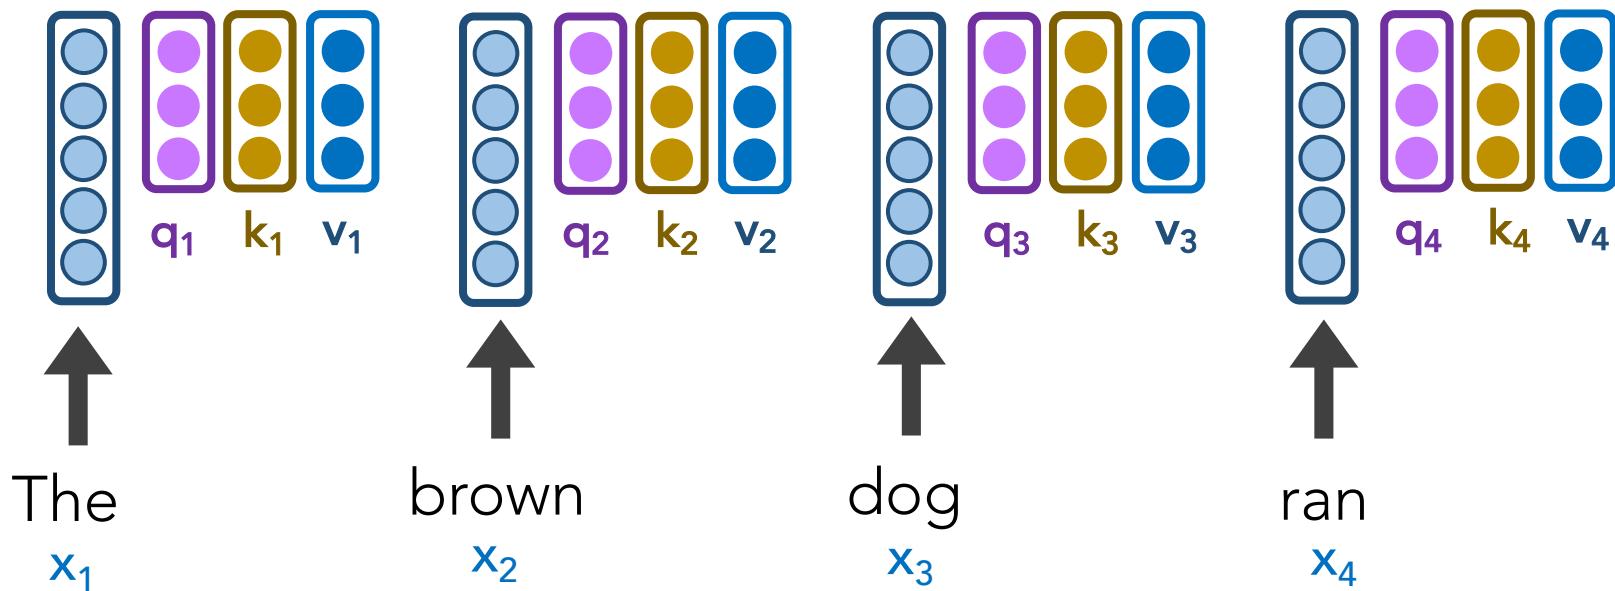
have those for all the input

Self-Attention

Step 1: Our Self-Attention Head has just 3 weight matrices W_q , W_k , W_v in total. These same 3 weight matrices are multiplied by each x_i to create all vectors:

$$\begin{aligned} q_i &= W_q x_i && \text{Same for different } i \\ k_i &= W_k x_i \\ v_i &= W_v x_i \end{aligned}$$

q_i , k_i , v_i are produced by the same matrices

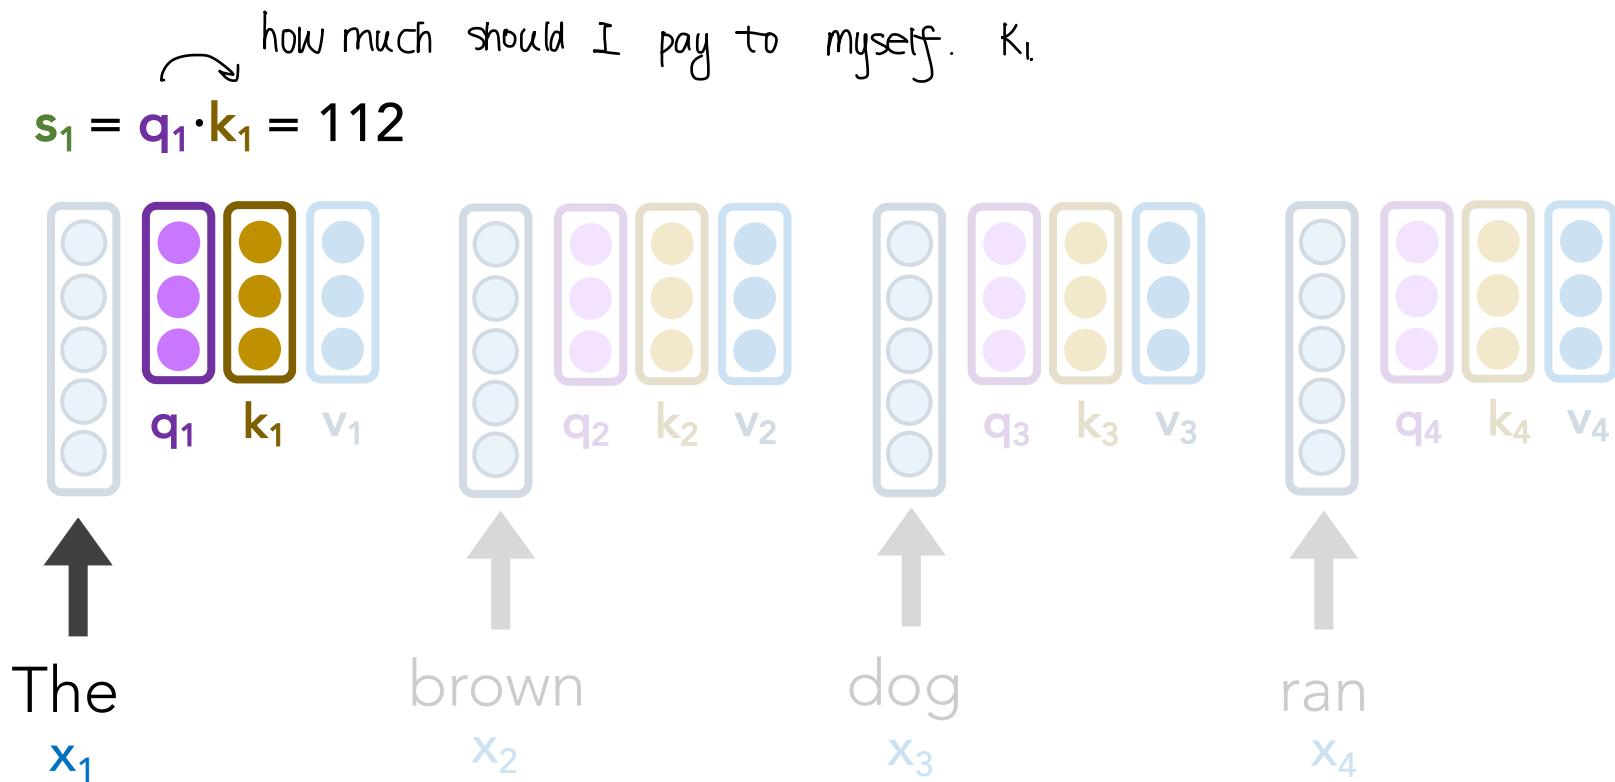


Under the hood, each x_i has 3 small, associated vectors. For example, x_1 has:

- Query q_1
- Key k_1
- Value v_1

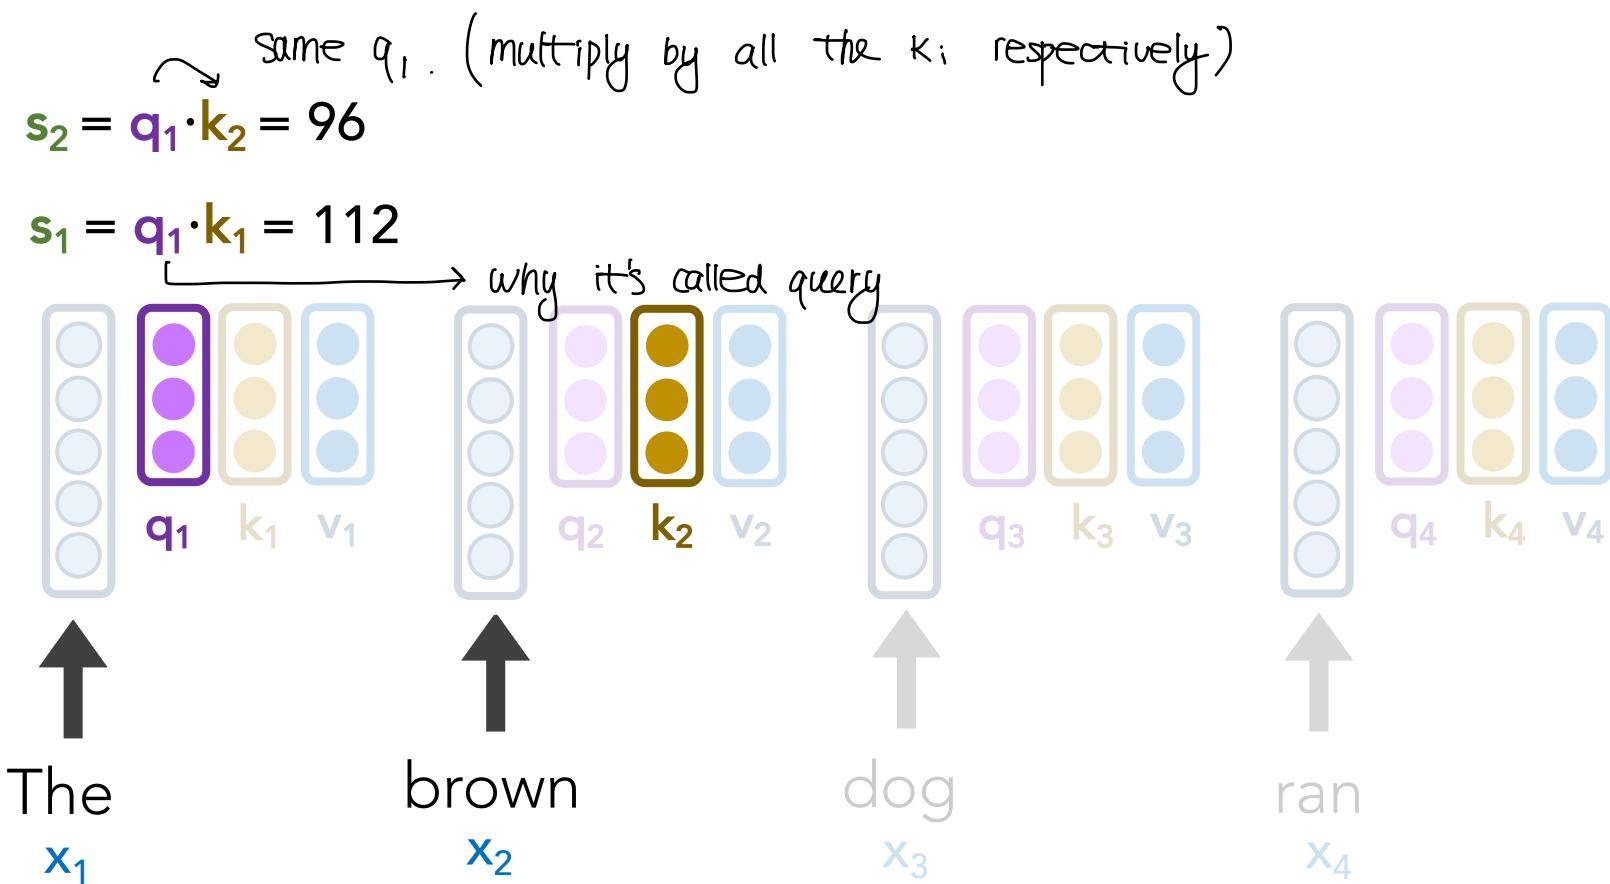
Self-Attention

Step 2: For word x_1 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i



Self-Attention

Step 2: For word x_1 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i



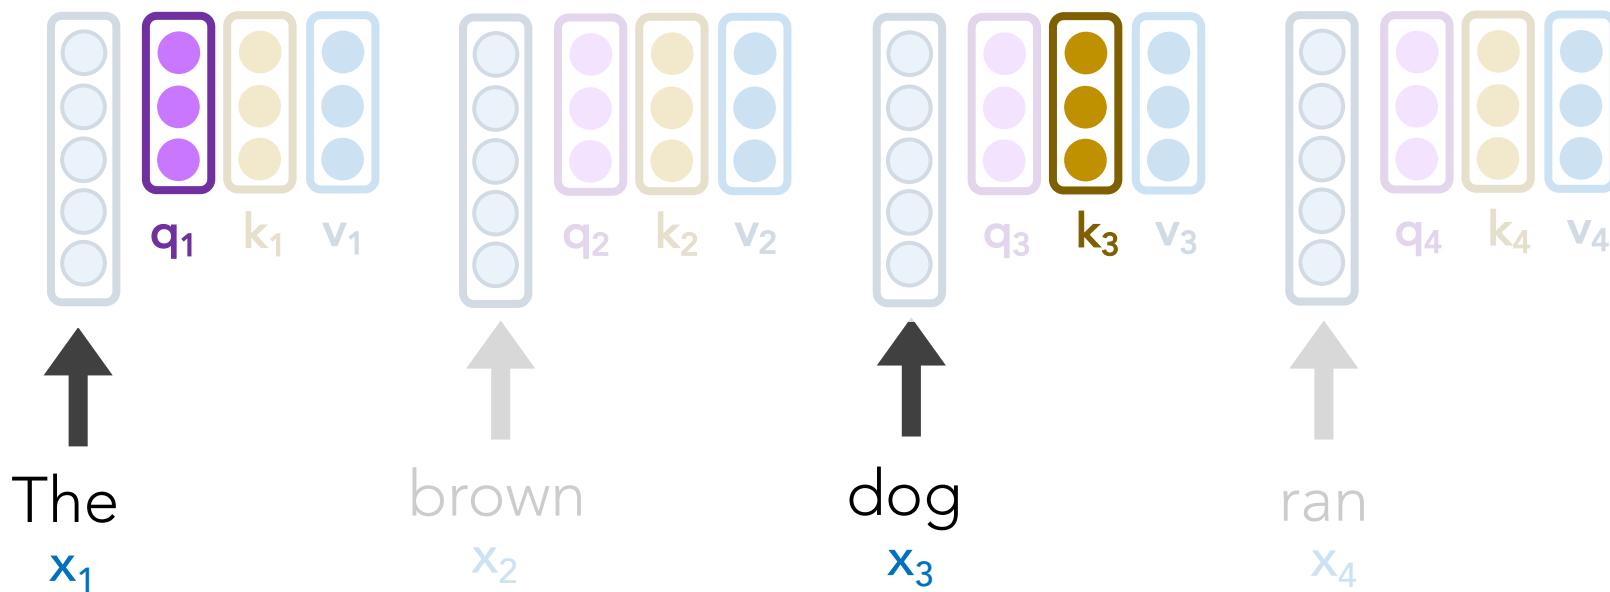
Self-Attention

Step 2: For word x_1 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i

$$s_3 = q_1 \cdot k_3 = 16$$

$$s_2 = q_1 \cdot k_2 = 96$$

$$s_1 = q_1 \cdot k_1 = 112$$



Self-Attention

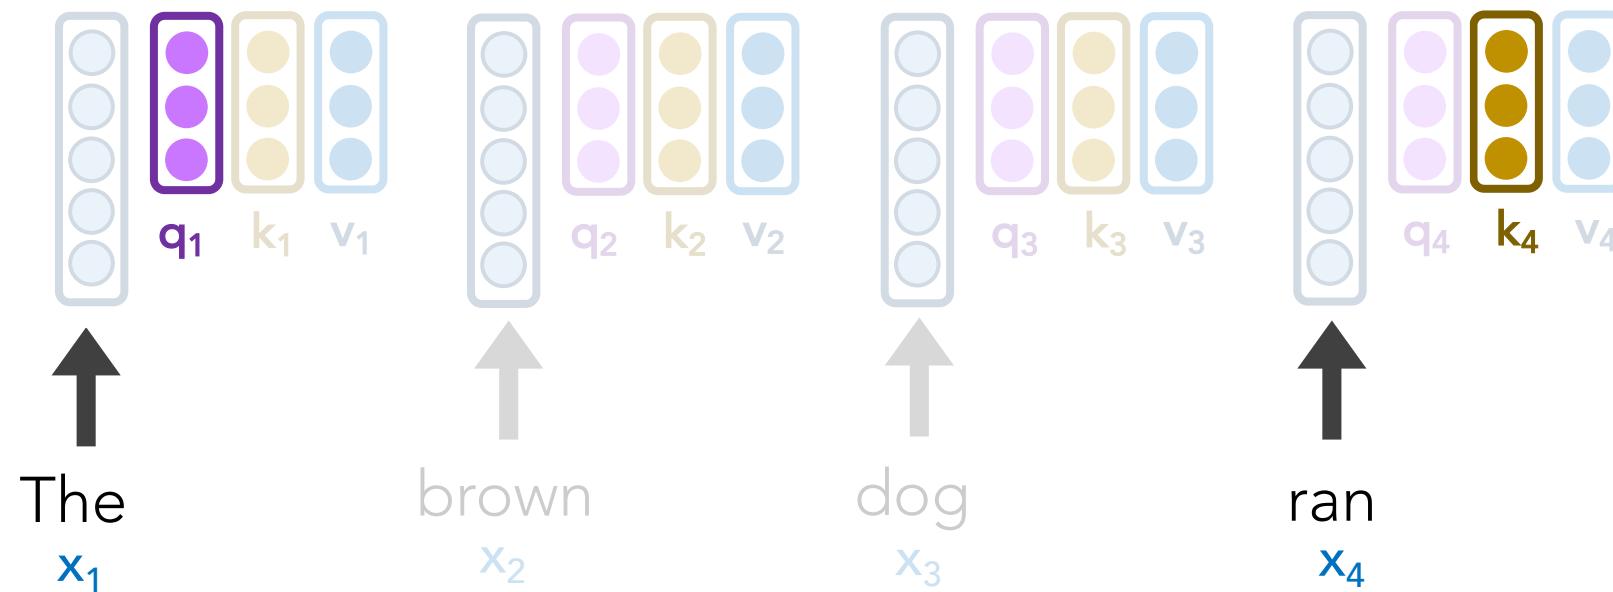
Step 2: For word x_1 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i

$$s_4 = q_1 \cdot k_4 = 8$$

$$s_3 = q_1 \cdot k_3 = 16$$

$$s_2 = q_1 \cdot k_2 = 96$$

$$s_1 = q_1 \cdot k_1 = 112$$



Self-Attention

Step 3: Our scores s_1, s_2, s_3, s_4 don't sum to 1. Let's divide by $\sqrt{\text{len}(k_i)}$ and softmax it

$$s_4 = q_1 \cdot k_4 = 8$$

$$a_4 = \sigma(s_4/8) = 0$$

$$s_3 = q_1 \cdot k_3 = 16$$

$$a_3 = \sigma(s_3/8) = .01$$

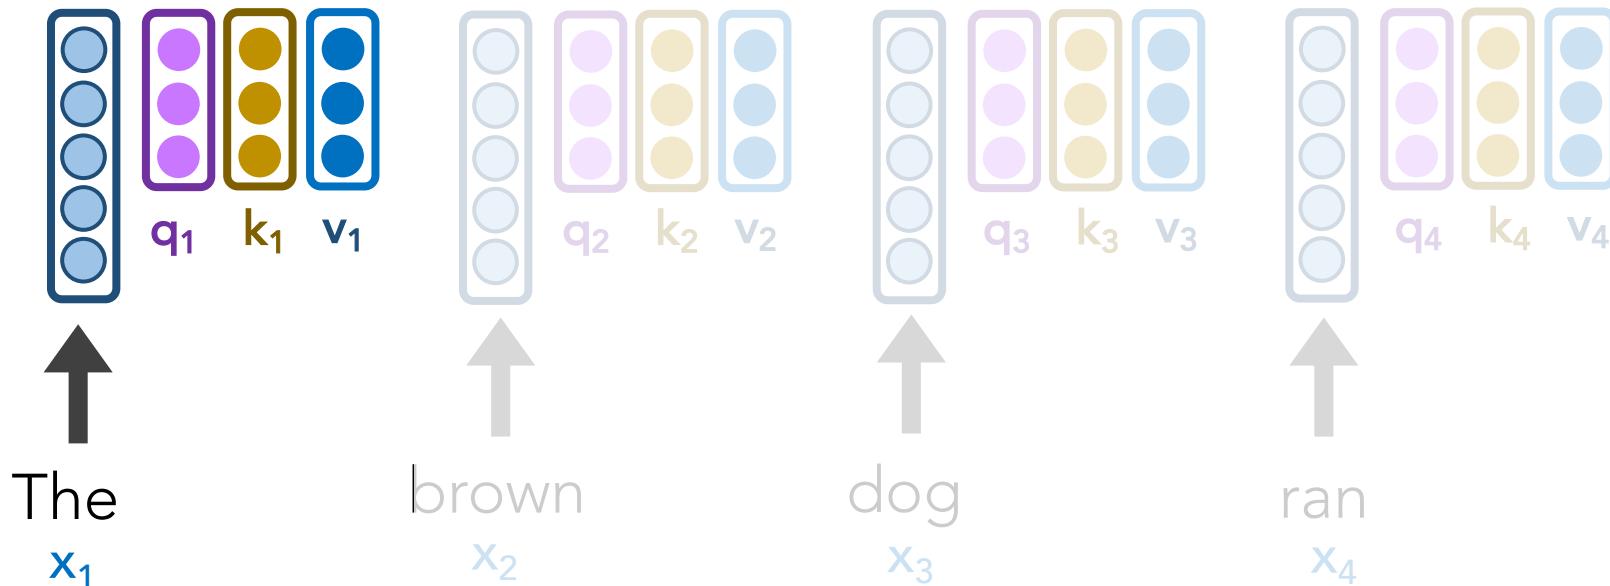
(how much to care about each of
the individual thing)

$$s_2 = q_1 \cdot k_2 = 96$$

$$a_2 = \sigma(s_2/8) = .12$$

$$s_1 = q_1 \cdot k_1 = 112$$

$$a_1 = \sigma(s_1/8) = .87$$



Self-Attention

Step 3: Our scores s_1, s_2, s_3, s_4 don't sum to 1. Let's divide by $\sqrt{\text{len}(k_i)}$ and softmax it

$$s_4 = q_1 \cdot k_4 = 8$$

$$a_4 = \sigma(s_4/8) = 0$$

$$s_3 = q_1 \cdot k_3 = 16$$

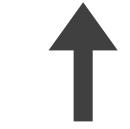
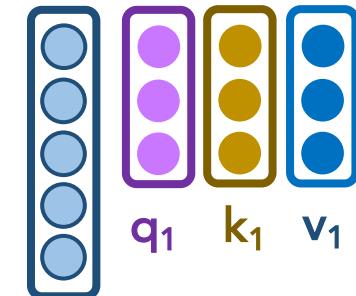
$$a_3 = \sigma(s_3/8) = .01$$

$$s_2 = q_1 \cdot k_2 = 96$$

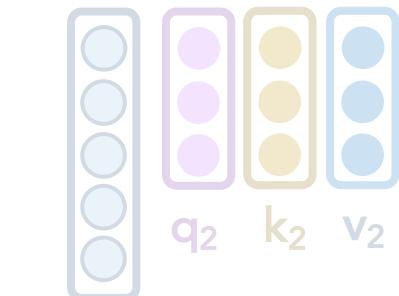
$$a_2 = \sigma(s_2/8) = .12$$

$$s_1 = q_1 \cdot k_1 = 112$$

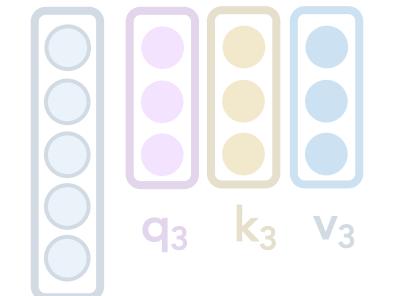
$$a_1 = \sigma(s_1/8) = .87$$



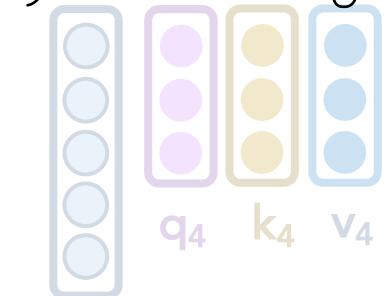
x_1



x_2
brown



x_3
dog



x_4
ran

⇒ sum to 1
self-attention weight

如果在这里 . vector 会
不一样. add positionality
brown 先 copy 一个再 add
positionality

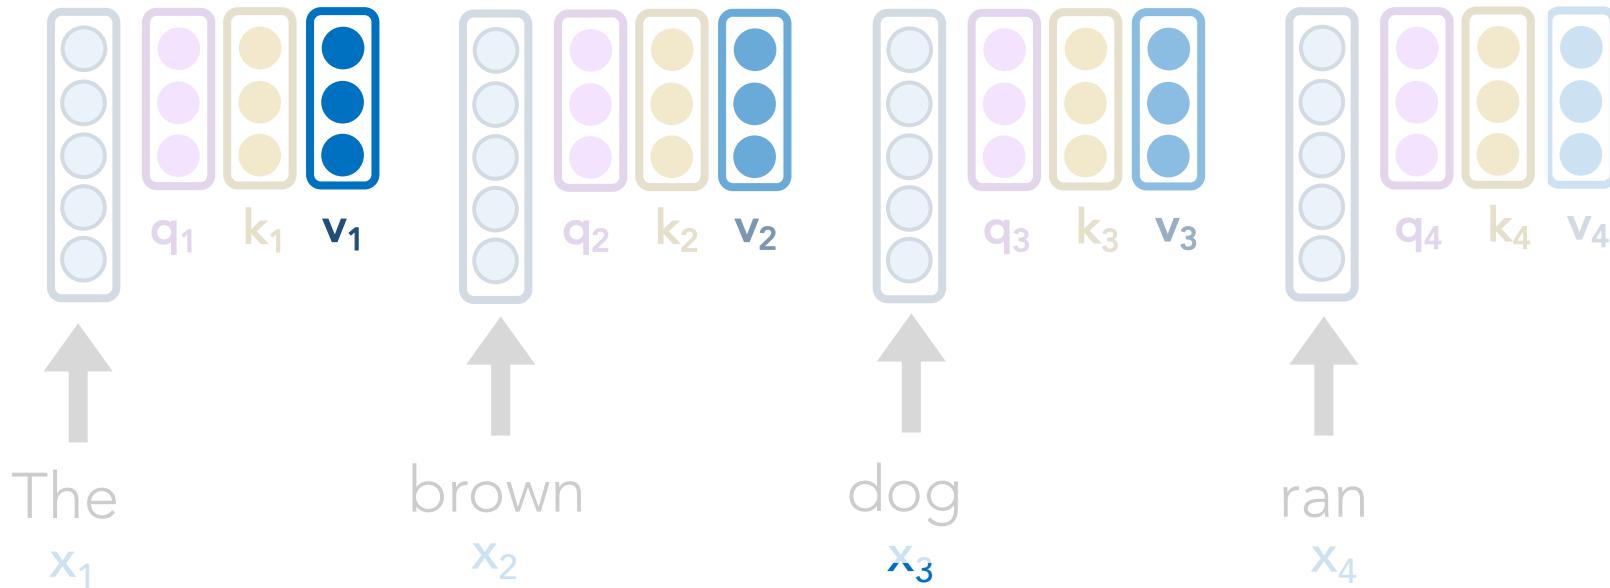
Self-Attention

Step 4: Let's weight our v_i vectors and simply sum them up!

z_1

$$\begin{aligned} z_1 &= a_1 \cdot v_1 + a_2 \cdot v_2 + a_3 \cdot v_3 + a_4 \cdot v_4 \\ &= 0.87 \cdot v_1 + 0.12 \cdot v_2 + 0.01 \cdot v_3 + 0 \cdot v_4 \end{aligned}$$

representatives . value factor



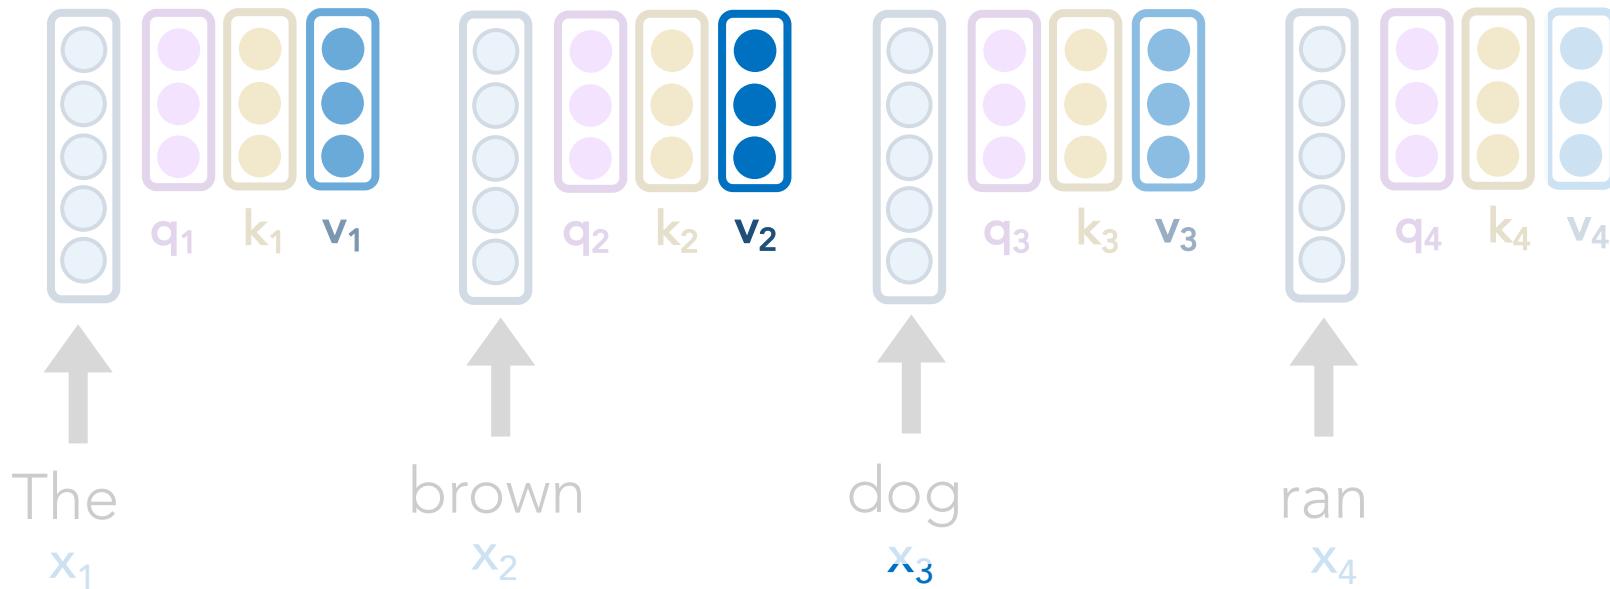
Self-Attention

Step 5: We repeat this for all other words, yielding us with great, new z_i representations!



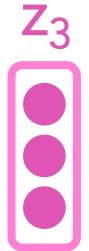
$$z_2 = a_1 \cdot v_1 + a_2 \cdot v_2 + a_3 \cdot v_3 + a_4 \cdot v_4$$

follow the same process

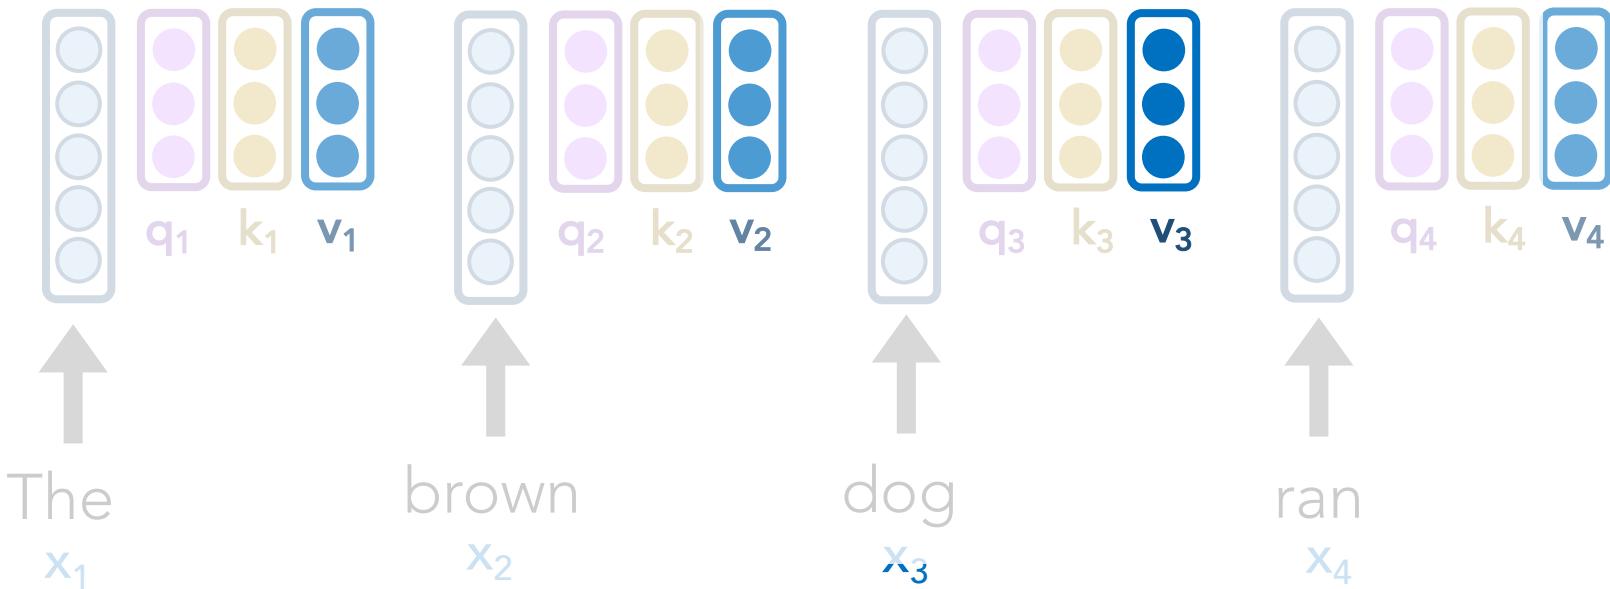


Self-Attention

Step 5: We repeat this for all other words, yielding us with great, new z_i representations!



$$z_3 = a_1 \cdot v_1 + a_2 \cdot v_2 + a_3 \cdot v_3 + a_4 \cdot v_4$$

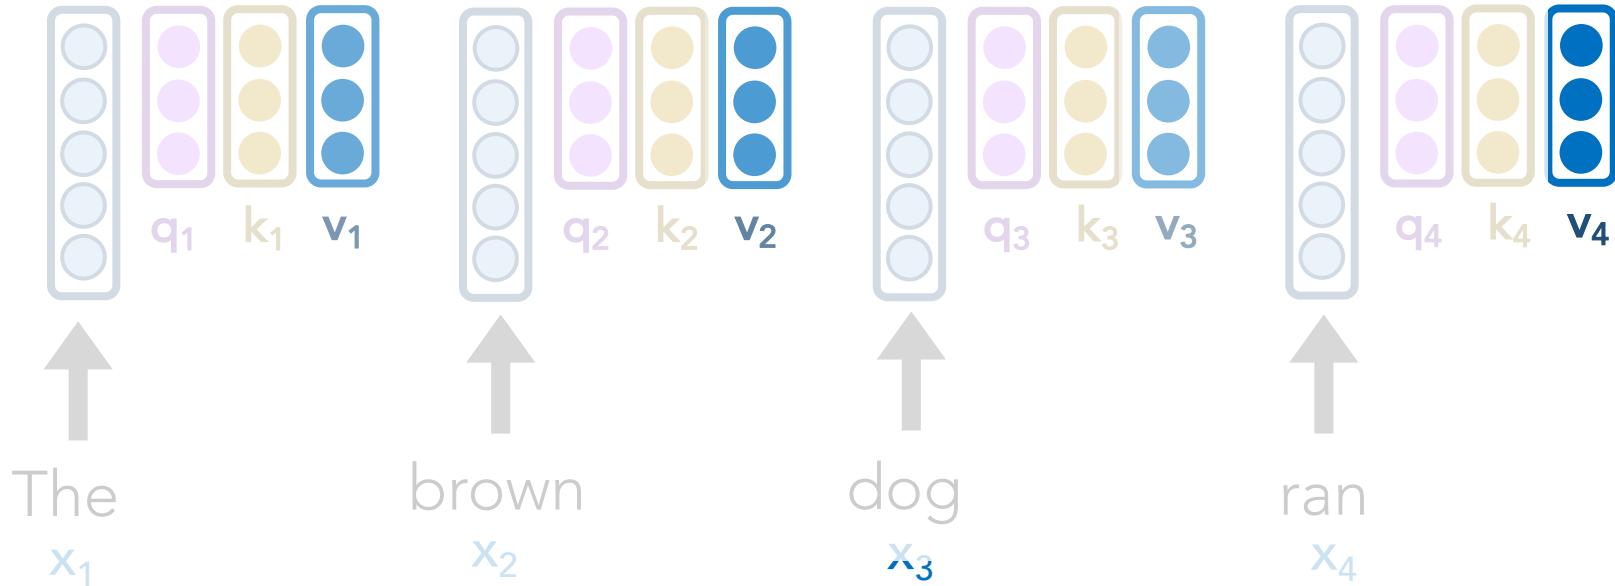


Self-Attention

Step 5: We repeat this for all other words, yielding us with great, new z_i representations!



$$z_4 = a_1 \cdot v_1 + a_2 \cdot v_2 + a_3 \cdot v_3 + a_4 \cdot v_4$$



Let's illustrate another example:



$$z_2 = a_1 \cdot v_1 + a_2 \cdot v_2 + a_3 \cdot v_3 + a_4 \cdot v_4$$

Remember, we use the same 3 weight matrices

W_q , W_k , W_v as we did for computing z_1 .

This gives us q_2, k_2, v_2

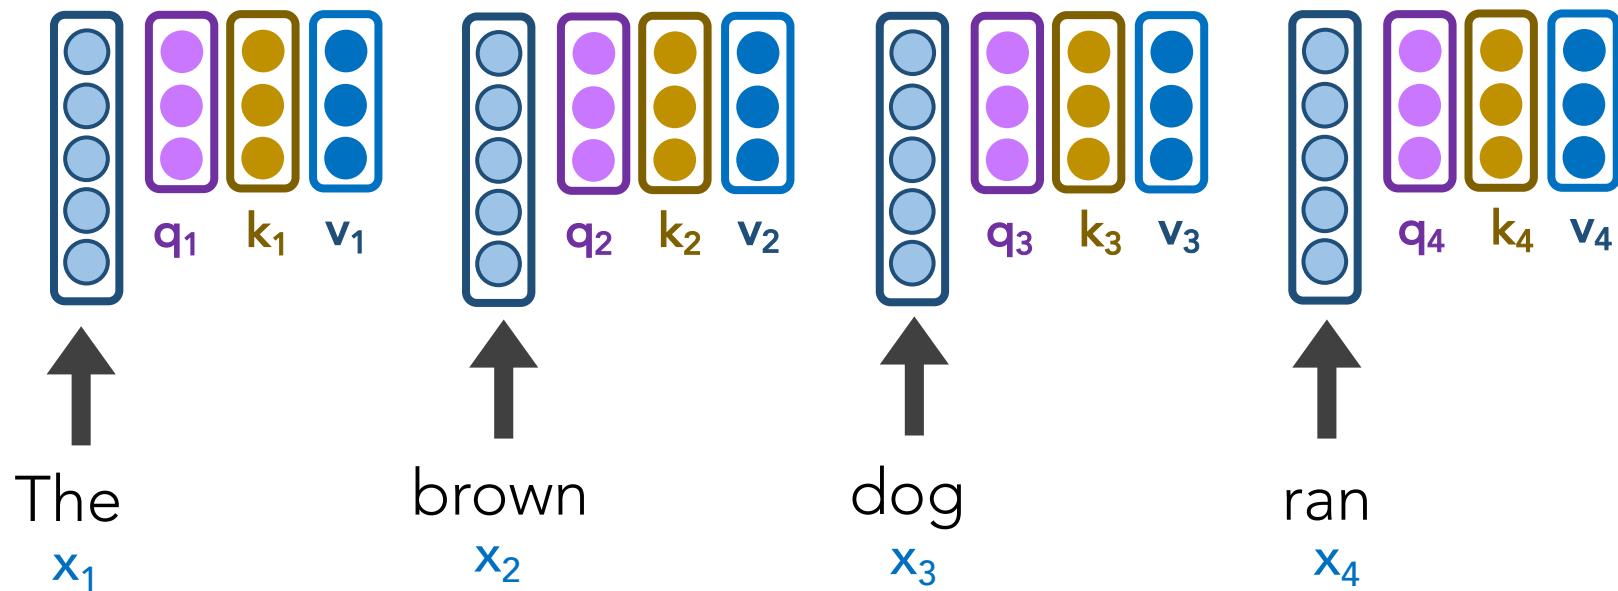
Self-Attention

Step 1: Our Self-Attention Head I has just 3 weight matrices W_q , W_k , W_v in total. These same 3 weight matrices are multiplied by each x_i to create all vectors:

$$q_i = w_q x_i$$

$$k_i = w_k x_i$$

$$v_i = w_v x_i$$



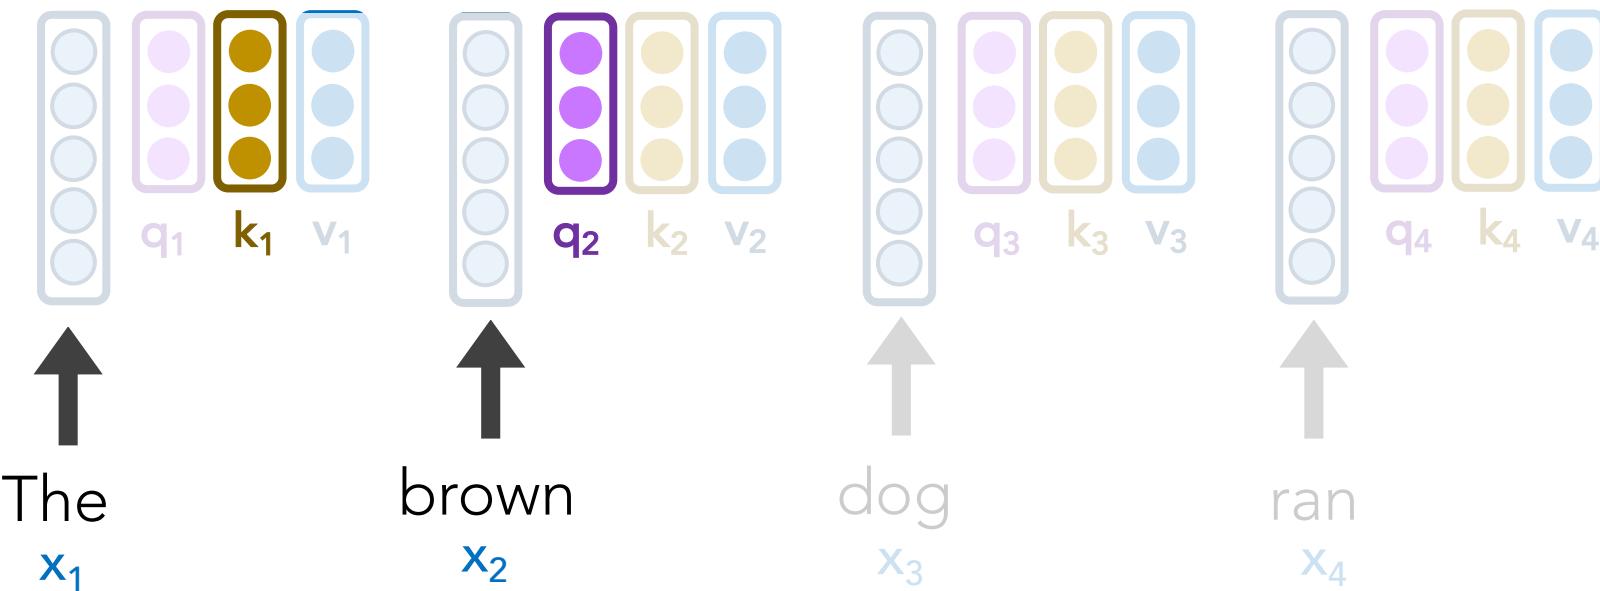
Under the hood, each x_i has 3 small, associated vectors. For example, x_1 has:

- Query q_1
- Key k_1
- Value v_1

Self-Attention

Step 2: For word x_2 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i

$$s_1 = q_2 \cdot k_1 = 92$$

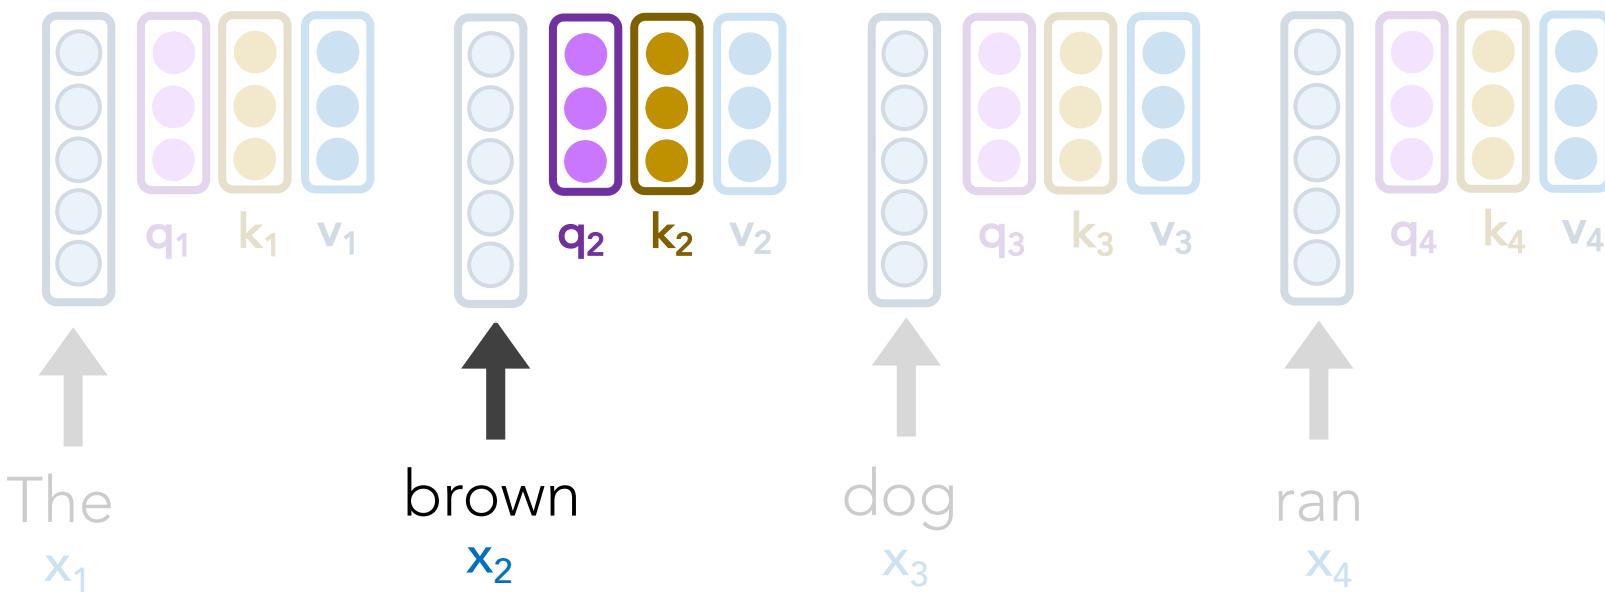


Self-Attention

Step 2: For word x_2 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i

$$s_2 = q_2 \cdot k_2 = 124$$

$$s_1 = q_2 \cdot k_1 = 92$$



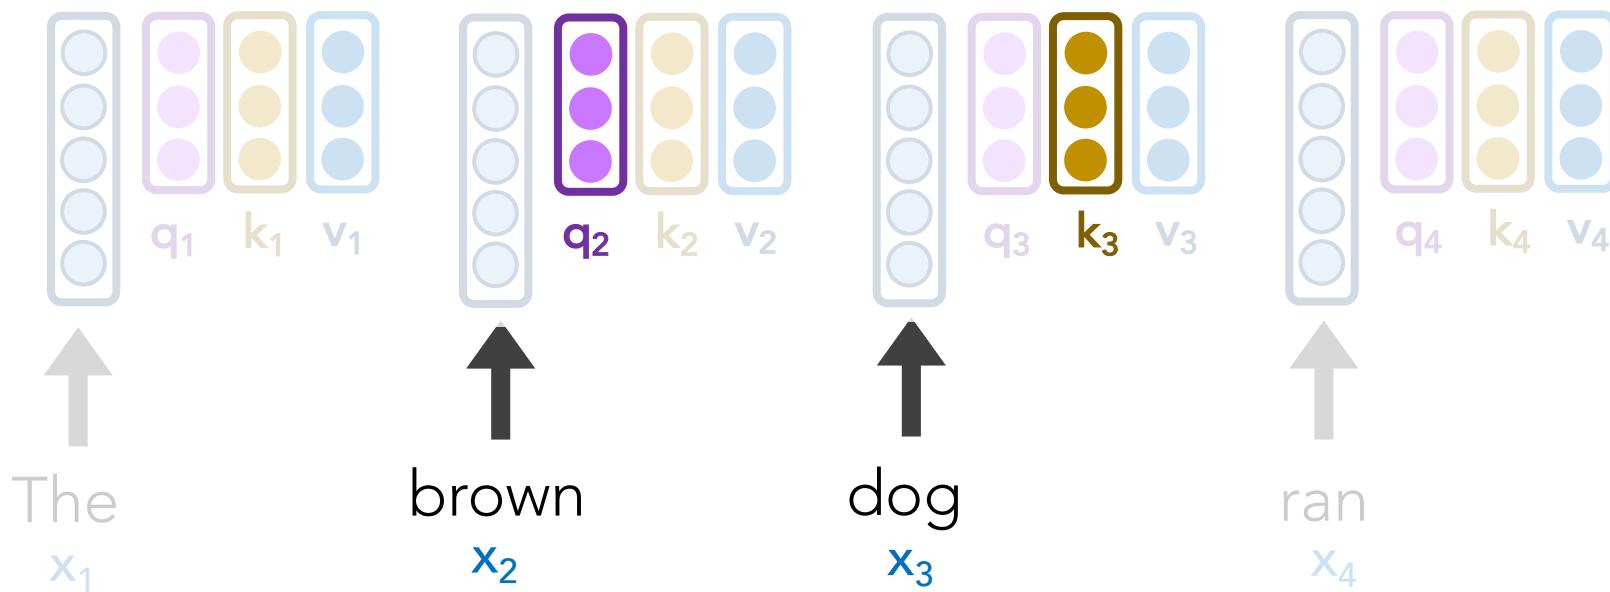
Self-Attention

Step 2: For word x_2 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i

$$s_3 = q_2 \cdot k_3 = 22$$

$$s_2 = q_2 \cdot k_2 = 124$$

$$s_1 = q_2 \cdot k_1 = 92$$



Self-Attention

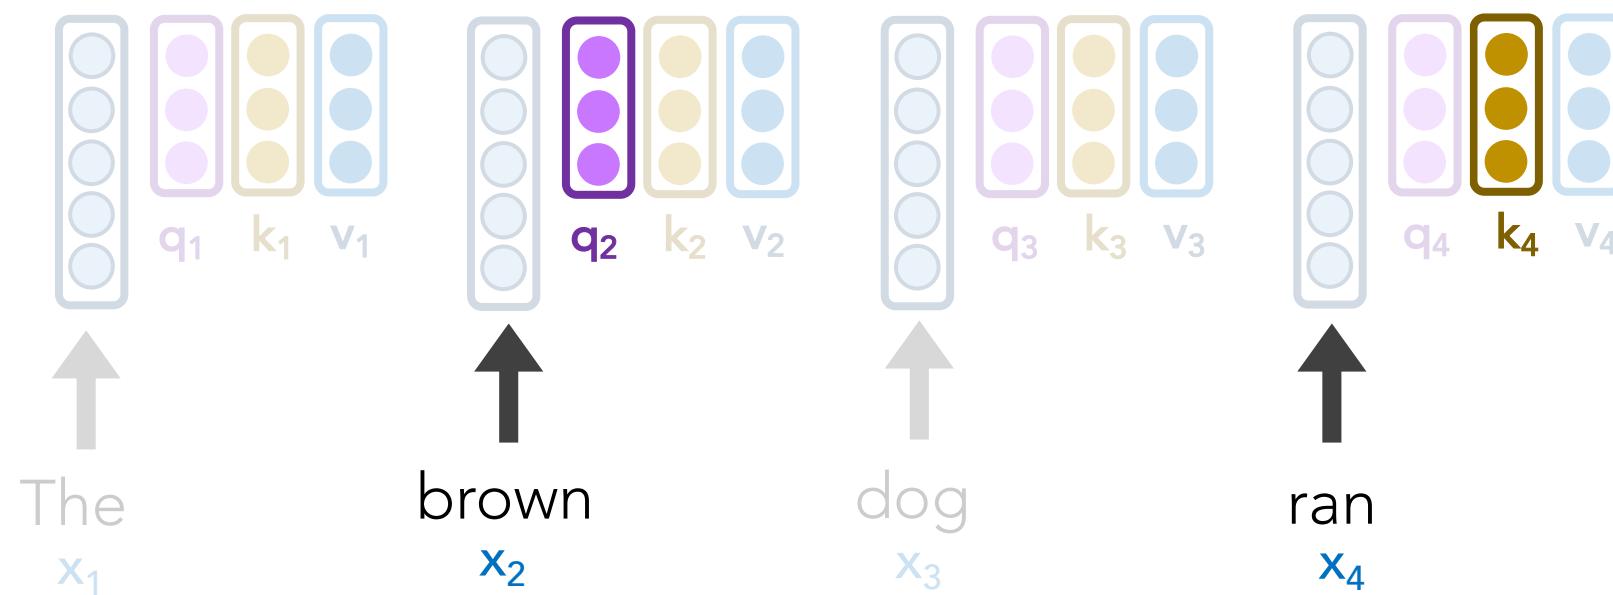
Step 2: For word x_2 , let's calculate the scores s_1, s_2, s_3, s_4 , which represent how much attention to pay to each respective "word" v_i

$$s_4 = q_2 \cdot k_4 = 8$$

$$s_3 = q_2 \cdot k_3 = 22$$

$$s_2 = q_2 \cdot k_2 = 124$$

$$s_1 = q_2 \cdot k_1 = 92$$



Self-Attention

Step 3: Our scores s_1, s_2, s_3, s_4 don't sum to 1. Let's divide by $\sqrt{\text{len}(k_i)}$ and **softmax** it

$$s_4 = q_2 \cdot k_4 = 8$$

$$a_4 = \sigma(s_4/8) = 0$$

$$s_3 = q_2 \cdot k_3 = 22$$

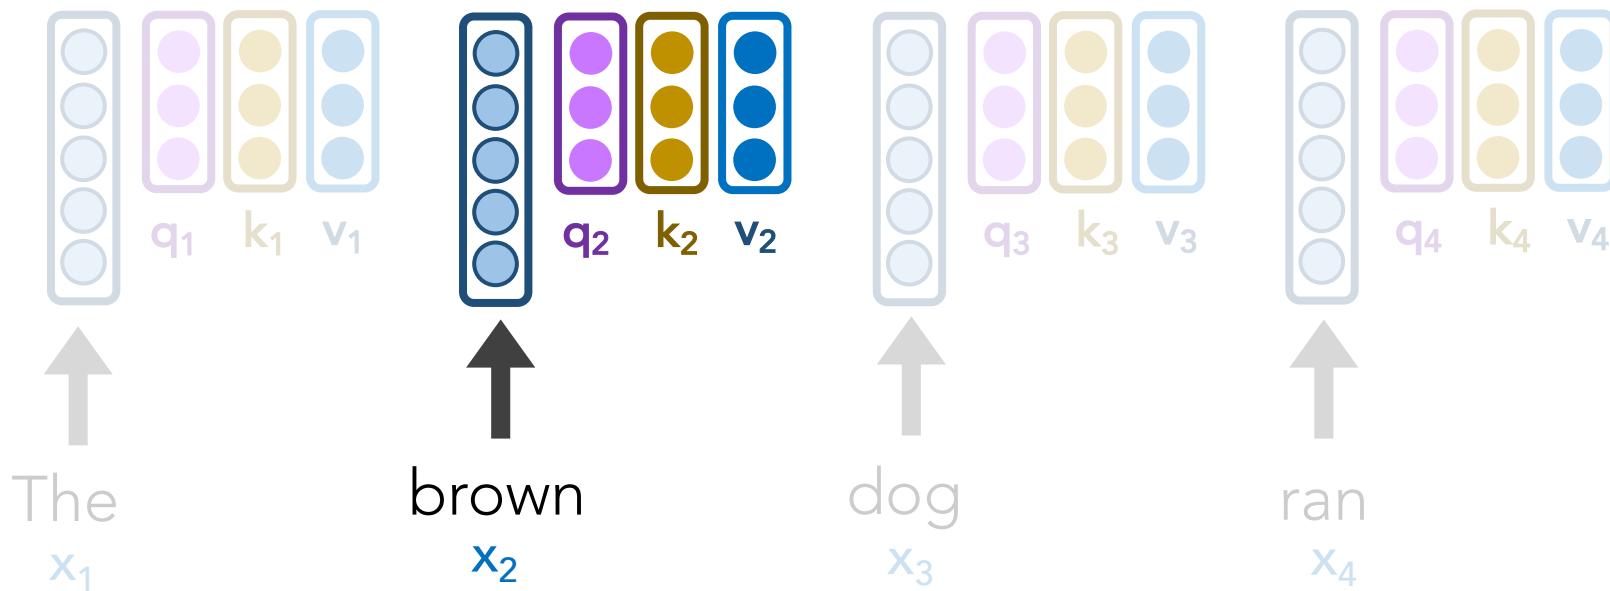
$$a_3 = \sigma(s_3/8) = .01$$

$$s_2 = q_2 \cdot k_2 = 124$$

$$a_2 = \sigma(s_2/8) = .91$$

$$s_1 = q_2 \cdot k_1 = 92$$

$$a_1 = \sigma(s_1/8) = .08$$



Self-Attention

Step 3: Our scores s_1, s_2, s_3, s_4 don't sum to 1. Let's divide by $\sqrt{\text{len}(k_i)}$ and softmax it

$$s_4 = q_2 \cdot k_4 = 8$$

$$a_4 = \sigma(s_4/8) = 0$$

$$s_3 = q_2 \cdot k_3 = 22$$

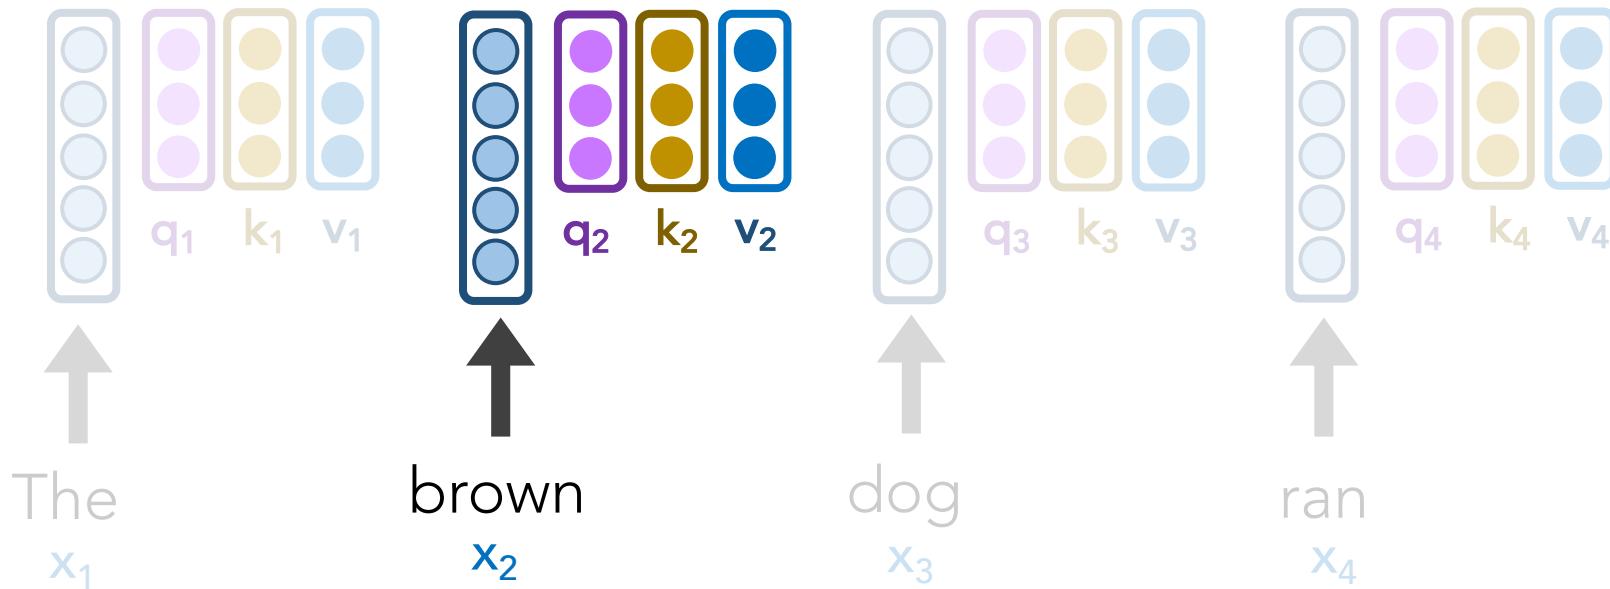
$$a_3 = \sigma(s_3/8) = .01$$

$$s_2 = q_2 \cdot k_2 = 124$$

$$a_2 = \sigma(s_2/8) = .91$$

$$s_1 = q_2 \cdot k_1 = 92$$

$$a_1 = \sigma(s_1/8) = .08$$



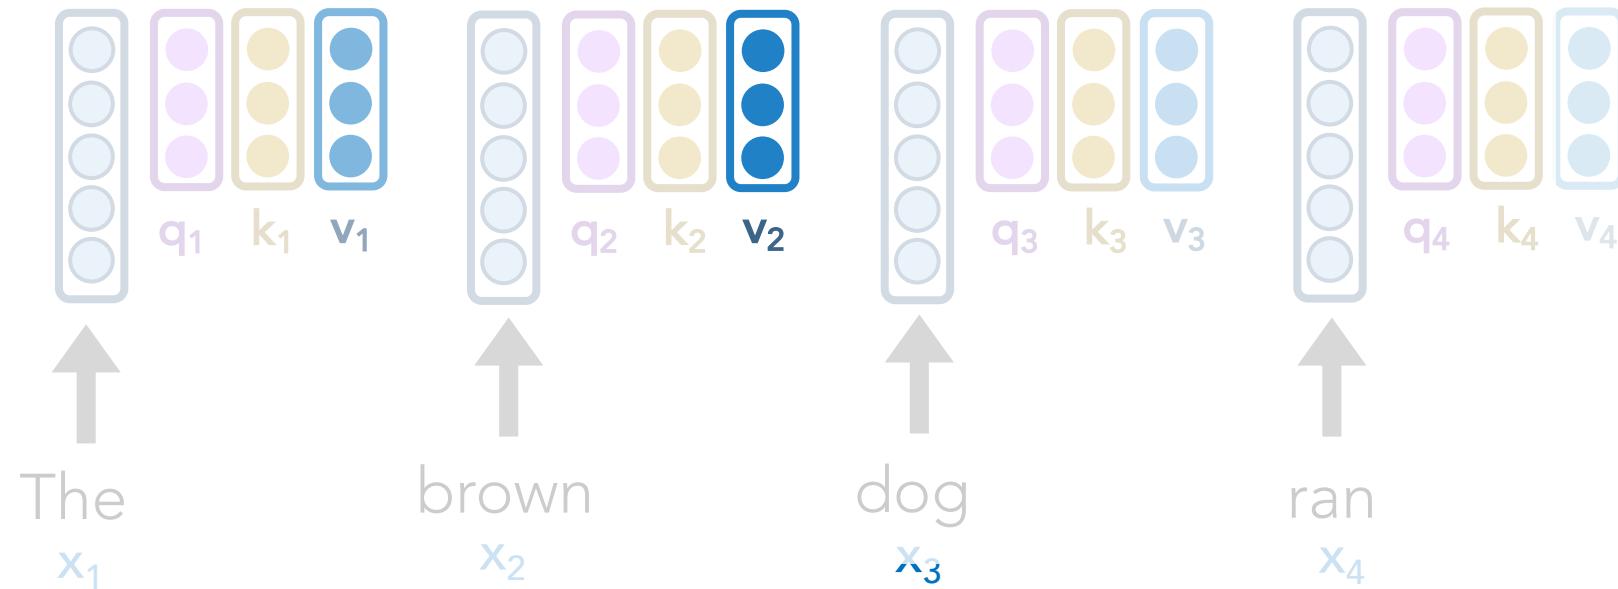
Instead of these a_i values directly weighting our original x_i word vectors, they directly weight our v_i vectors.

Self-Attention

Step 4: Let's weight our v_i vectors and simply sum them up!

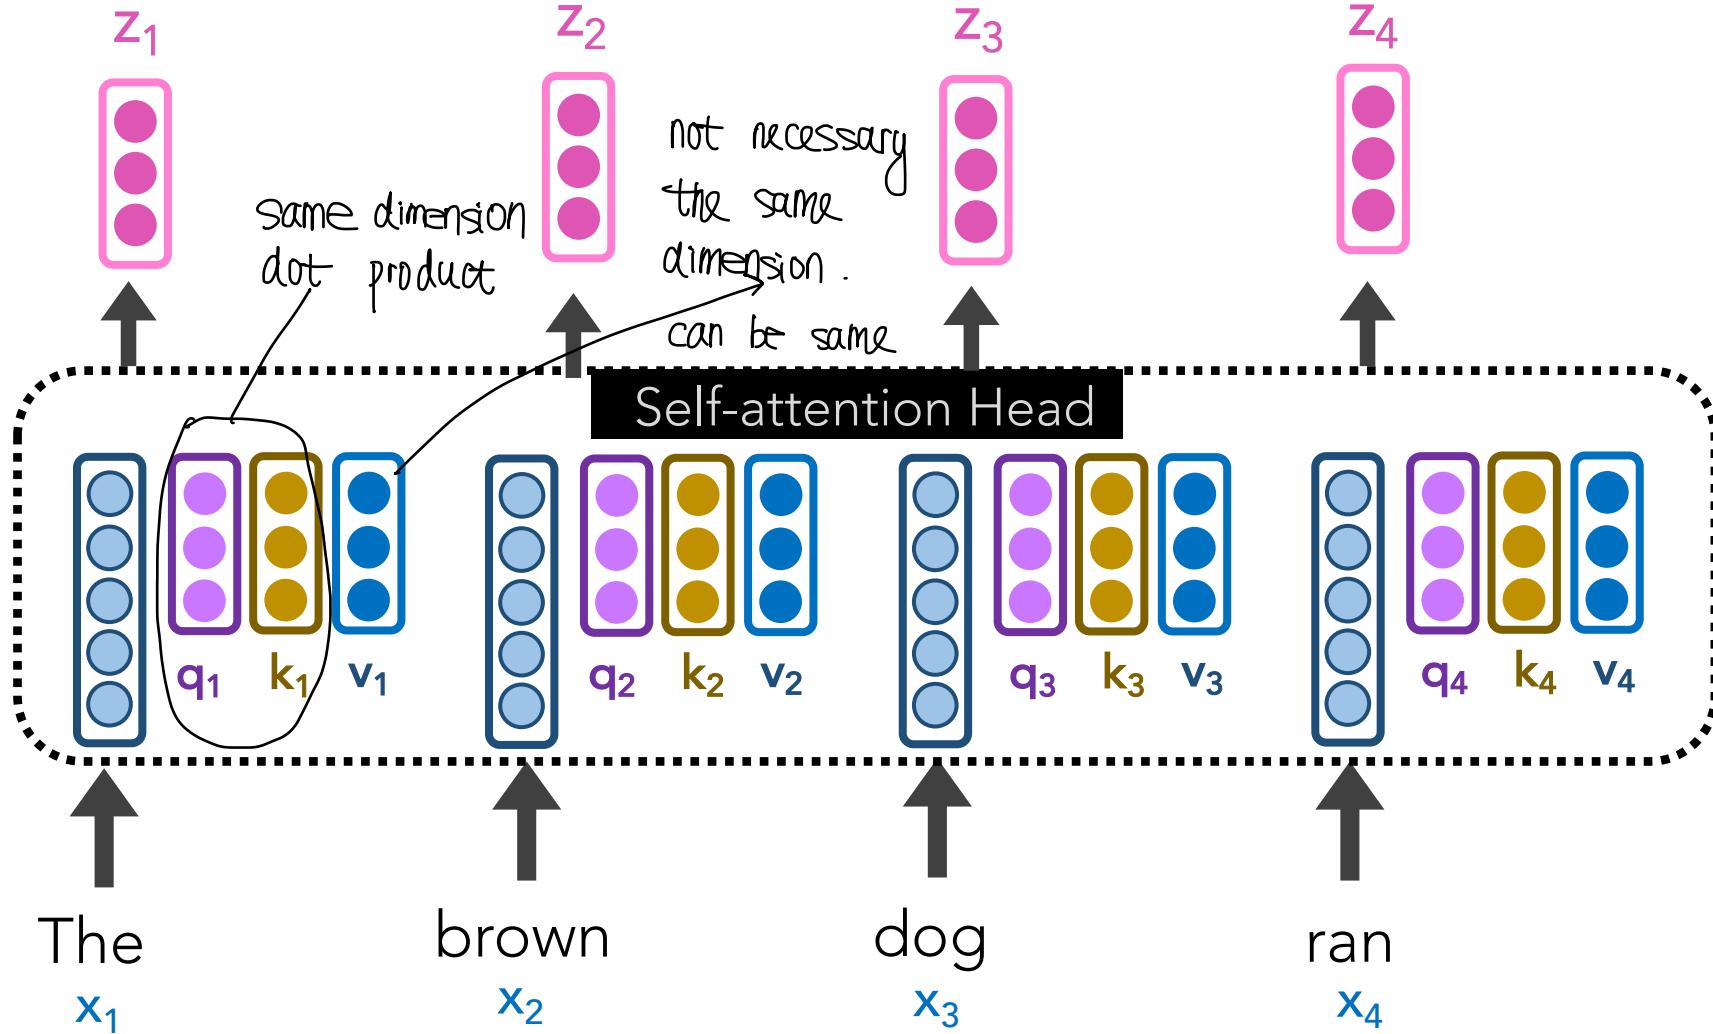
The diagram illustrates the final step of self-attention. On the left, four input vectors x_1 , x_2 , x_3 , and x_4 are shown, each with four components. These are processed by weight matrices q_1, k_1, v_1 and q_2, k_2, v_2 to produce query (q), key (k), and value (v) vectors respectively. The k and v vectors have three components each. An arrow points from the v vector of x_2 to a vertical stack of three circles labeled z_2 . To the right of this stack is the equation:

$$\begin{aligned} z_2 &= a_1 \cdot v_1 + a_2 \cdot v_2 + a_3 \cdot v_3 + a_4 \cdot v_4 \\ &= 0.08 \cdot v_1 + 0.91 \cdot v_2 + 0.01 \cdot v_3 + 0 \cdot v_4 \end{aligned}$$



Self-Attention

Tada! Now we have great, new representations z_i via a self-attention head



Self-Attention

Tada! Now

z



Takeaway:

Self-Attention is powerful; allows us to
create great, context-aware
representations



$q_1 \ k_1 \ v_1$



The
 x_1

brown
 x_2

dog
 x_3

ran
 x_4

$k_2 \ v_2$

$q_3 \ k_3 \ v_3$

$q_4 \ k_4 \ v_4$

隨
意

Self-Attention may seem strikingly
like **Attention** in seq2seq models

$$s_4 = h_1^D * h_4^E$$

$$a_4 = \sigma(s_4)$$

$$s_3 = h_1^D * h_3^E$$

$$a_3 = \sigma(s_3)$$

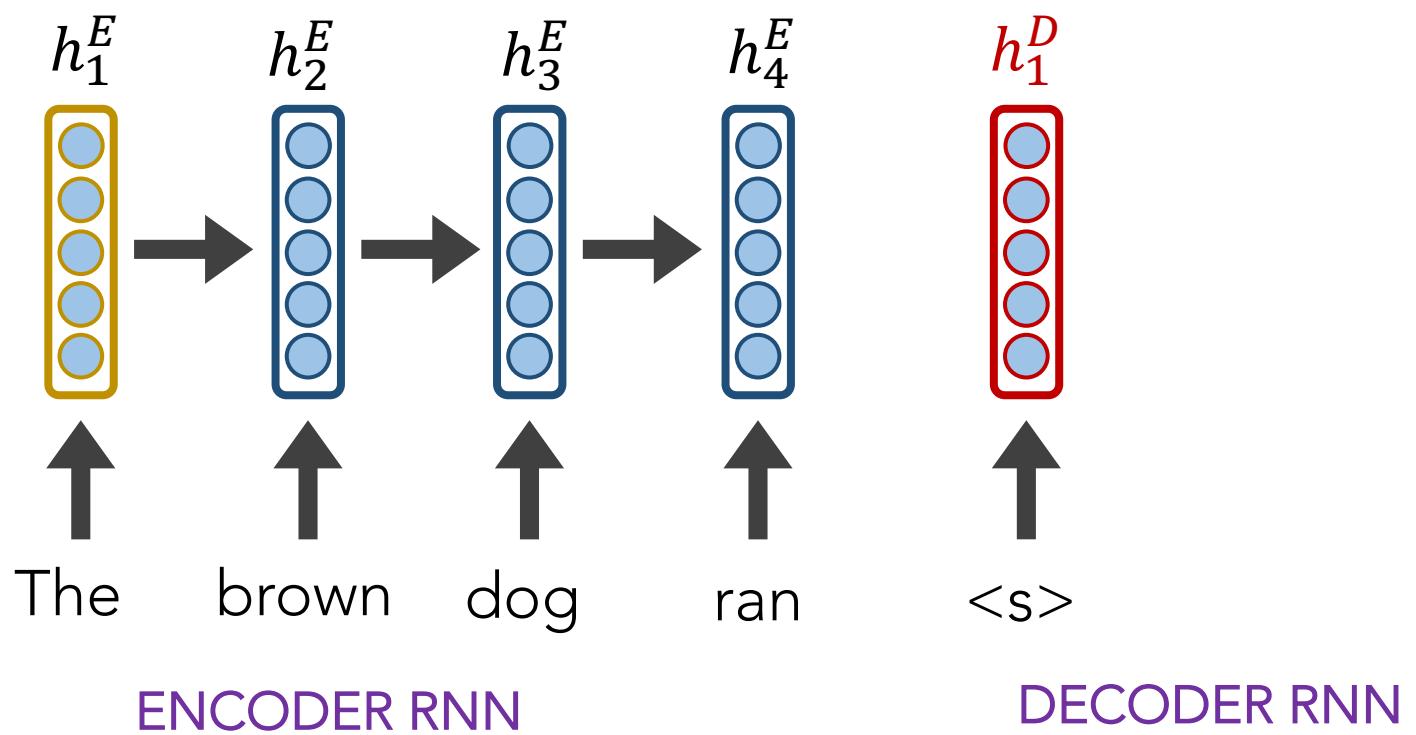
$$s_2 = h_1^D * h_2^E$$

$$a_2 = \sigma(s_2)$$

$$s_1 = h_1^D * h_1^E$$

$$a_1 = \sigma(s_1)$$

Attention



$$s_4 = h_1^D * h_4^E$$

$$a_4 = \sigma(s_4)$$

$$s_3 = h_1^D * h_3^E$$

$$a_3 = \sigma(s_3)$$

$$s_2 = h_1^D * h_2^E$$

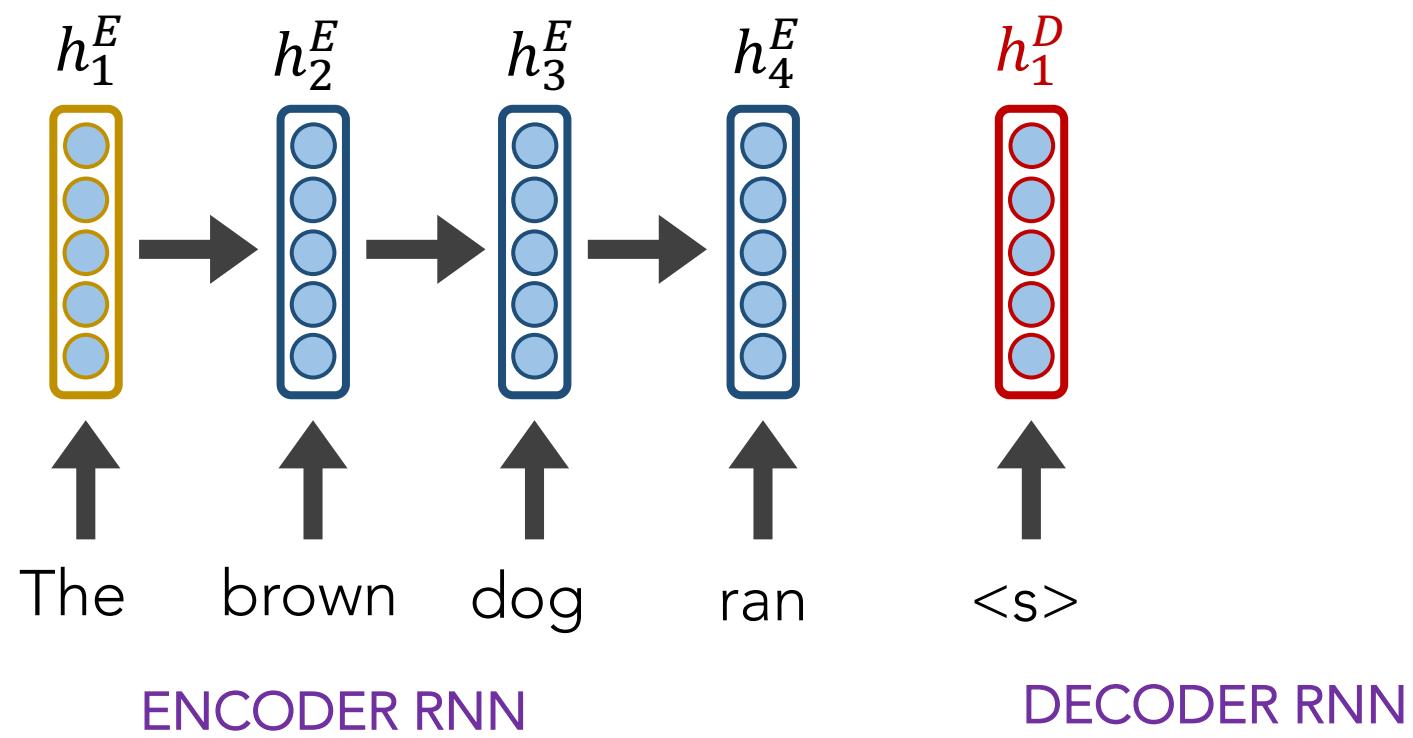
$$a_2 = \sigma(s_2)$$

$$s_1 = h_1^D * h_1^E$$

$$a_1 = \sigma(s_1)$$

We multiply each encoder's hidden layer by its a_i^1 attention weights to create a context vector c_1^D

Attention



$$s_4 = h_1^D * h_4^E$$

$$a_4 = \sigma(s_4)$$

$$s_3 = h_1^D * h_3^E$$

$$a_3 = \sigma(s_3)$$

$$s_2 = h_1^D * h_2^E$$

$$a_2 = \sigma(s_2)$$

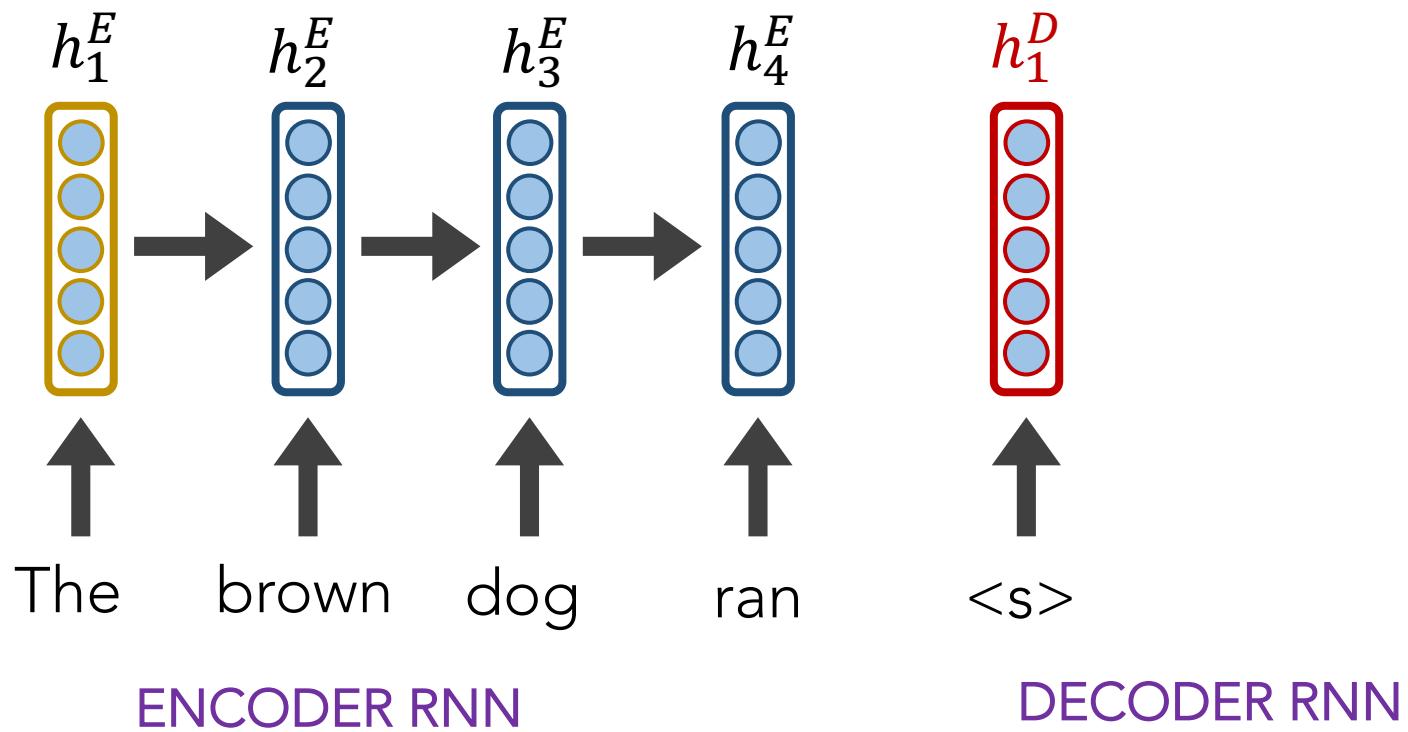
$$s_1 = h_1^D * h_1^E$$

$$a_1 = \sigma(s_1)$$

We multiply each encoder's hidden layer by its a_i^1 attention weights to create a context vector c_1^D

$$c_1^D = a_1 \cdot h_1^E + a_2 \cdot h_2^E + a_3 \cdot h_3^E + a_4 \cdot h_4^E$$

Attention



$$s_4 = \mathbf{q}_2 \cdot \mathbf{k}_4$$

$$a_4 = \sigma(s_4/8)$$

$$s_3 = \mathbf{q}_2 \cdot \mathbf{k}_3$$

$$a_3 = \sigma(s_3/8)$$

$$s_2 = \mathbf{q}_2 \cdot \mathbf{k}_2$$

$$a_2 = \sigma(s_2/8)$$

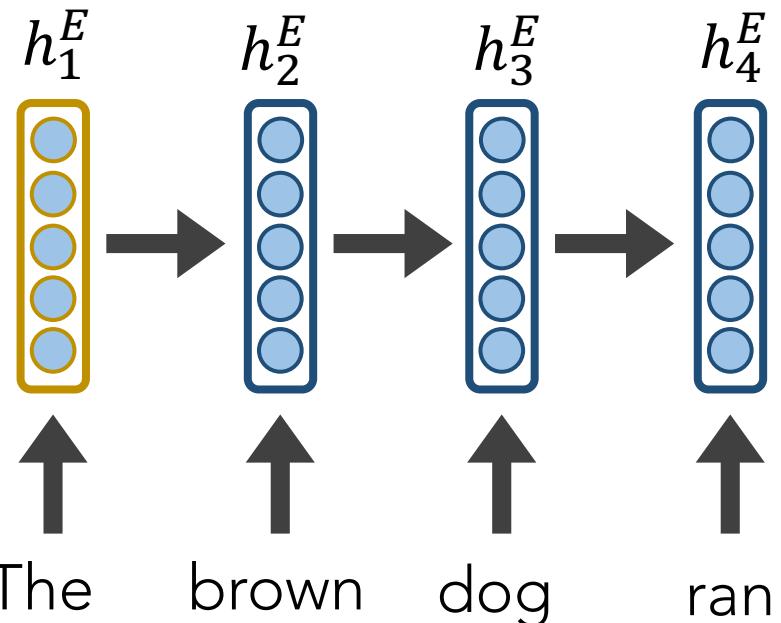
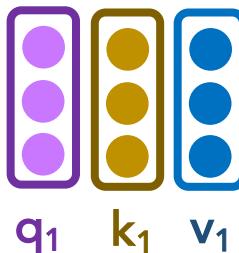
$$s_1 = \mathbf{q}_2 \cdot \mathbf{k}_1$$

$$a_1 = \sigma(s_1/8)$$

We multiply each word's value vector by its a_i^1 attention weights to create a better vector z_1

$$z_1 = a_1 \cdot v_1^E + a_2 \cdot v_2^E + a_3 \cdot v_3^E + a_4 \cdot v_4^E$$

Self-Attention



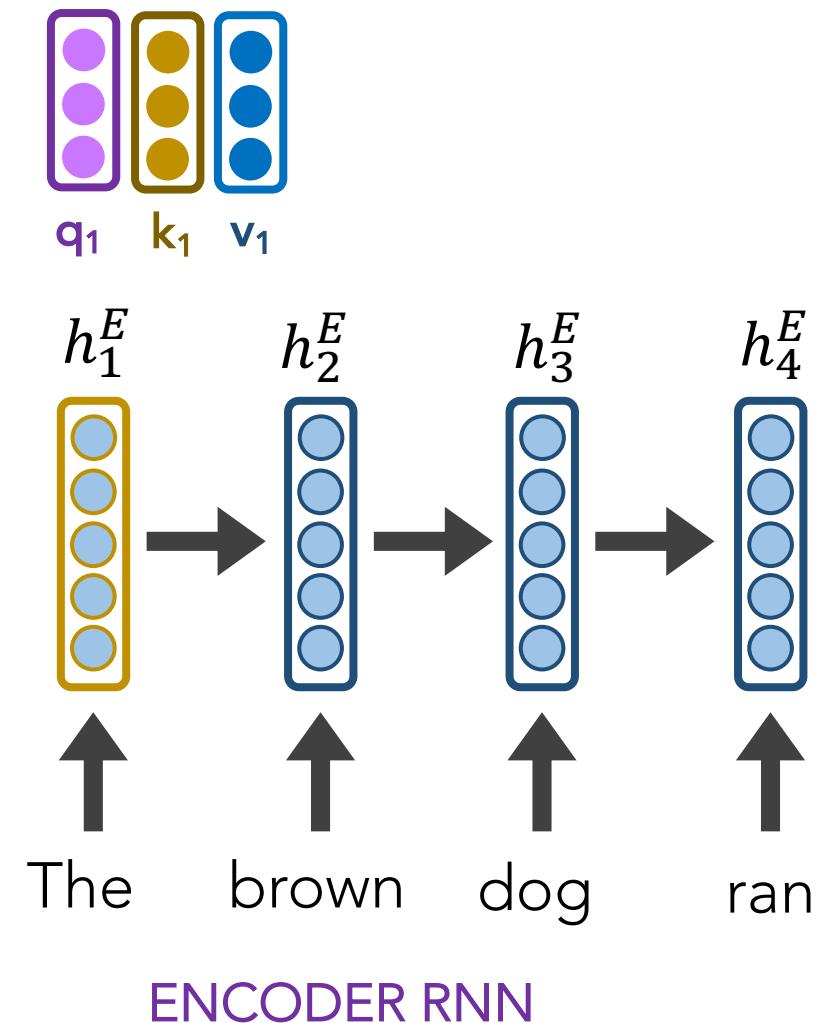
ENCODER RNN

Self-Attention

Self-Attention	Attention	Description
q_i	h_i^D	the probe
k_i	h_i^E	item being compared
v_i	h_i^E	item being weighted

vector by its a_i^1 attention weights to
create a better vector z_1

$$z_1 = a_1 \cdot v_1^E + a_2 \cdot v_2^E + a_3 \cdot v_3^E + a_4 \cdot v_4^E$$

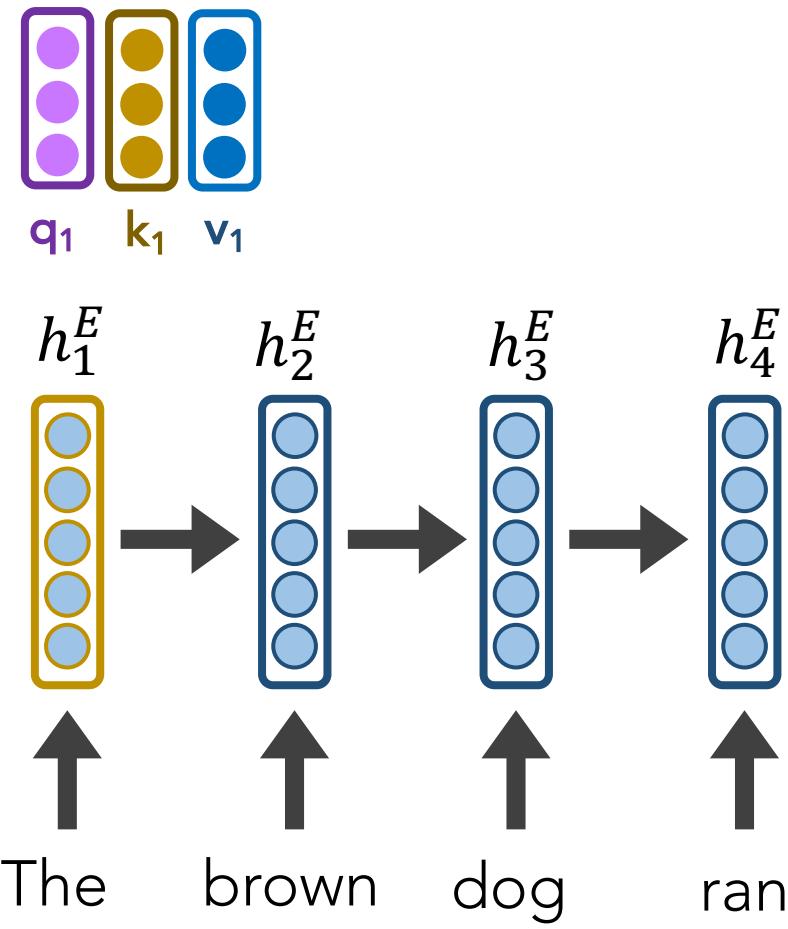


Self-Attention	Attention	Description
q_i	h_i^D	the probe
k_i	h_i^E	item being compared
v_i	h_i^E	item being weighted

All of these are like surrogates/proxies/abstractions.
 替代,代理 抽象,抽取
 This provides **flexibility** and fewer constraints.

More room for rich **abstractions**.

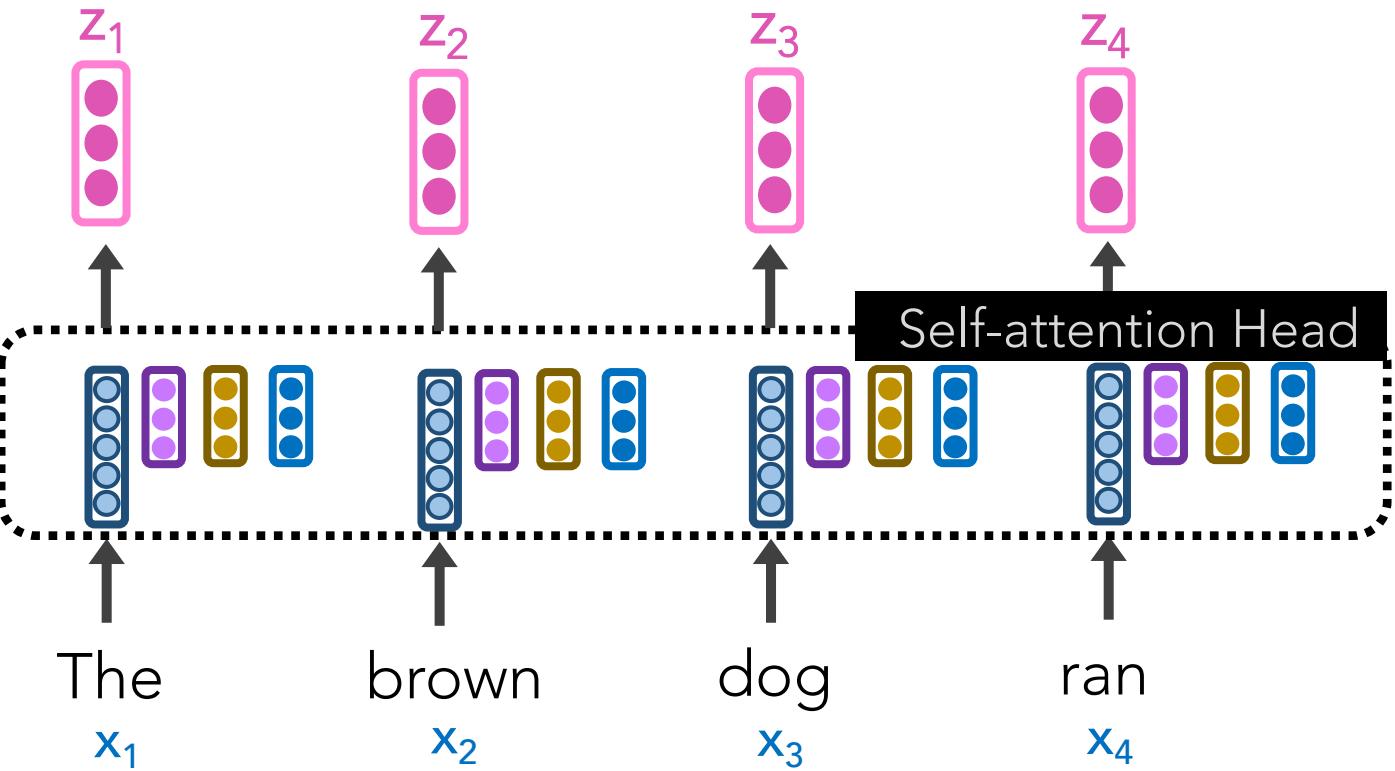
Self-Attention



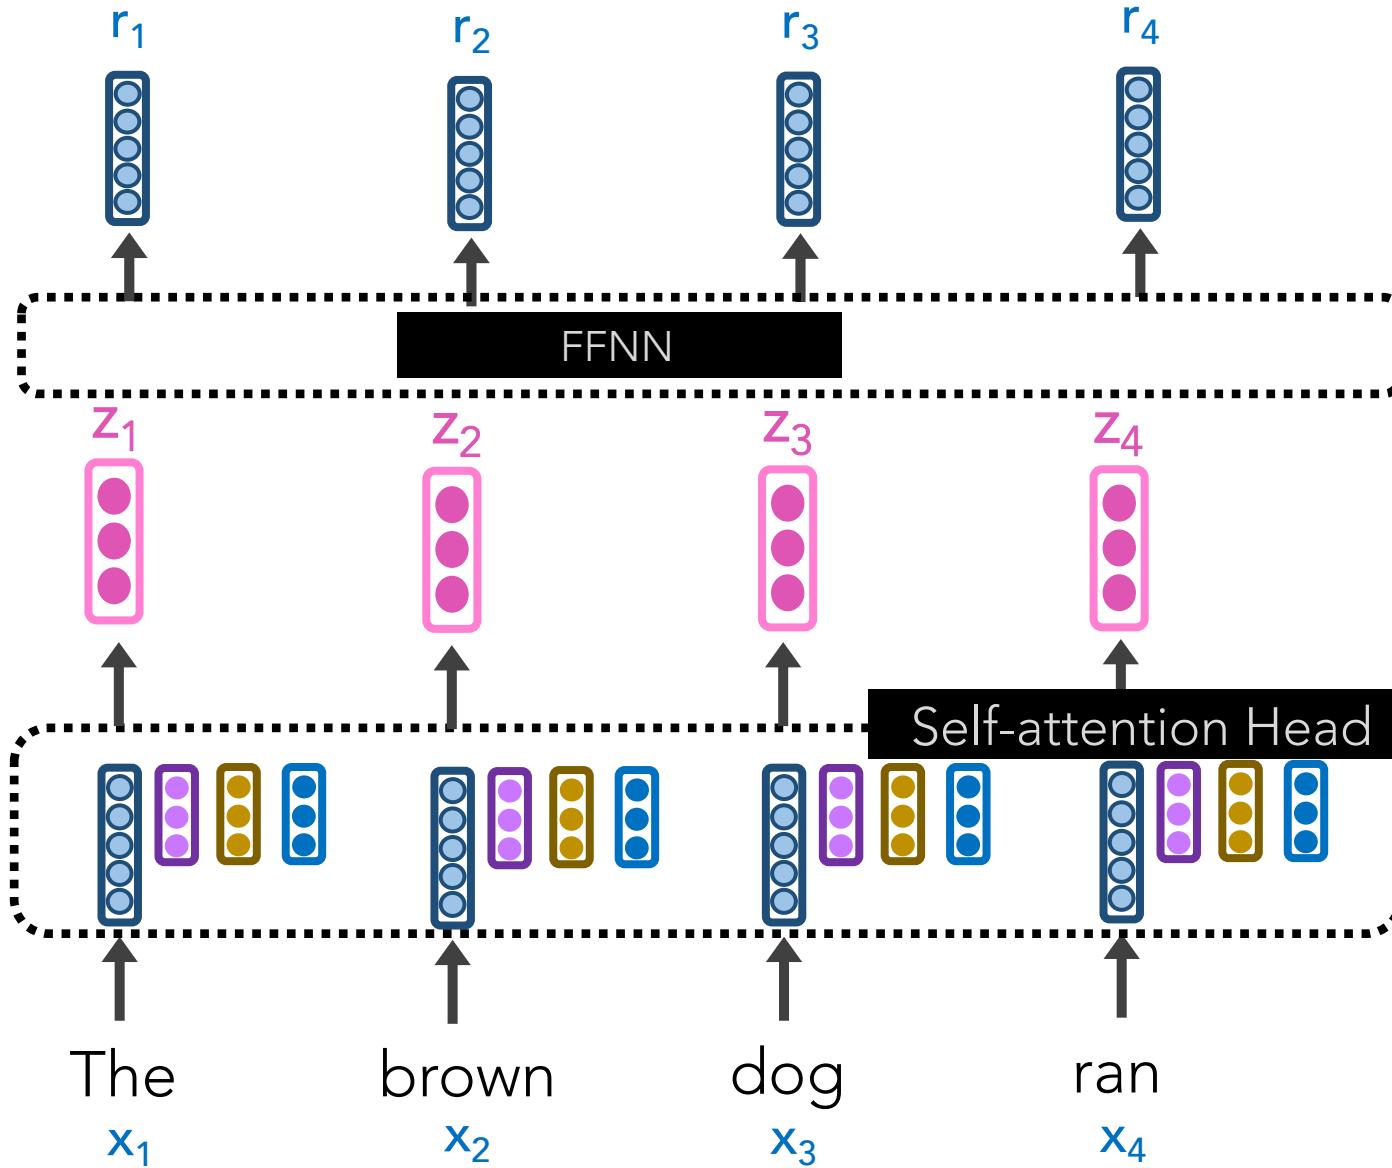
ENCODER RNN

Self-Attention

Let's further pass each z_i through a FFNN

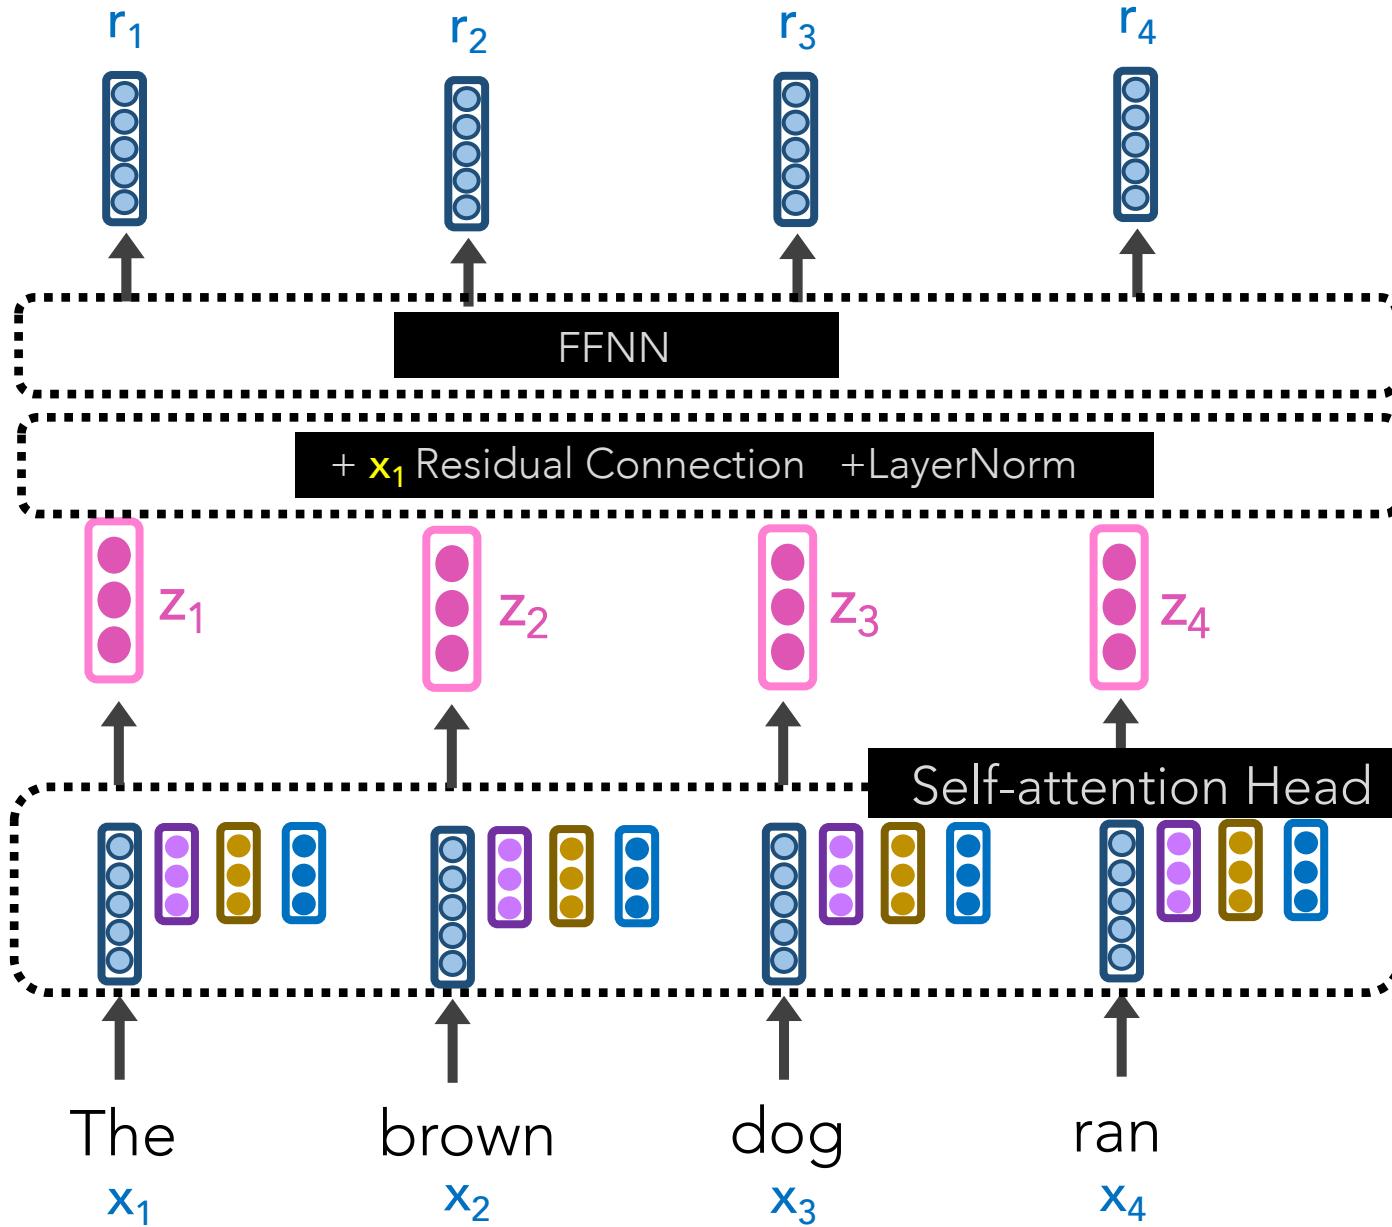


Self-Attention + FFNN



Let's further pass each z_i through a FFNN

Self-Attention + FFNN

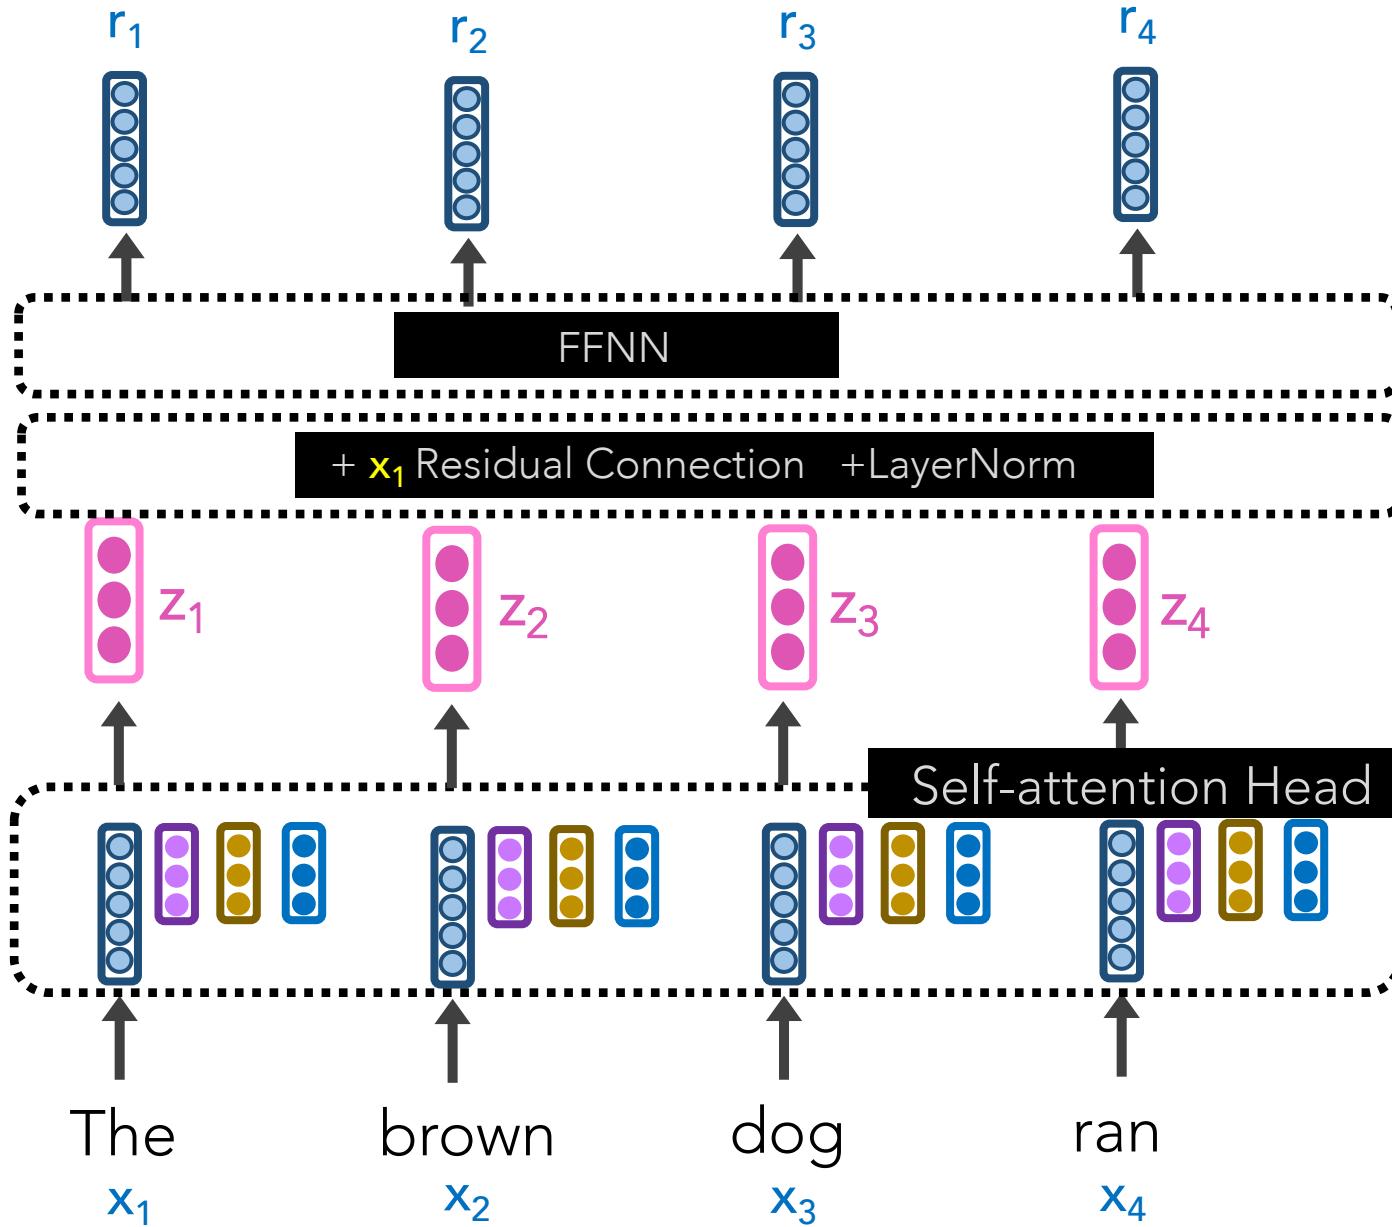


Let's further pass each z_i through a FFNN

We concat w/ a **residual connection** to help ensure relevant info is getting forward passed.

We perform **LayerNorm** to stabilize the network and allow for proper gradient flow.

Self-Attention + FFNN



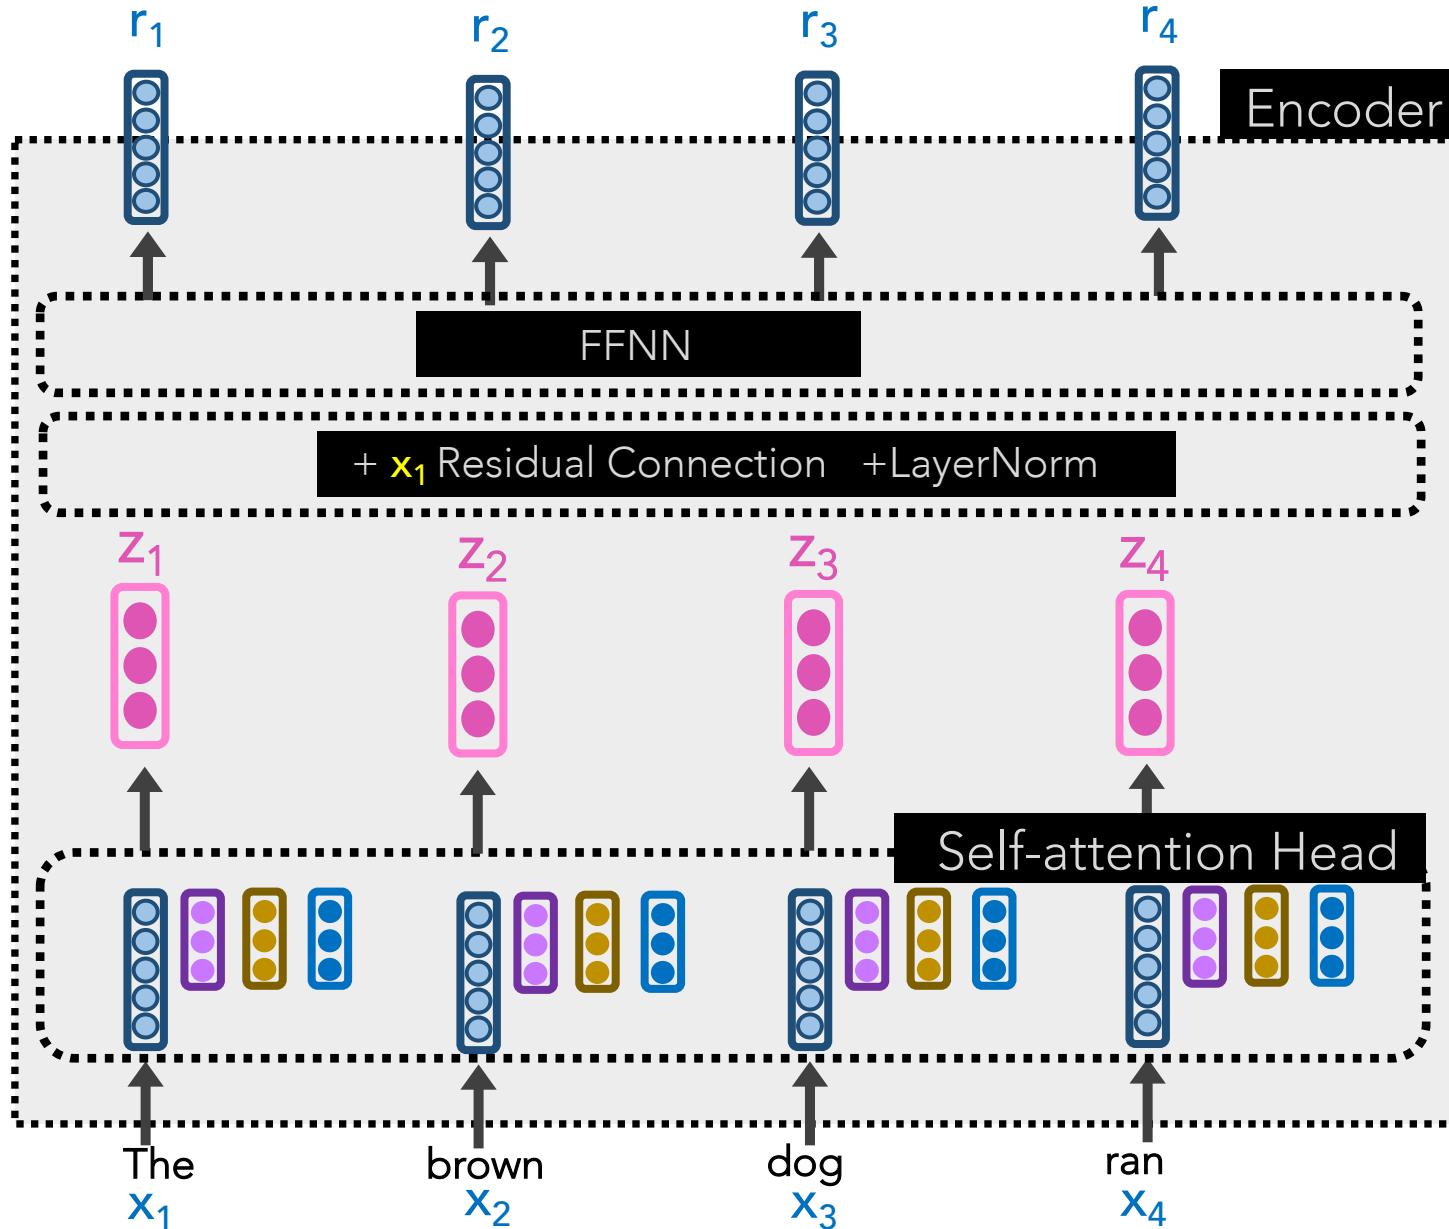
Let's further pass each z_i through a FFNN

We concat w/ a **residual connection** to help ensure relevant info is getting forward passed.

We perform **LayerNorm** to stabilize the network and allow for proper gradient flow.

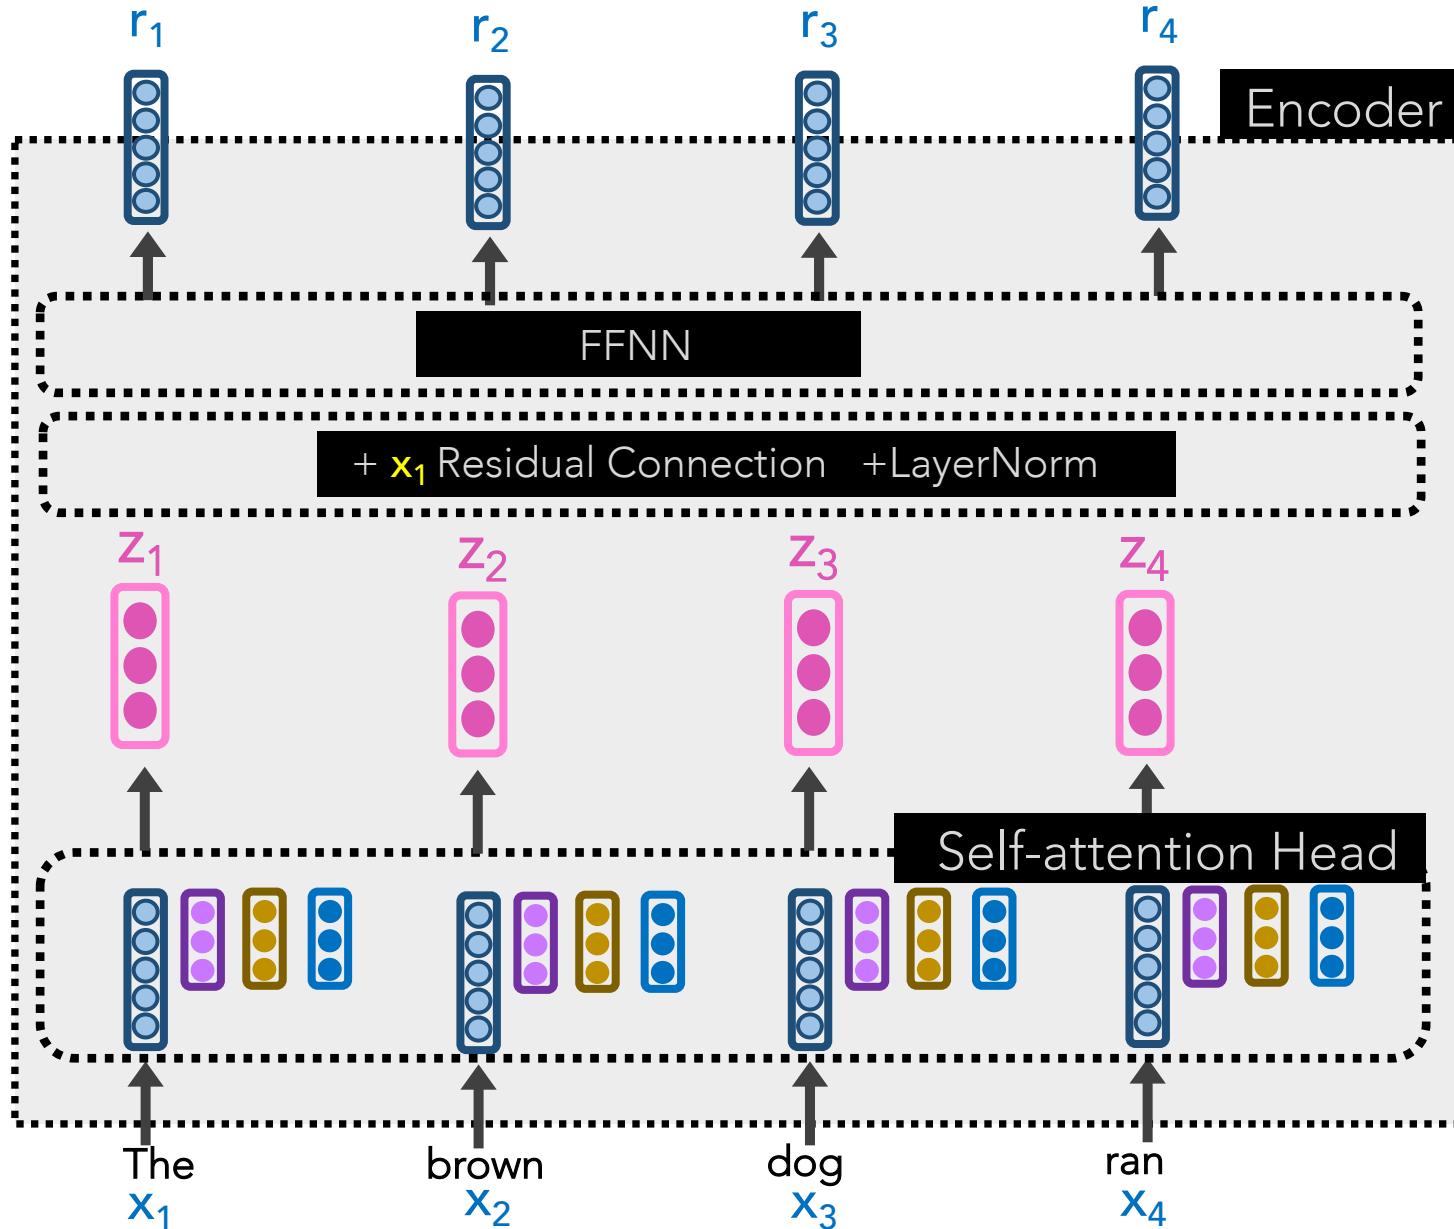
Each z_i can be computed in **parallel**, unlike LSTMs!

Transformer Encoder



Yay! Our r_i vectors are our new representations, and this entire process is called a **Transformer Encoder**

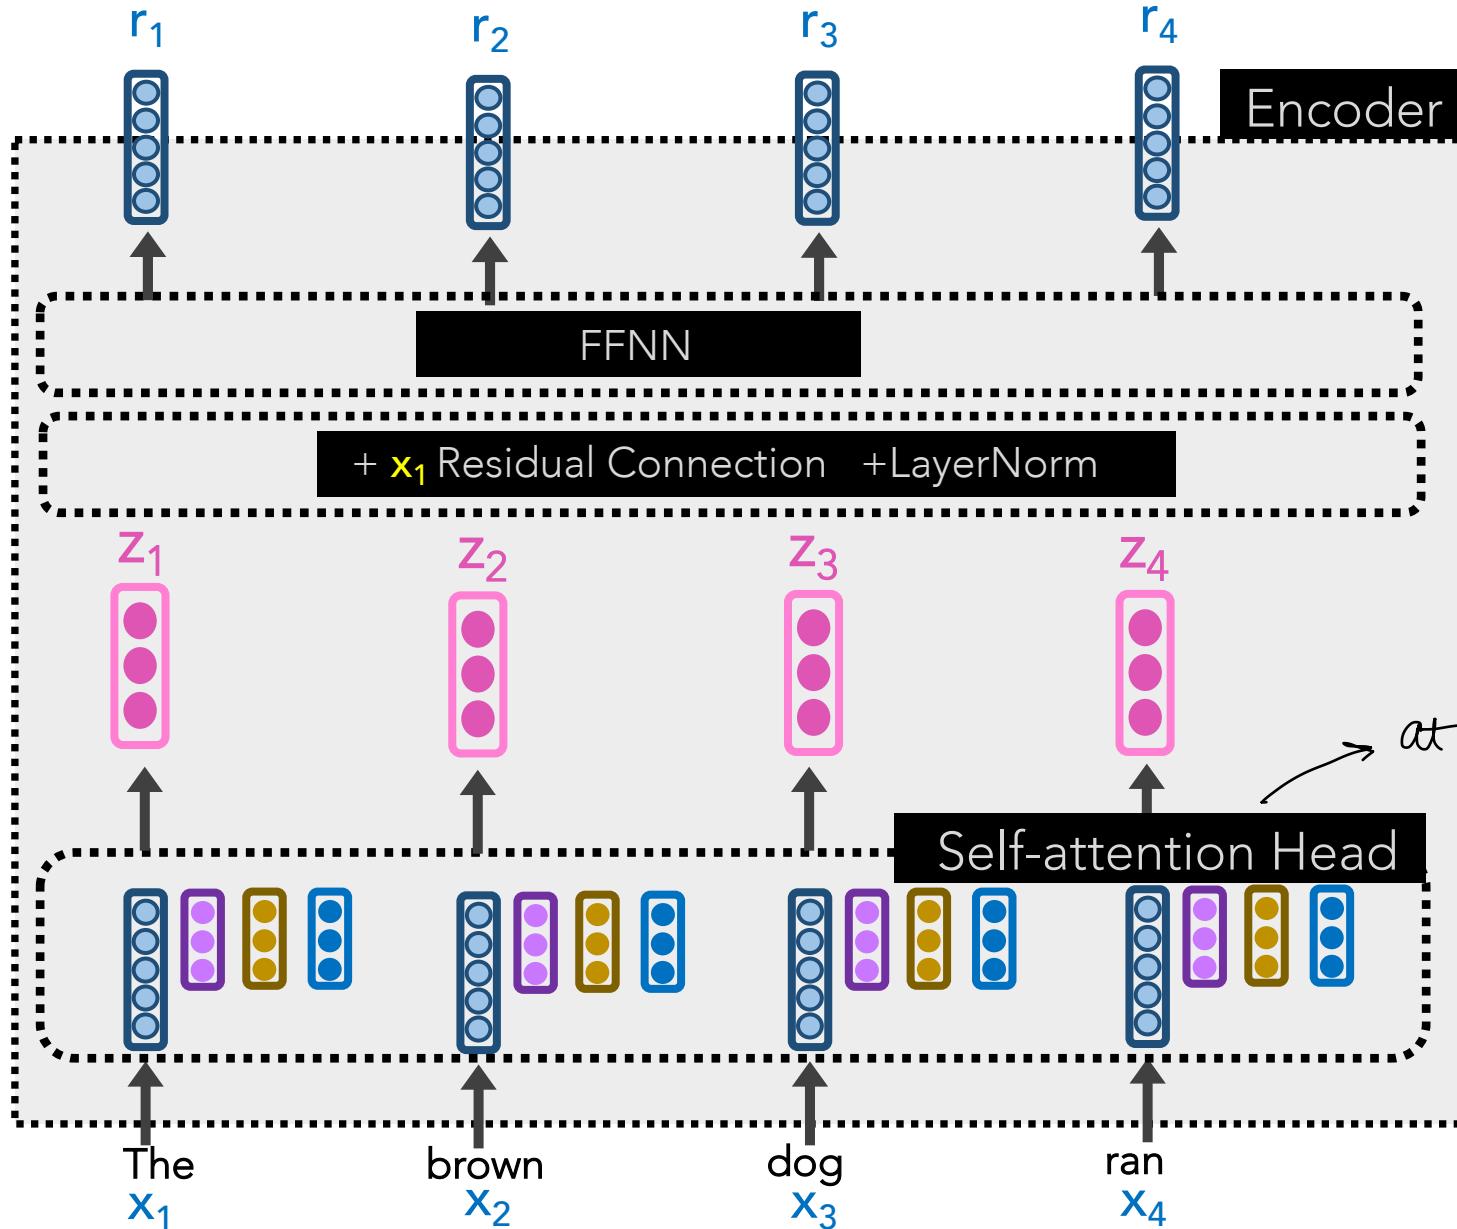
Transformer Encoder



Yay! Our r_i vectors are our new representations, and this entire process is called a **Transformer Encoder**

Problem: there is no concept of positionality. Words are weighted as if a “bag of words”

Transformer Encoder



Yay! Our r_i vectors are our new representations, and this entire process is called a **Transformer Encoder**

Problem: there is no concept of positionality. Words are weighted as if a “bag of words”

beginning - they are not the same
Solution: append each input word x_i with a **positional encoding**: $\sin(i) \cos(i)$

↑
original paper
positional encoding
natural moving value

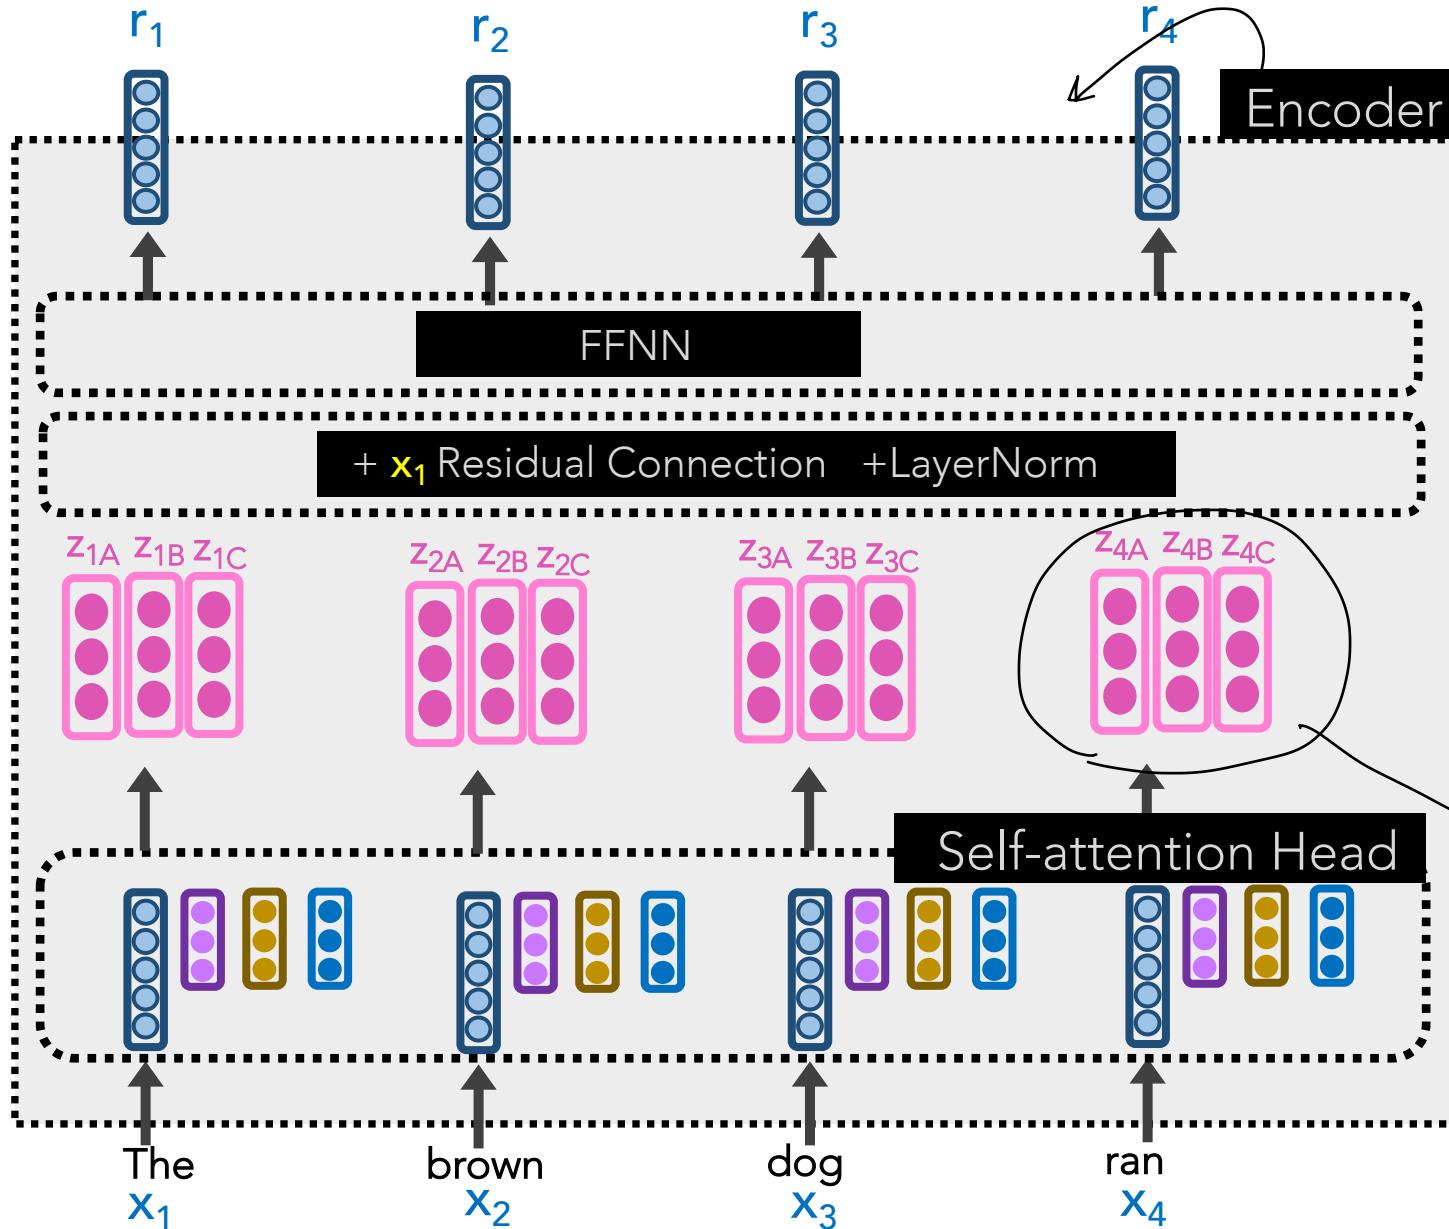
A **Self-Attention Head** has just one set of query/key/value weight matrices w_q , w_k , w_v

one head means three matrices
 w_q , w_k , w_v

Words can relate in many ways, so it's restrictive to rely on just one Self-Attention Head in the system.

Let's create Multi-headed Self-Attention
(do these multiple times)

Transformer Encoder

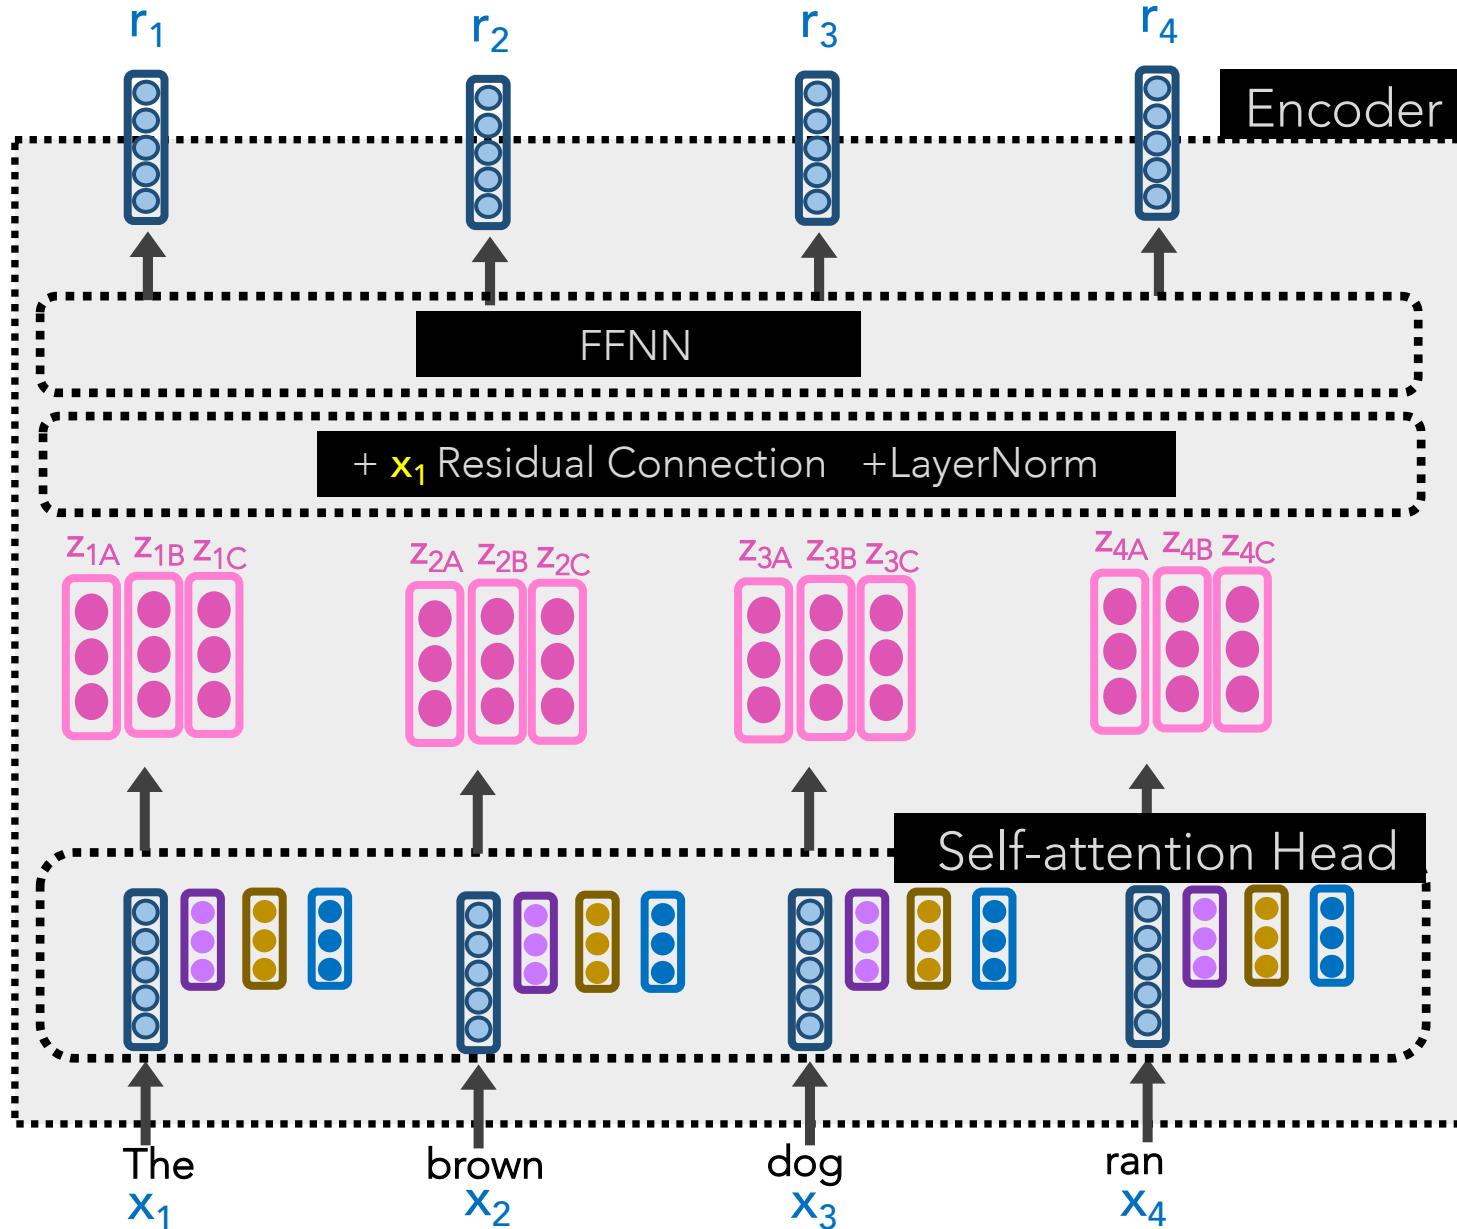


Each **Self-Attention Head** produces a z_i vector.

We can, in parallel, use **multiple heads** and concat the z_i 's.

Concatenate together
(explicitly learning a separate representation)

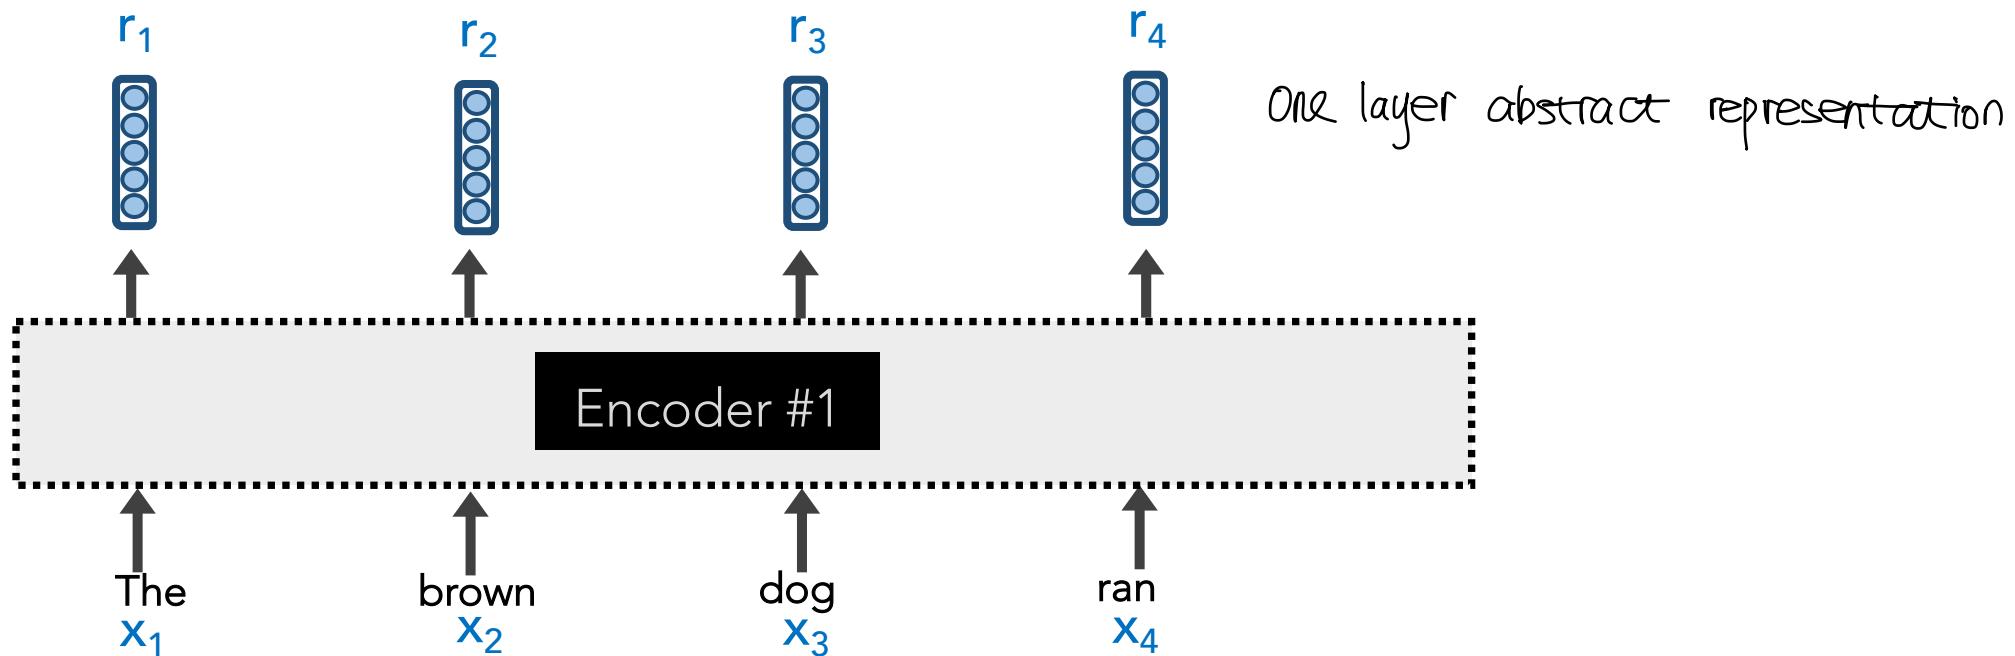
Transformer Encoder



To recap: all of this looks fancy, but ultimately it's just producing a very good **contextualized embedding** r_i of each word x_i

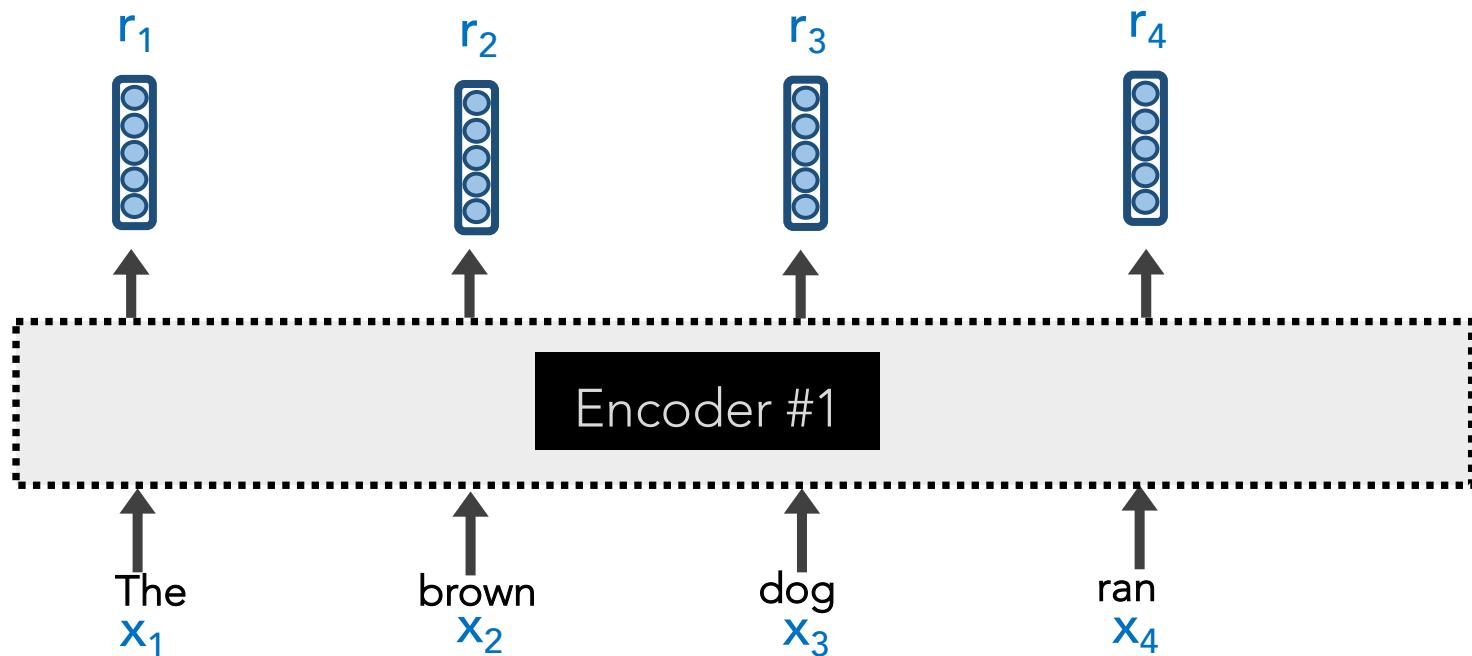
Transformer Encoder

To recap: all of this looks fancy, but ultimately it's just producing a very good **contextualized embedding** r_i of each word x_i



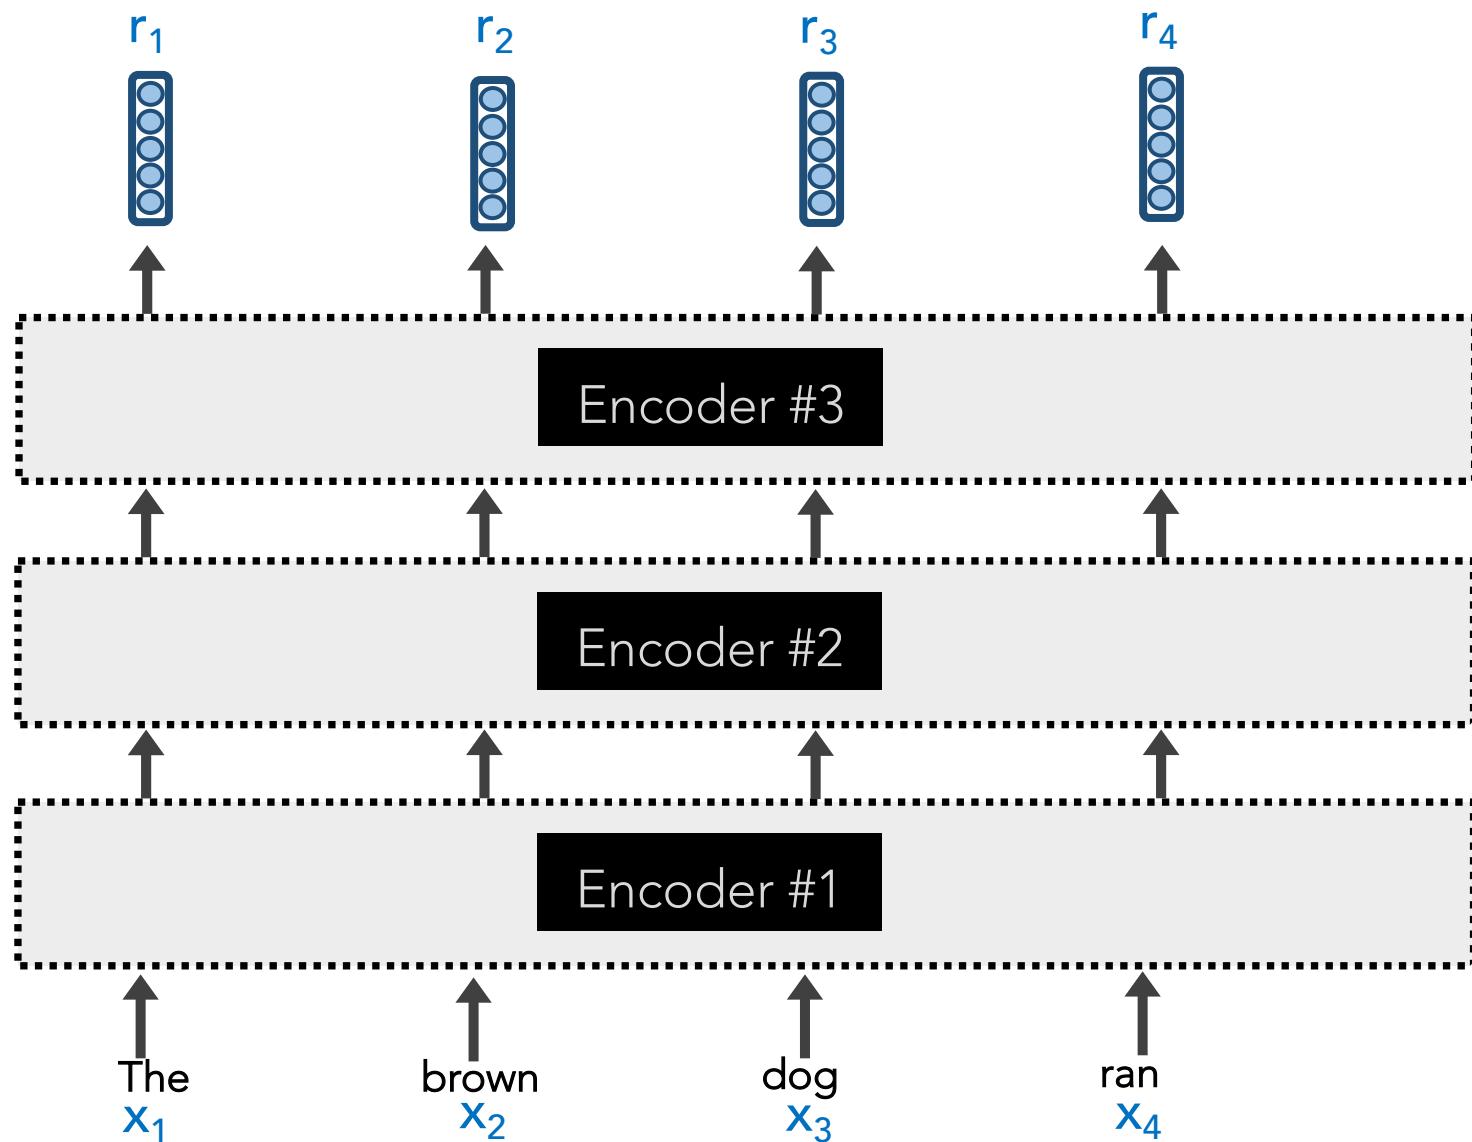
Transformer Encoder

To recap: all of this looks fancy, but ultimately it's just producing a very good **contextualized embedding** r_i of each word x_i



Why stop with just 1
Transformer Encoder?
We could stack several!

Transformer Encoder

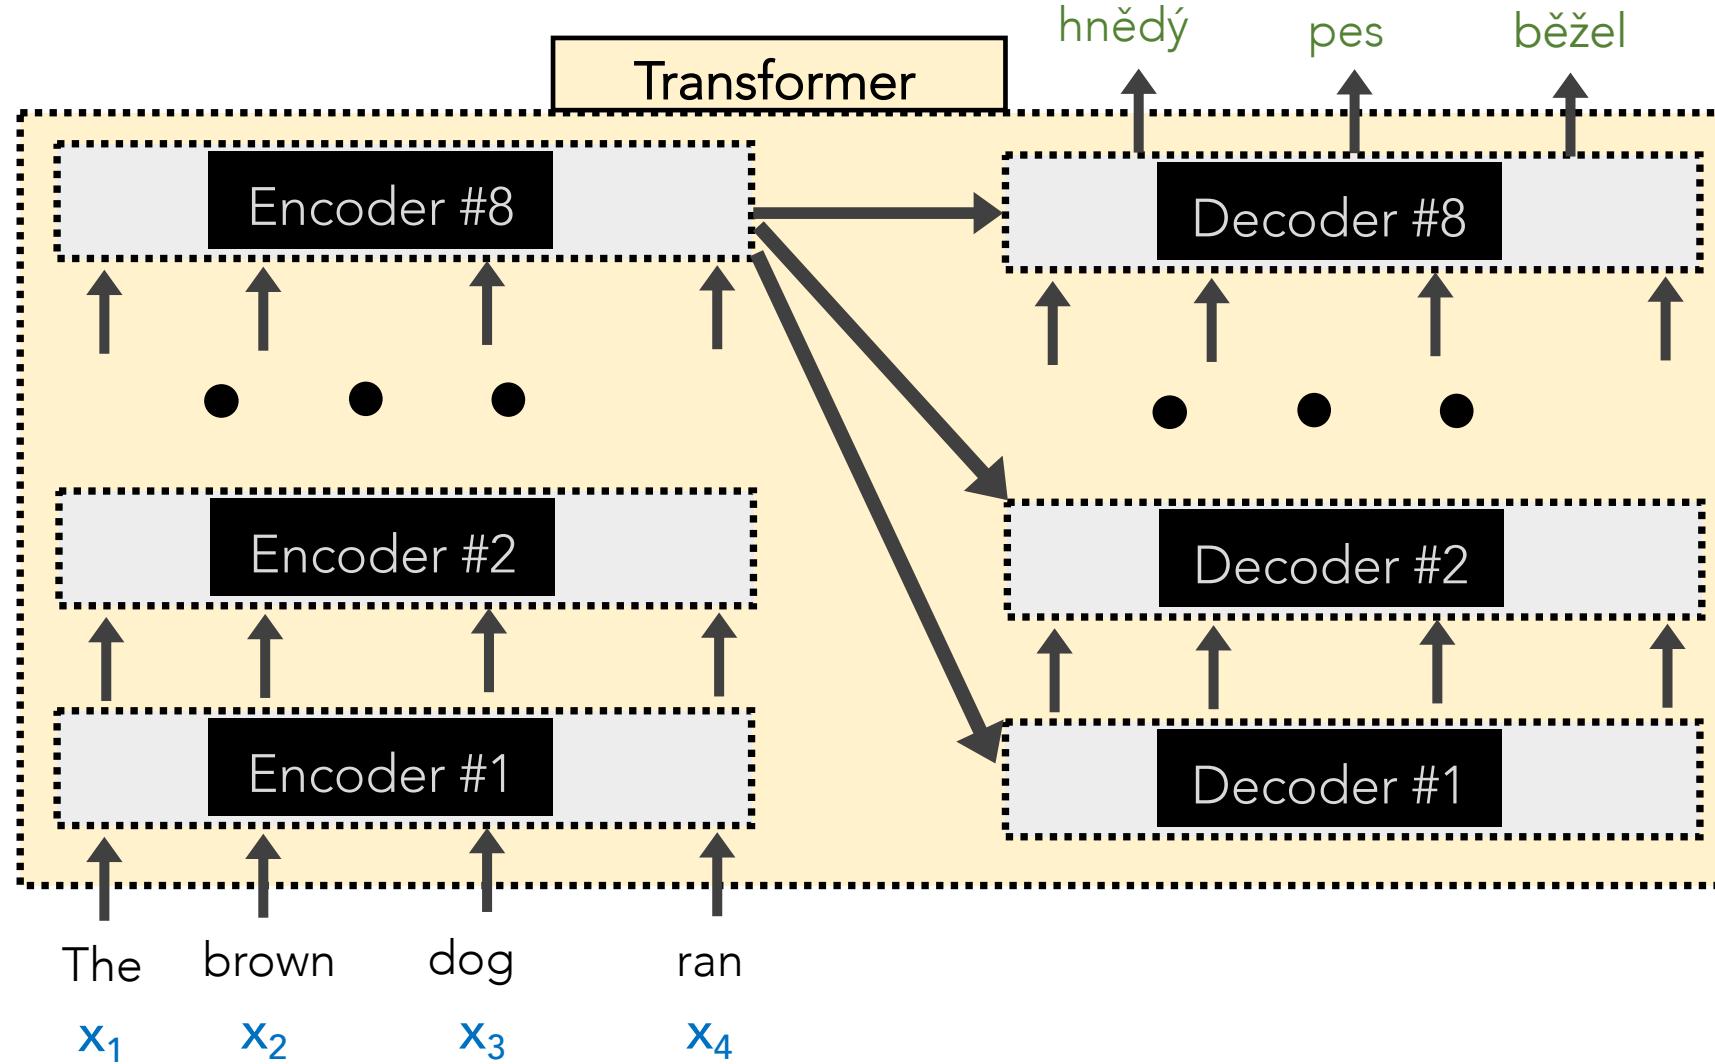


To recap: all of this looks fancy, but ultimately it's just producing a very good **contextualized embedding** r_i of each word x_i

Why stop with just 1
Transformer Encoder?
We could stack several!

The original Transformer model was intended for Machine Translation, so it had **Decoders**, too

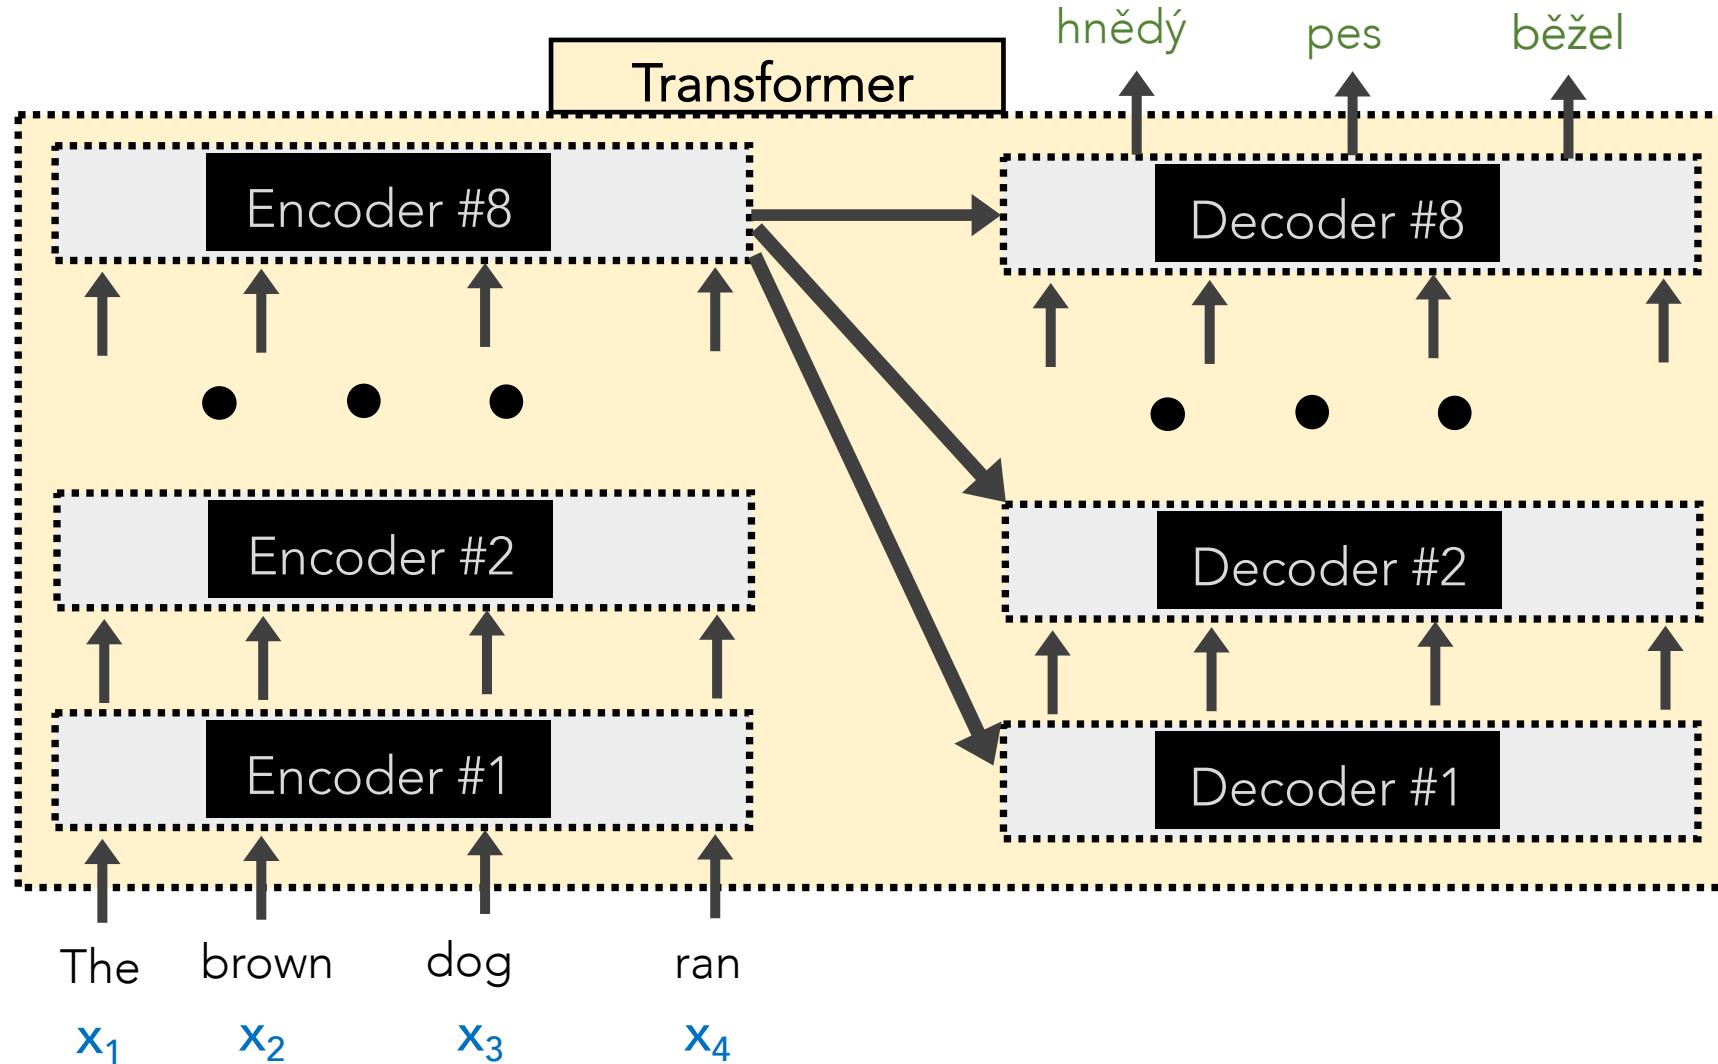
Transformer Encoders and Decoders



Transformer Encoders produce **contextualized embeddings** of each word

Transformer Decoders generate new sequences of text

Transformer Encoders and Decoders



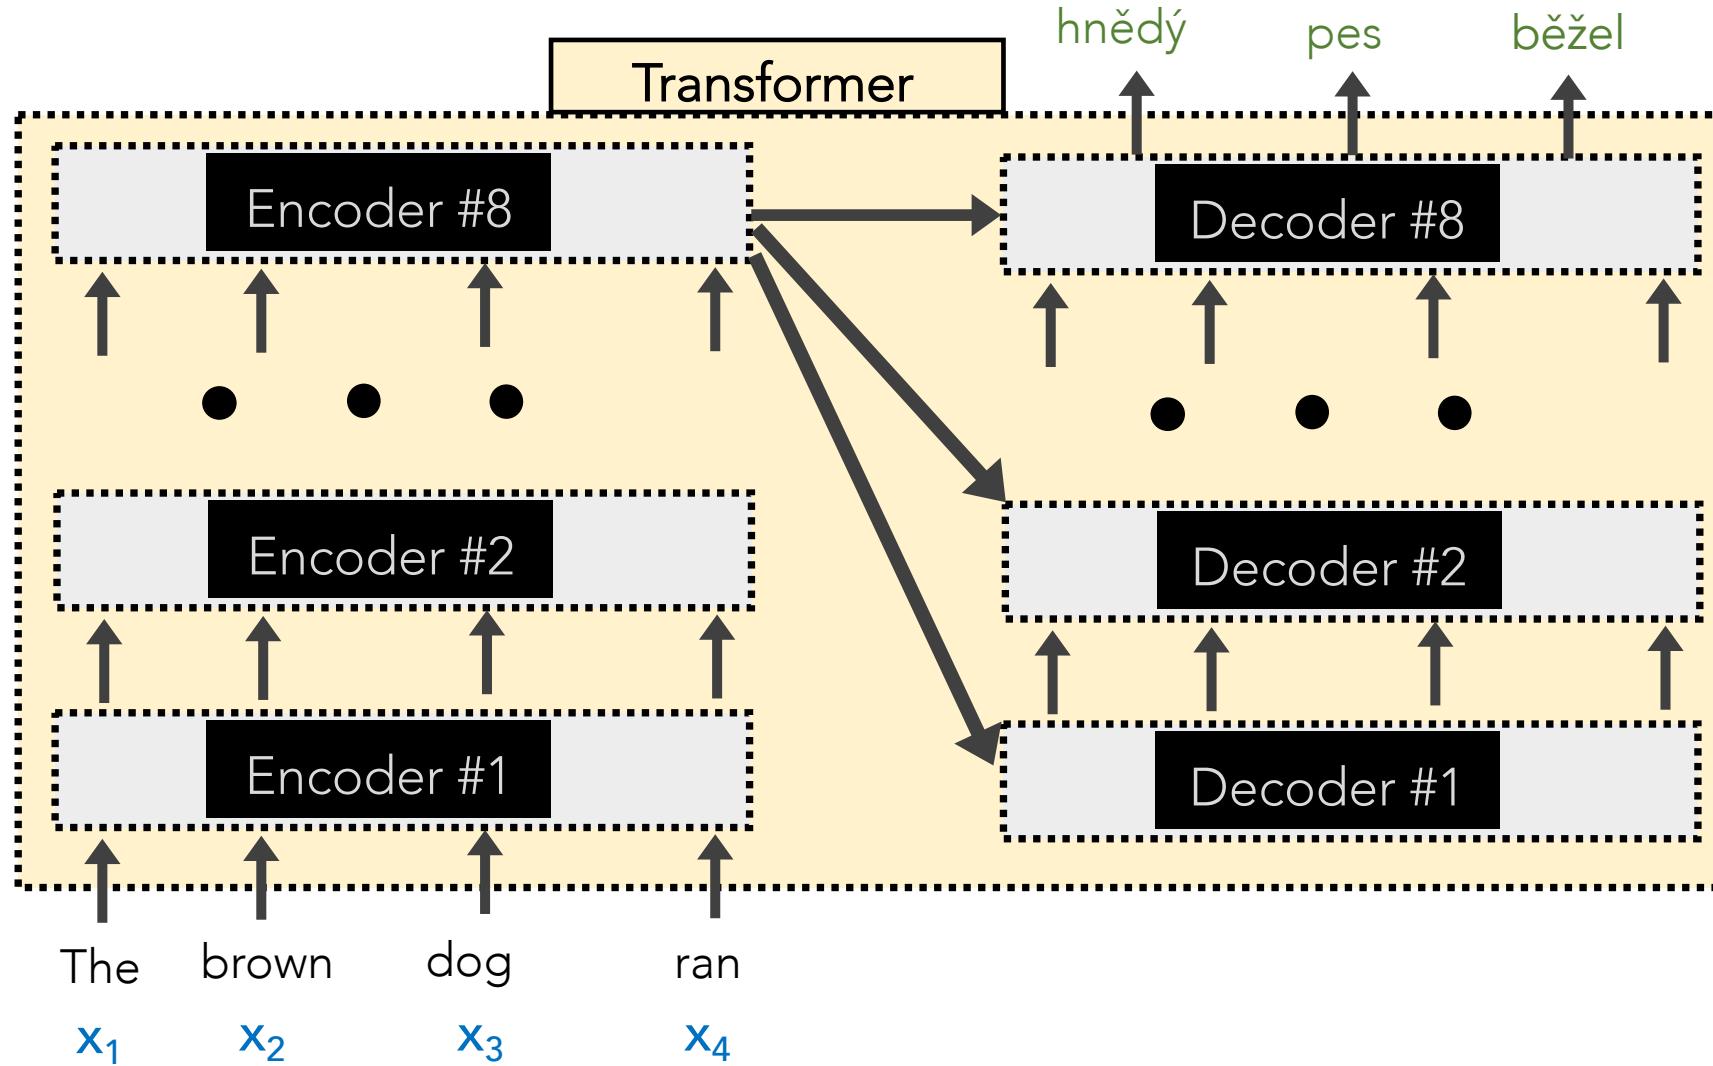
(regular attention
besides self-attention)

NOTE

Transformer Decoders are identical to the Encoders, except they have an additional **Attention Head** in between the Self-Attention and FFNN layers.

This additional **Attention Head** focuses on parts of the encoder's representations.

Transformer Encoders and Decoders

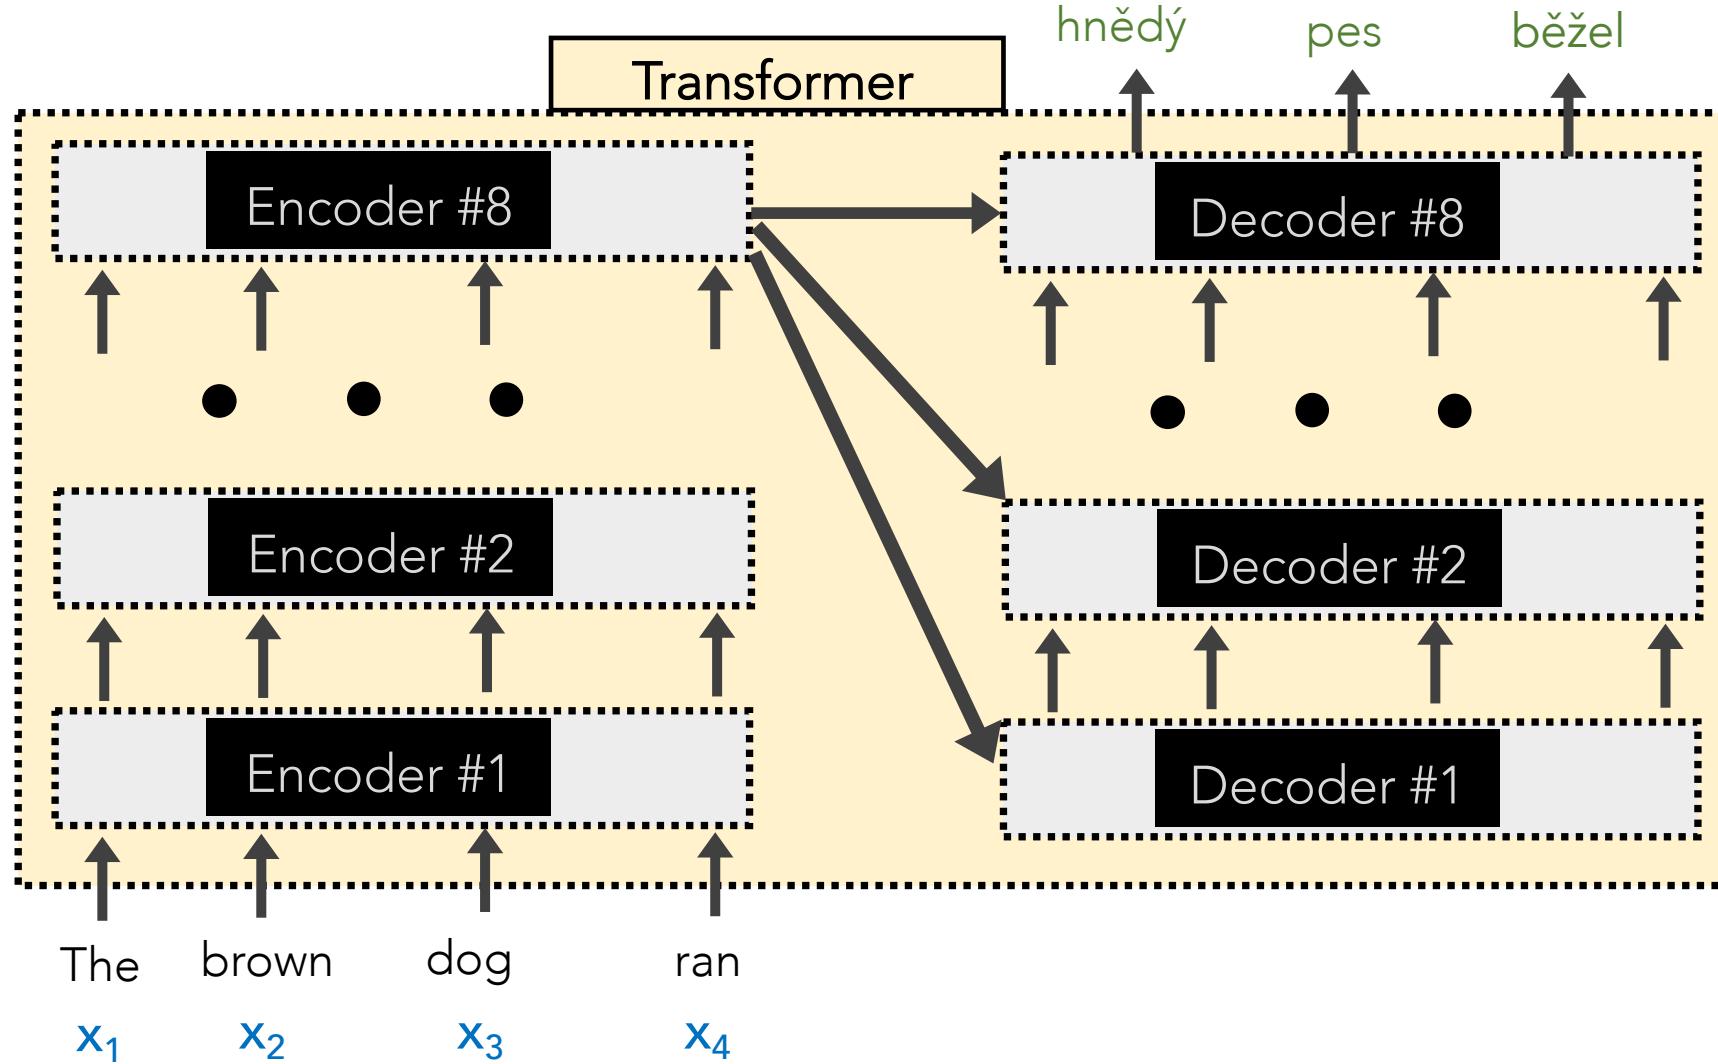


NOTE

The **query** vector for a Transformer Decoder's **Attention Head** (not Self-Attention Head) is from the output of the previous decoder layer.

However, the **key** and **value** vectors are from the **Transformer Encoders'** outputs.

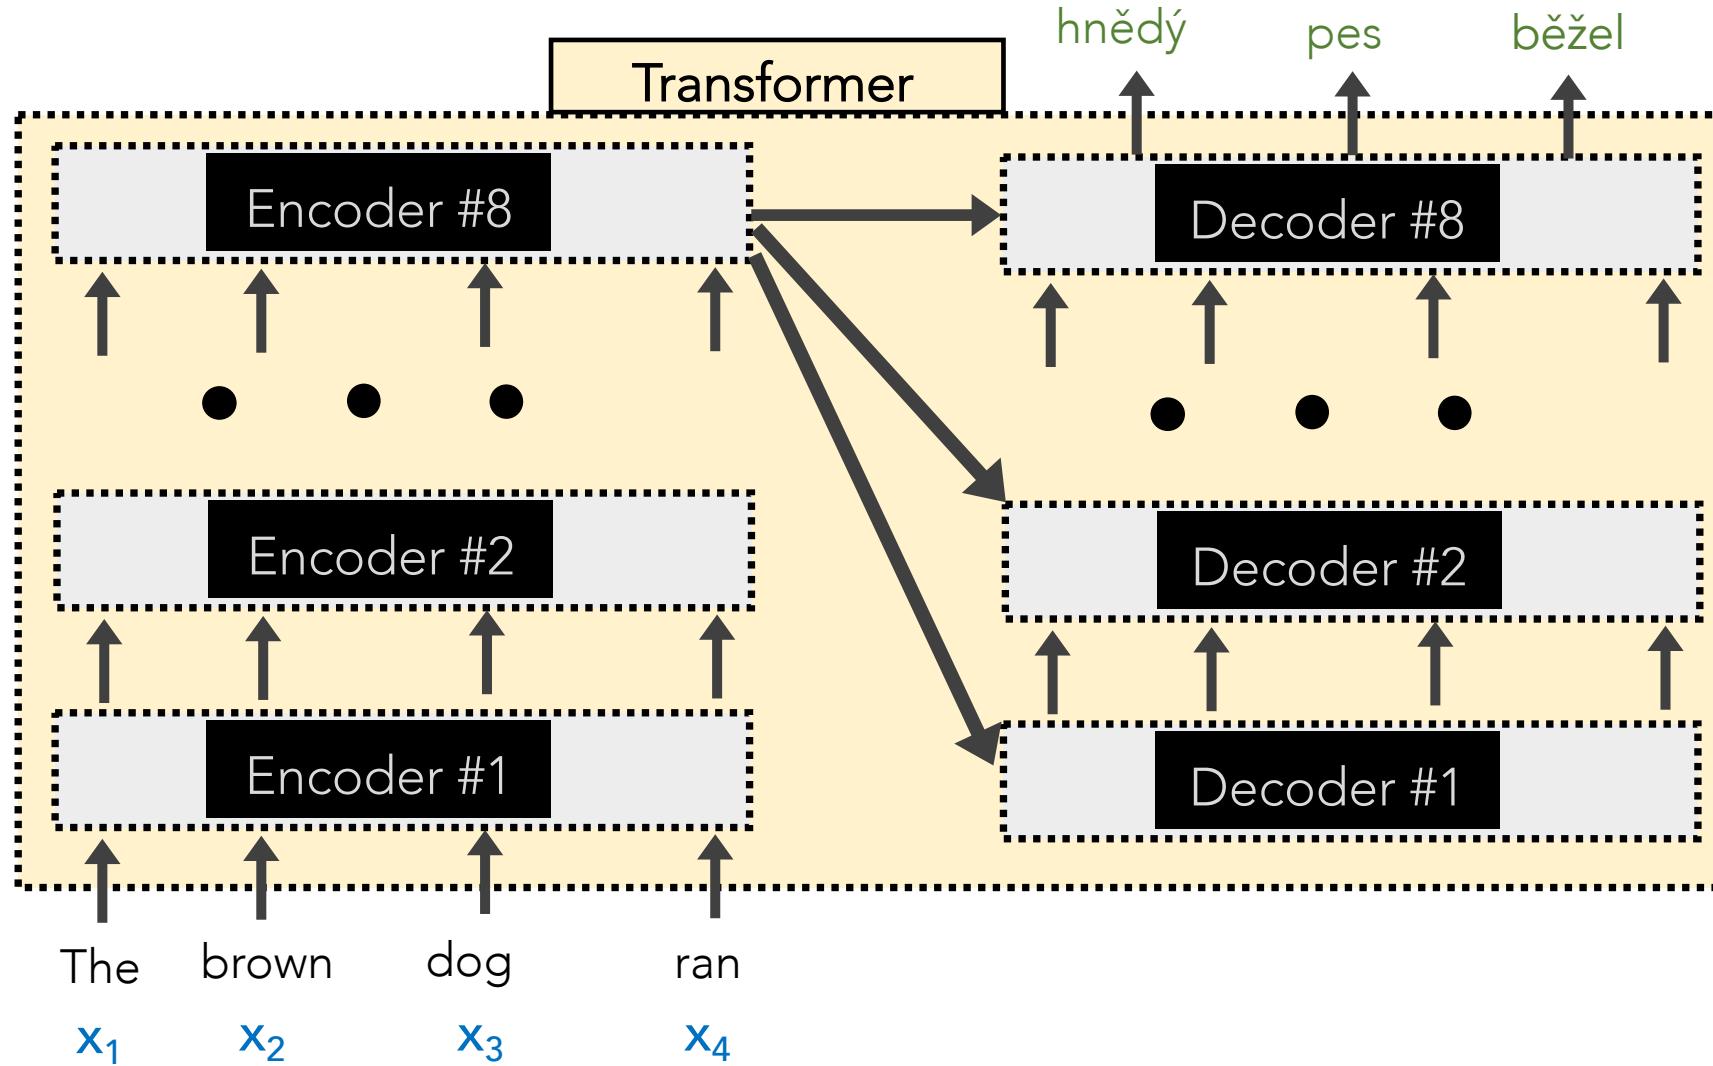
Transformer Encoders and Decoders



NOTE

The **query**, **key**, and **value** vectors for a Transformer Decoder's **Self-Attention Head** (not Attention Head) are all from the output of the previous decoder layer.

Transformer Encoders and Decoders



IMPORTANT

The Transformer **Decoders** have **positional embeddings**, too, just like the **Encoders**.

Critically, each position is **only allowed to attend to the previous indices**. This **masked Attention** preserves it as being an auto-regressive LM.

Loss Function: cross-entropy (predicting translated word)

Training Time: ~4 days on (8) GPUs

Layer Type	Complexity per Layer	Sequential Operations	Maximum Path Length
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$

Machine Translation results: state-of-the-art (at the time)

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1		$3.3 \cdot 10^{18}$
Transformer (big)	28.4	41.8		$2.3 \cdot 10^{19}$

Machine Translation results: state-of-the-art (at the time)

You can train to translate from Language A to Language B.

Then train it to translate from Language B. to Language C.

Then, without training, it can translate from Language A to Language C

- What if we don't want to decode/translate?
- Just want to perform a particular task (e.g., classification)
- Want even more robust, flexible, rich representation!
- Want **positionality** to play a more explicit role, while not being restricted to a particular form (e.g., CNNs)

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Bidirectional Encoder Representations from Transformers



Bidirectional Encoder Representations from Transformers

Like *Bidirectional LSTMs*, let's look in **both** directions



Bidirectional Encoder Representations from Transformers

Let's only use Transformer *Encoders*, no Decoders

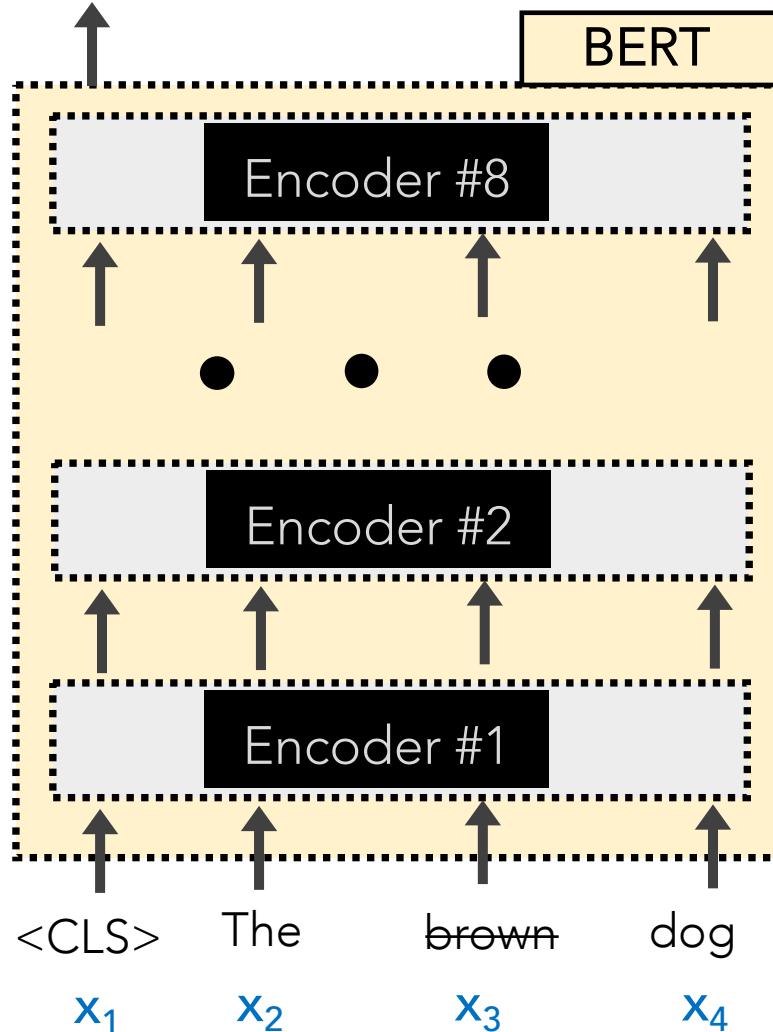


Bidirectional Encoder Representations from Transformers

It's a language model that builds rich representations



brown 0.92
lazy 0.05
playful 0.03



BERT has 2 training objectives:

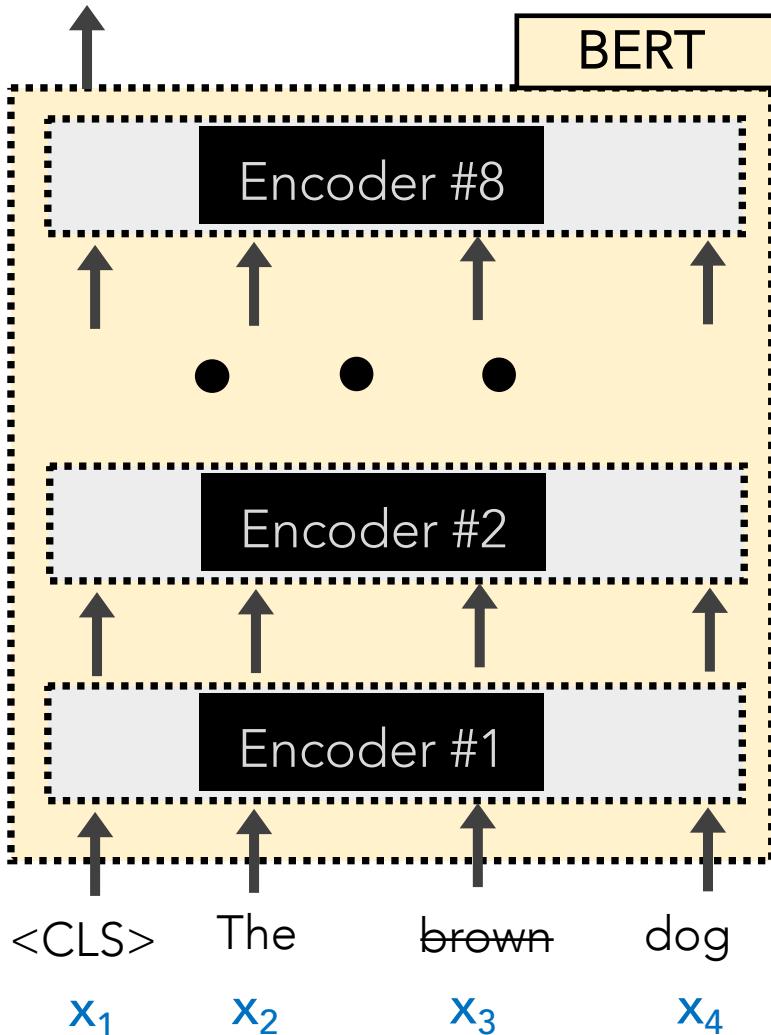
1. Predict the Masked word (a la CBOW)

15% of all input words are randomly masked.

- 80% become [MASK]
- 10% become revert back
- 10% become are deliberately corrupted as wrong words

BERT

brown 0.92
lazy 0.05
playful 0.03



BERT has 2 training objectives:

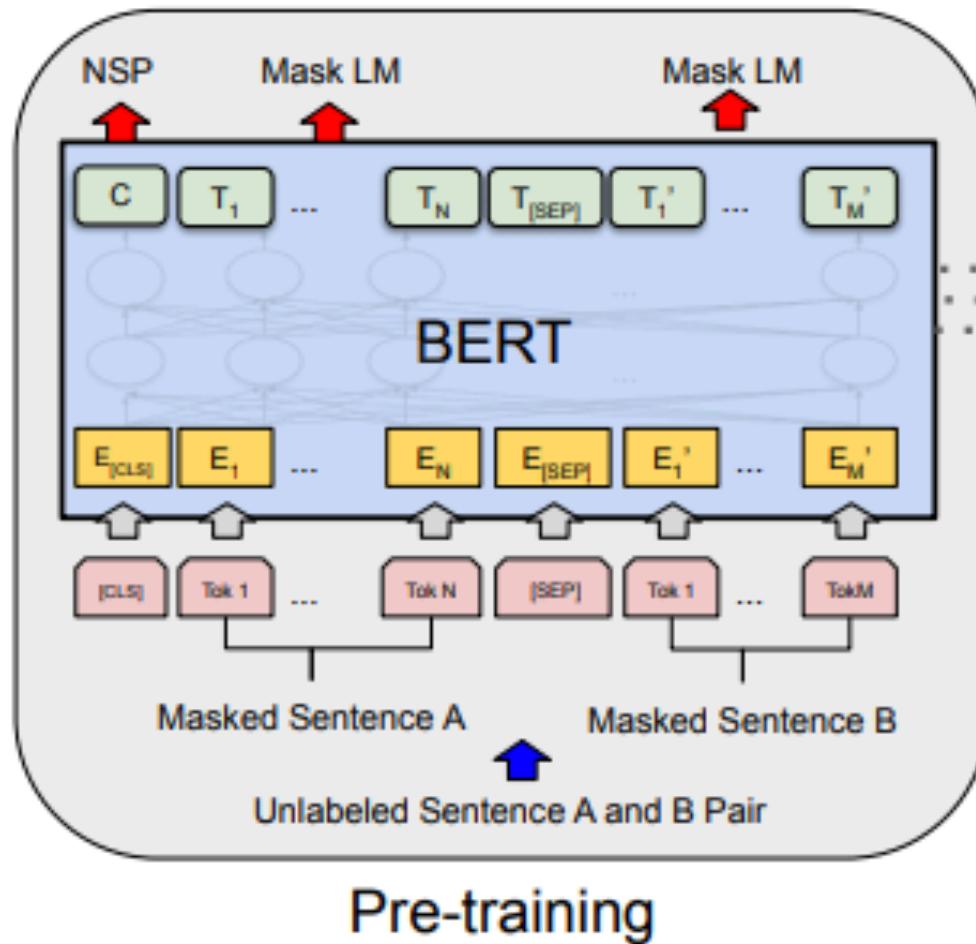


2. Two sentences are fed in at a time. Predict the if the second sentence of input truly follows the first one or not.

50% random sentences

50% truly follows the sentences

BERT

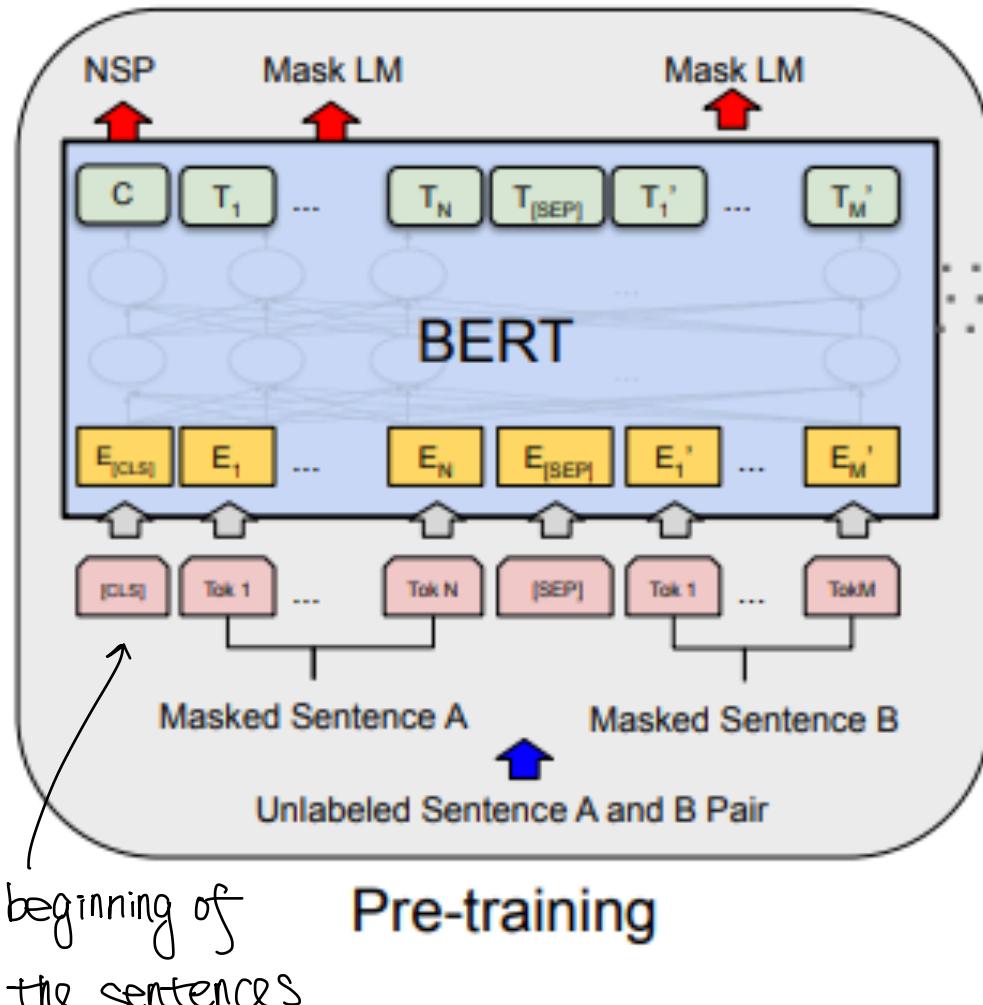


Every two sentences are separated by a $\langle \text{SEP} \rangle$ token.

50% of the time, the 2nd sentence is a randomly selected sentence from the corpus.

50% of the time, it truly follows the first sentence in the corpus.

BERT

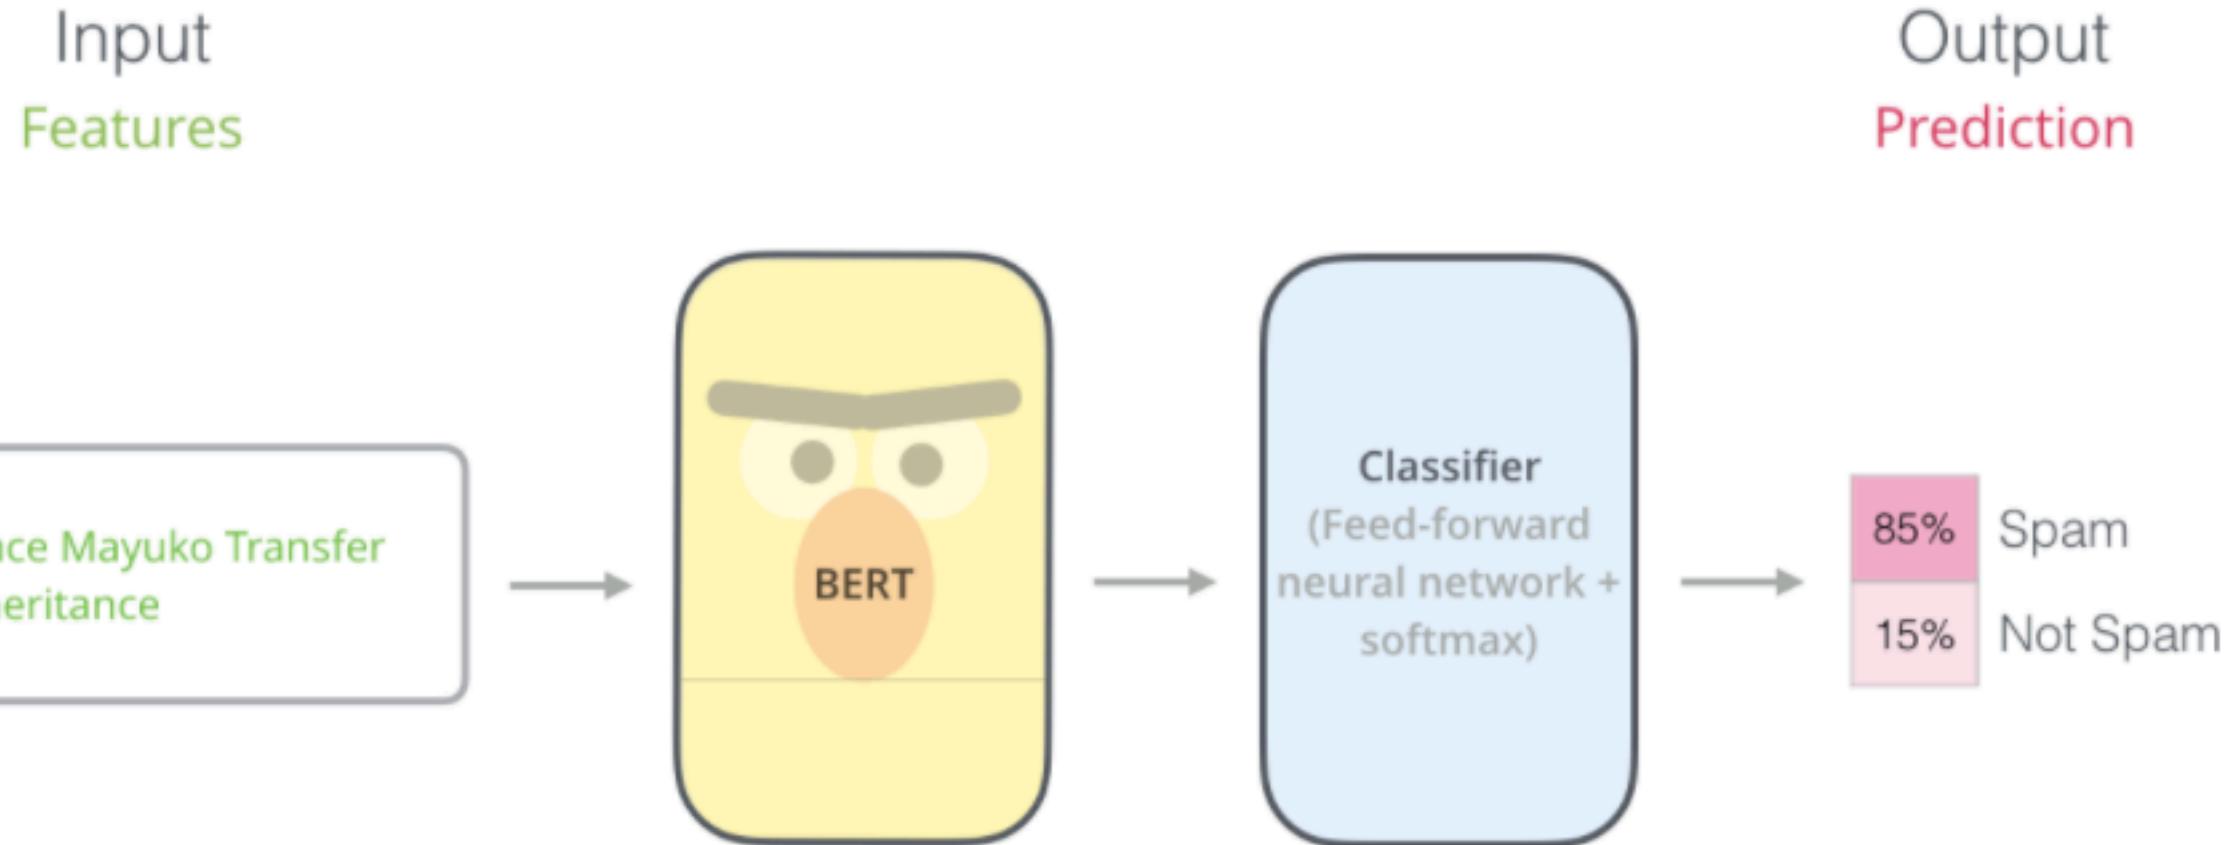


NOTE: BERT also embeds the inputs by their **WordPiece** embeddings.

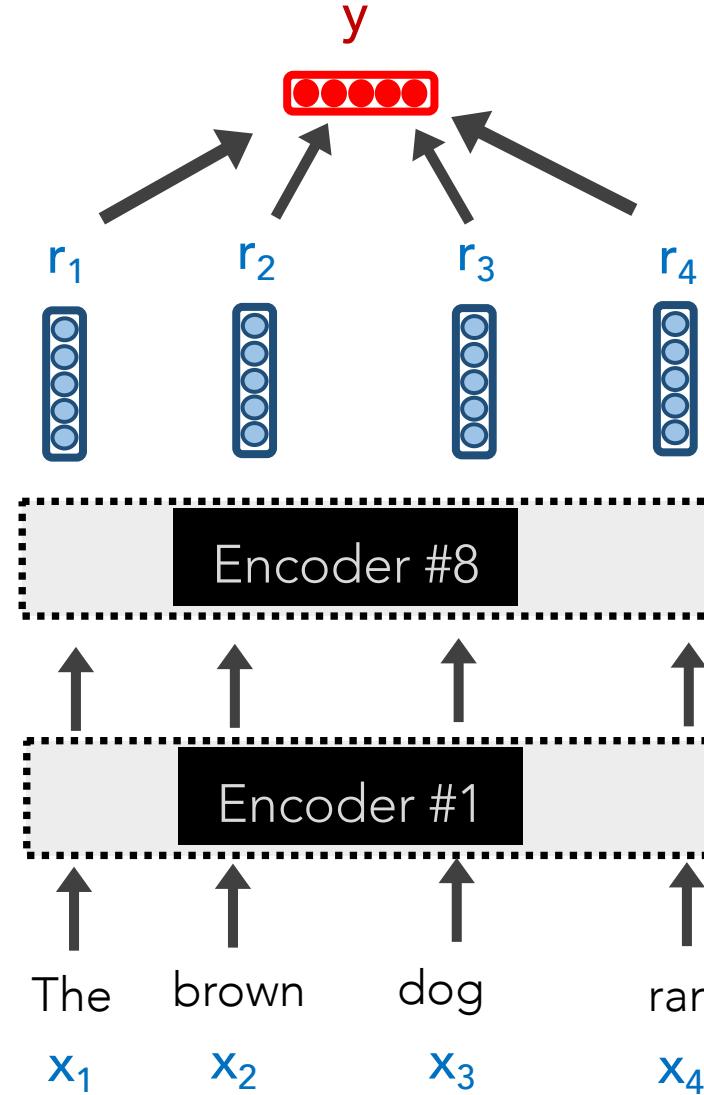
WordPiece is a sub-word tokenization learns to merge and use characters based on which pairs maximize the likelihood of the training data if added to the vocab.

BERT

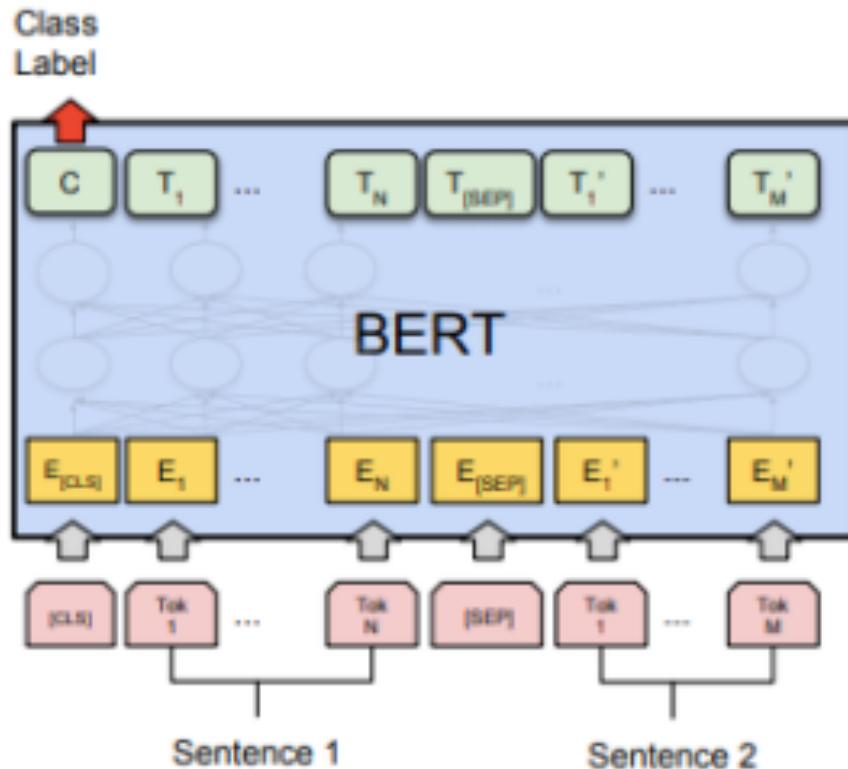
Most of the time, people use BERT
to fine-tune on a separate task



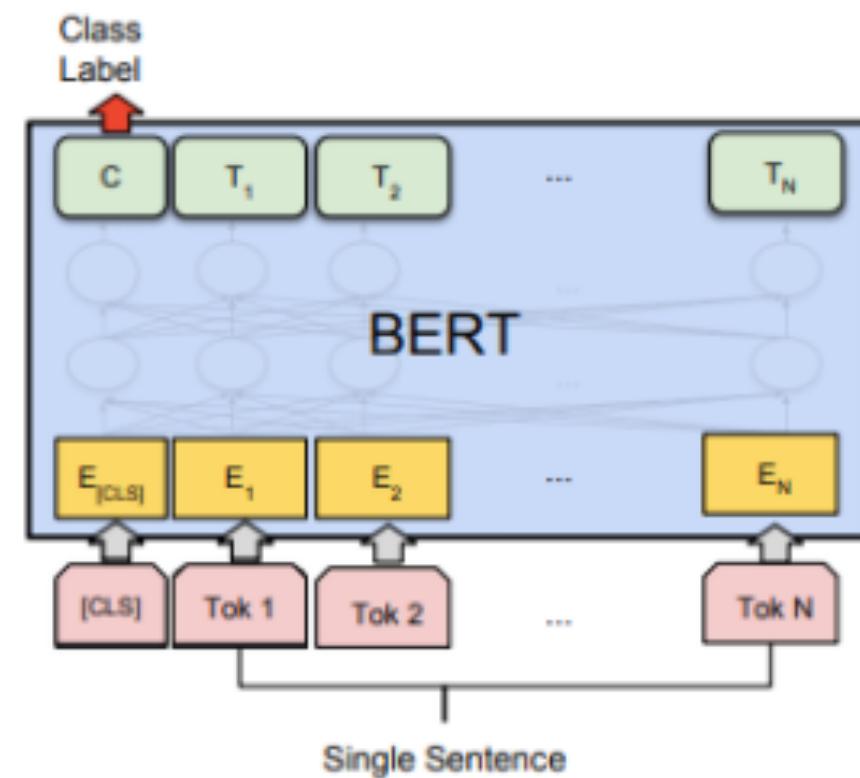
Most of the time, people use BERT
to fine-tune on a separate task



Most of the time, people use BERT
to fine-tune on a separate task

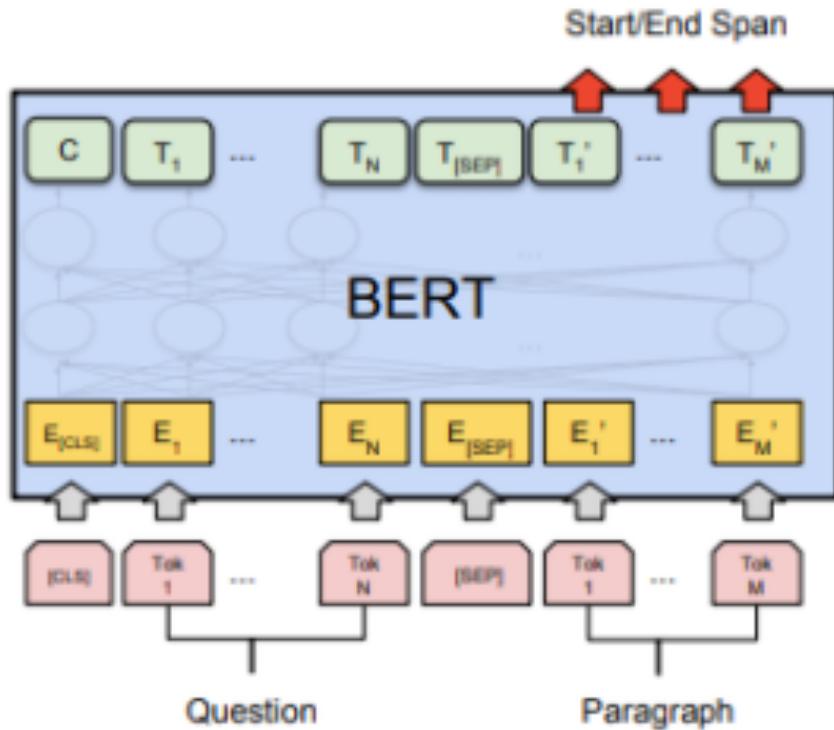


(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG

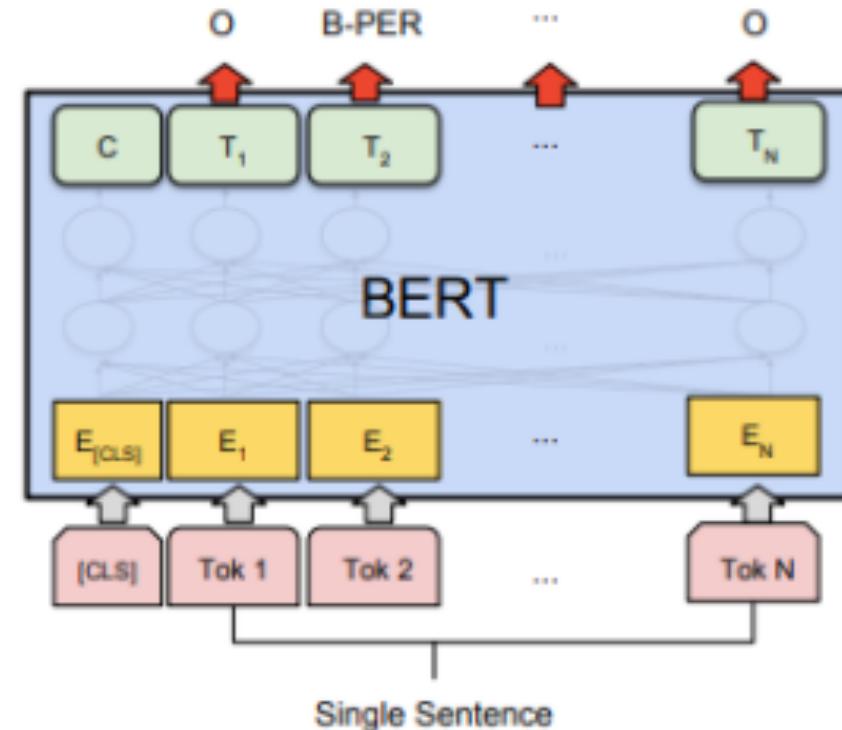


(b) Single Sentence Classification Tasks:
SST-2, CoLA

Most of the time, people use BERT
to fine-tune on a separate task

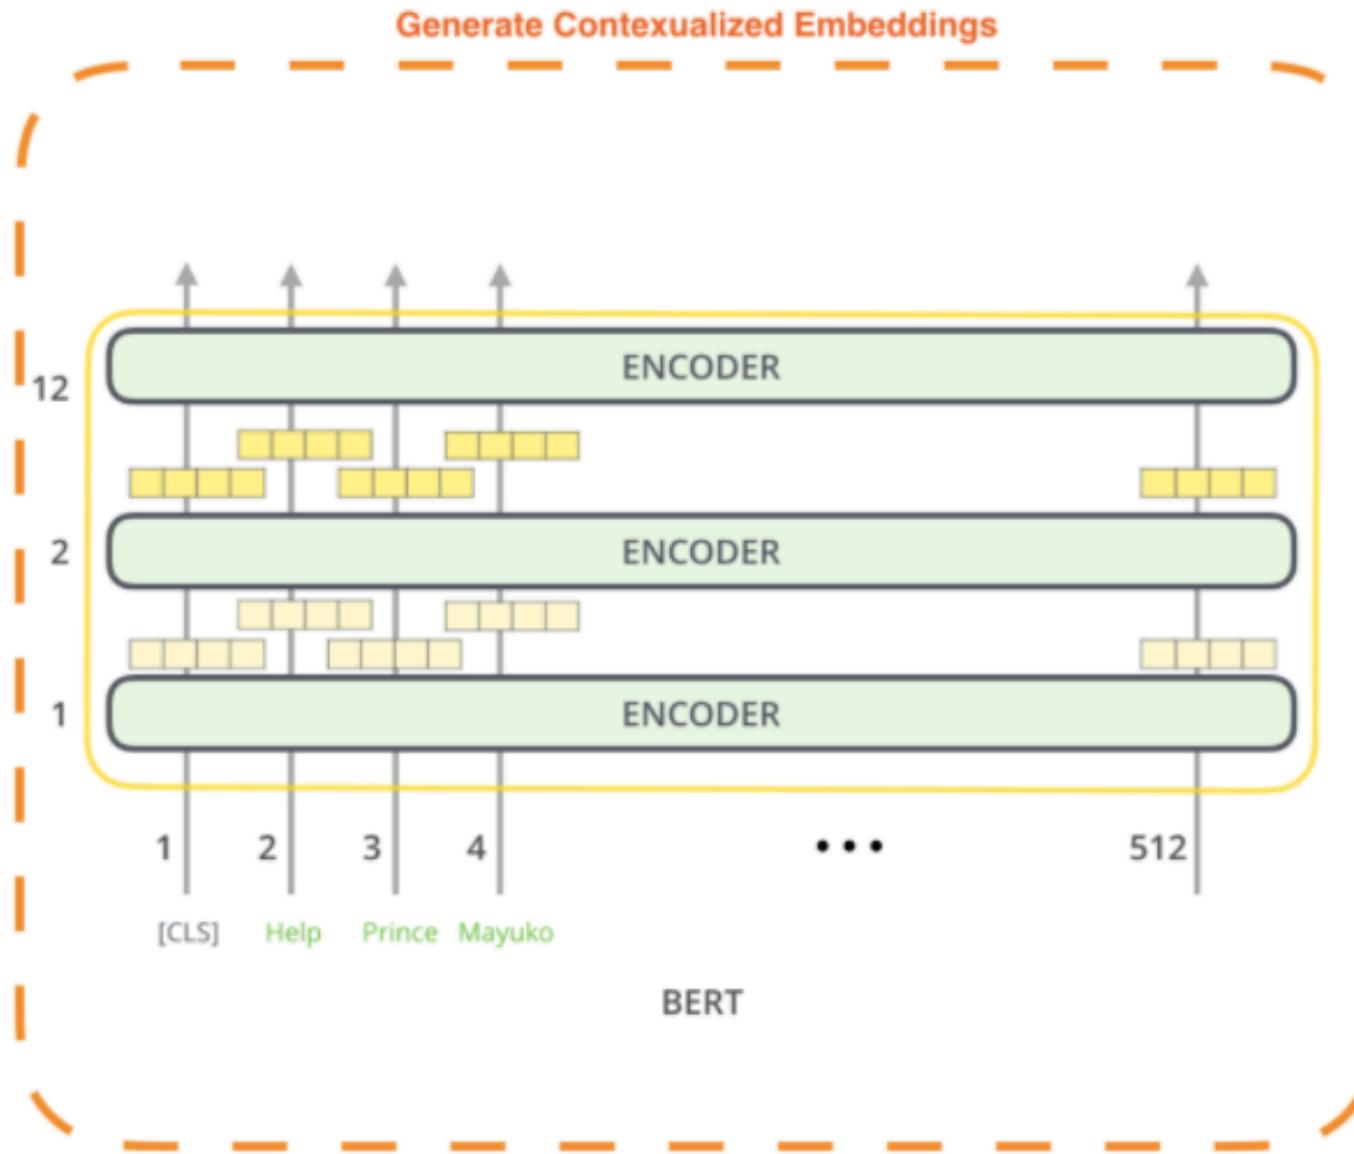


(c) Question Answering Tasks:
SQuAD v1.1

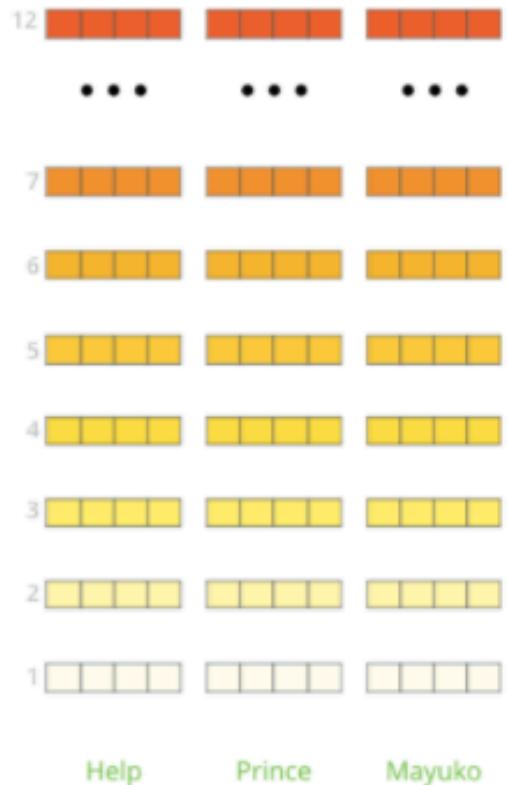


(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

One could also extract the **contextualized embeddings**

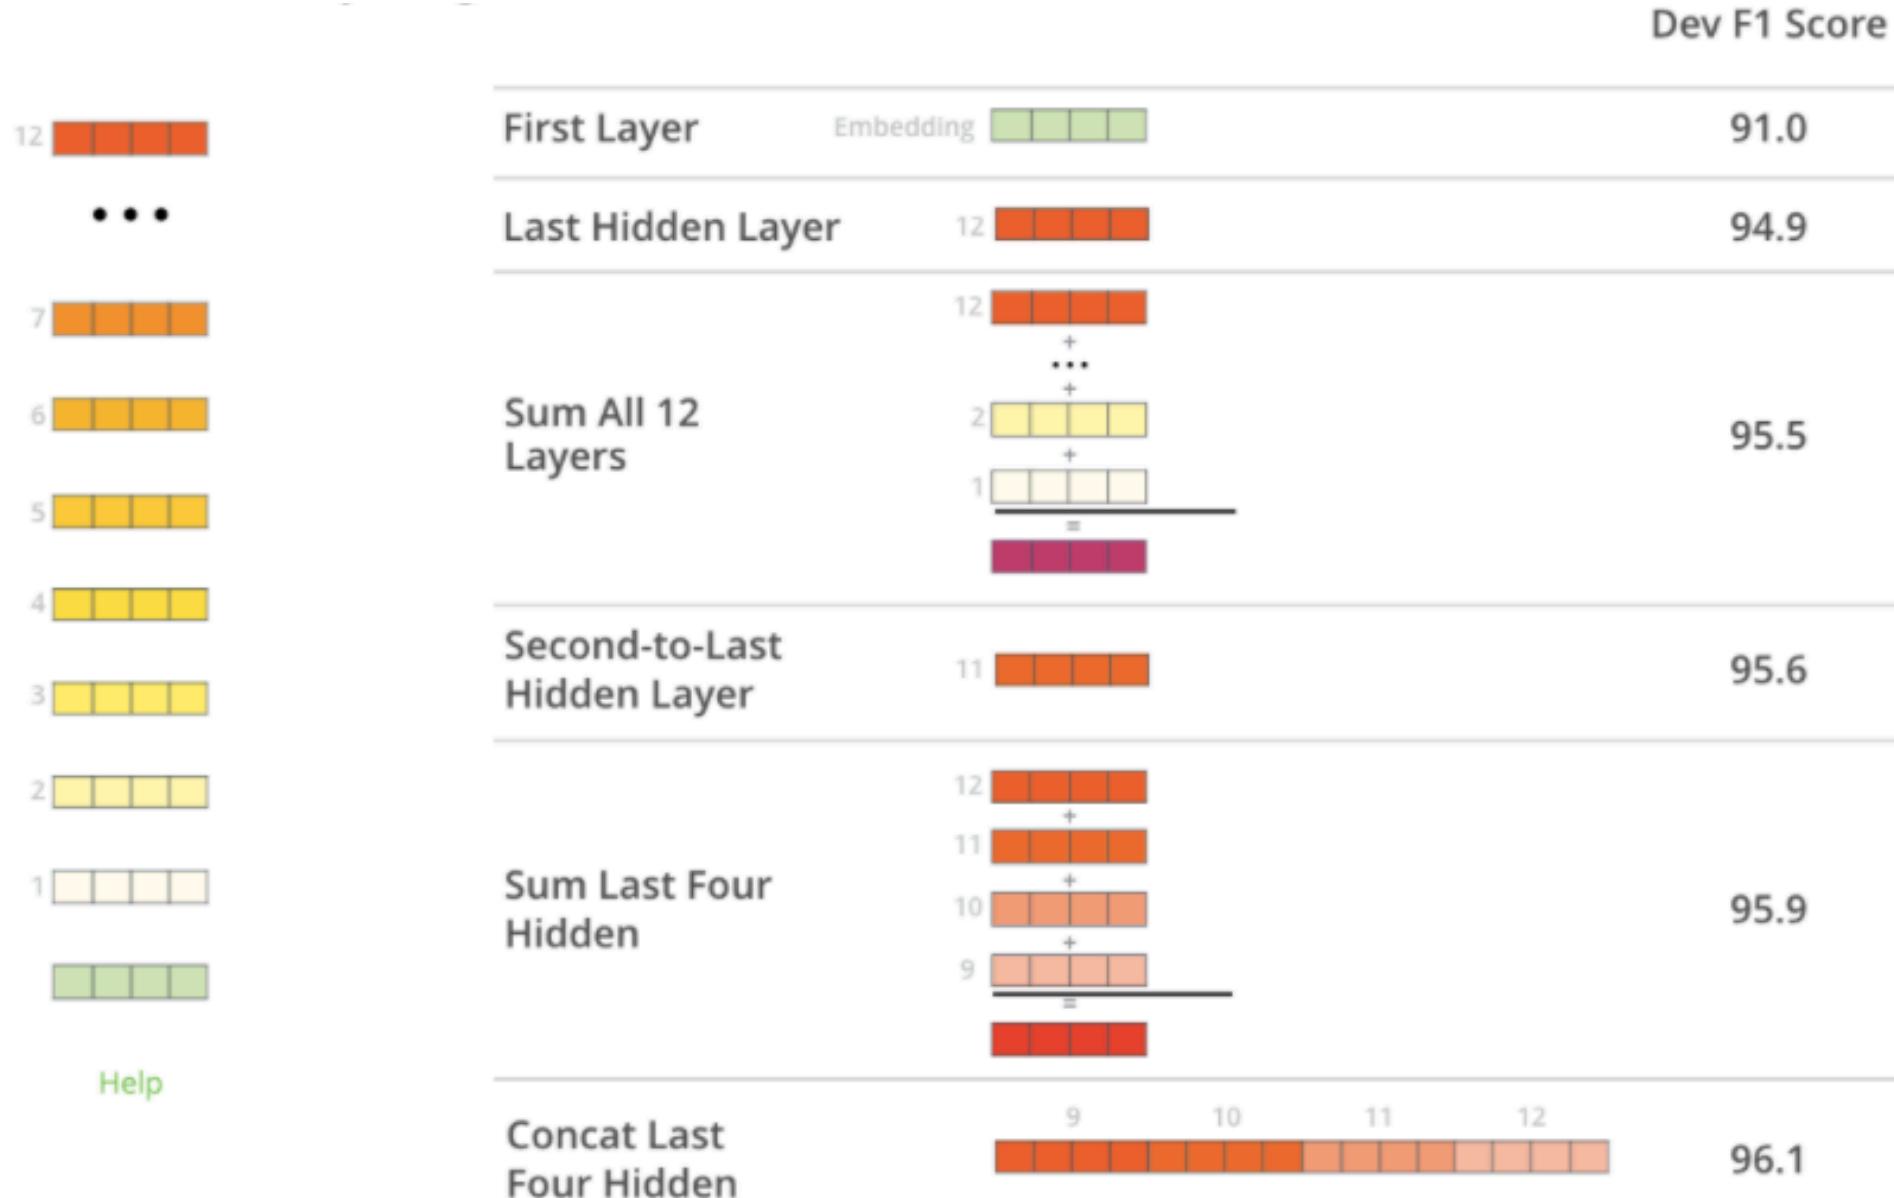


The output of each encoder layer along each token's path can be used as a feature representing that token.



But which one should we use?

Later layers have the best contextualized embeddings



BERT

BERT yields state-of-the-art (SOTA) results on many tasks

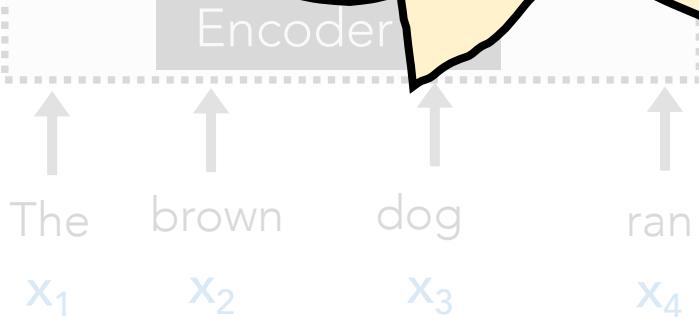
System	MNLI-(m/mm)	QQP	QNLI	SST-2	CoLA	STS-B	MRPC	RTE	Average
	392k	363k	108k	67k	8.5k	5.7k	3.5k	2.5k	-
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

Table 1: GLUE Test results, scored by the evaluation server (<https://gluebenchmark.com/leaderboard>).

Takeaway

BERT is incredible for learning contextualized embeddings of words and using transfer learning for other tasks (e.g., classification).

Can't generate new sentences though,
due to no decoders.



Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Transformer

What if we want to generate a new output sequence?

GPT-2 model to the rescue!

Generative Pre-trained Transformer 2

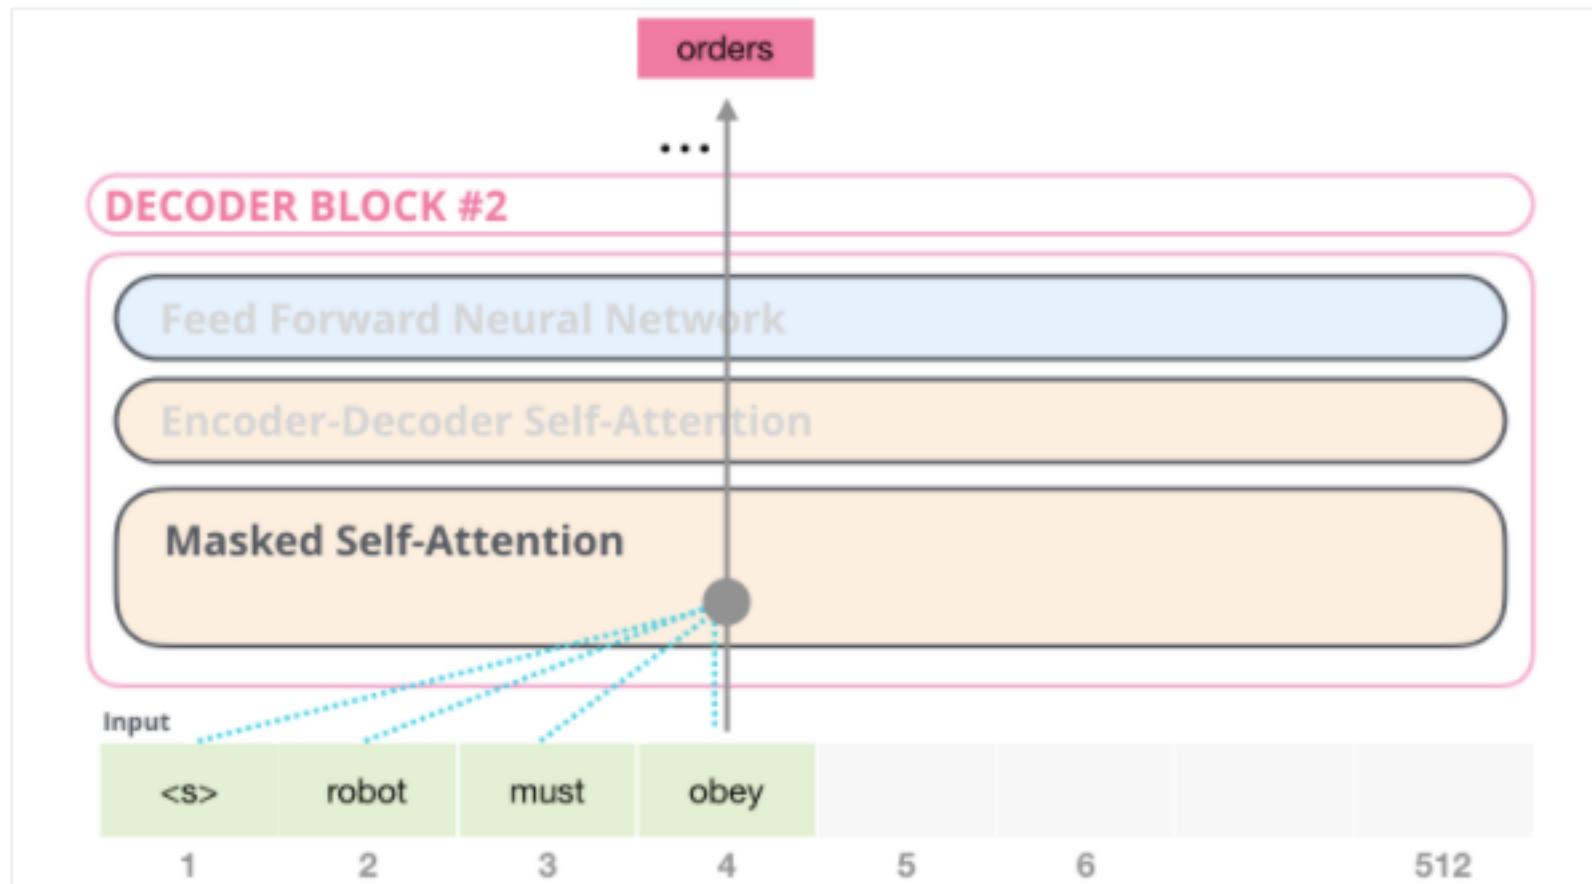
GPT-2 (a Transformer variant)

- GPT-2 uses only **Transformer Decoders** (no Encoders) to generate new sequences (from scratch or from a starting sequence)
- Oddly, there is **no Attention** (since there is no **Transformer Encoder** to attend to). So, there is only **Self-Attention**.
- As it processes each word/token, it cleverly **masks** the “future” words and conditions itself on the previous words

proper autoregressive model

GPT-2 (a Transformer variant)

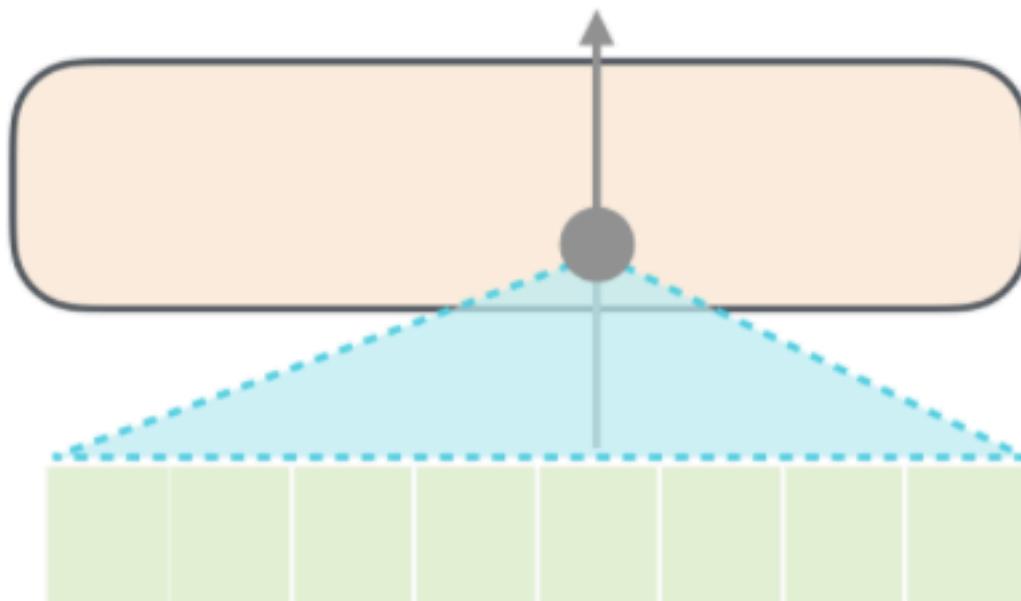
As it processes each word/token, it cleverly **masks** the “future” words and conditions itself on the previous words



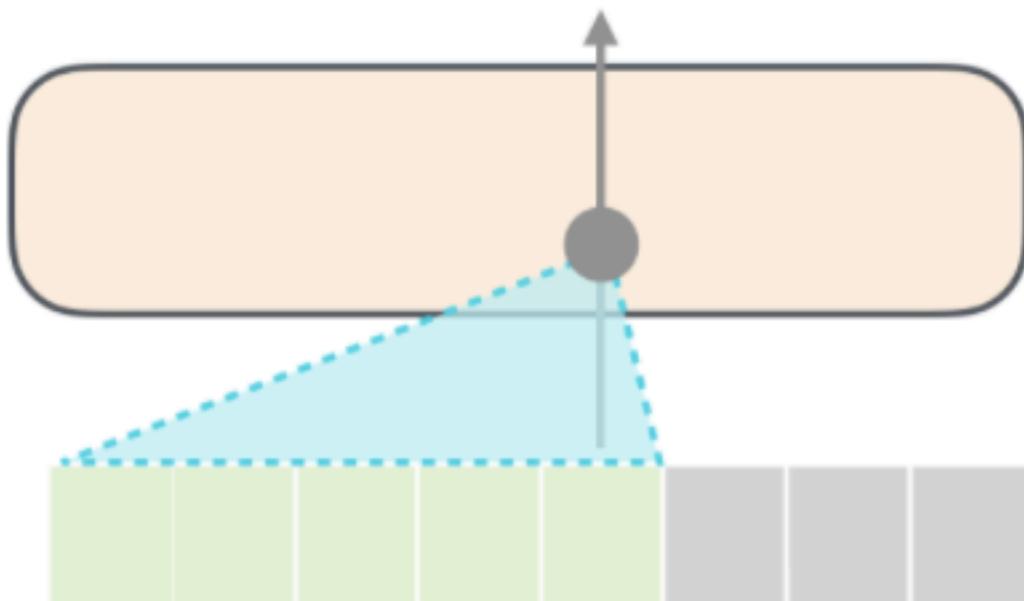
GPT-2 (a Transformer variant)

As it processes each word/token, it cleverly **masks** the “future” words and conditions itself on the previous words

Self-Attention



Masked Self-Attention



GPT-2 (a Transformer variant)

- Technically, it doesn't use words as input but **Byte Pair Encodings** (sub-words), similar to BERT's WordPieces.
- Includes **positional embeddings** as part of the input, too.
- Easy to fine-tune on your own dataset (language)

GPT-2 Results

Easy to fine-tune on your own dataset (language)

SYSTEM PROMPT (HUMAN-WRITTEN)

A train carriage containing controlled nuclear materials was stolen in Cincinnati today. Its whereabouts are unknown.

MODEL COMPLETION (MACHINE-WRITTEN, FIRST TRY)

The incident occurred on the downtown train line, which runs from Covington and Ashland stations.

In an email to Ohio news outlets, the U.S. Department of Energy said it is working with the Federal Railroad Administration to find the thief.

“The theft of this nuclear material will have significant negative consequences on public and environmental health, our workforce and the economy of our nation,” said Tom Hicks, the U.S. Energy Secretary, in a statement. “Our top priority is to secure the theft and ensure it doesn’t happen again.” The stolen material was taken from the University of Cincinnati’s Research Triangle Park nuclear research site, according to a news release from Department officials.

The Nuclear Regulatory Commission did not immediately release any information.

GPT-2 Results

Question	Generated Answer	Correct	Probability
Who wrote the book the origin of species?	Charles Darwin	✓	83.4%
Who is the founder of the ubuntu project?	Mark Shuttleworth	✓	82.0%
Who is the quarterback for the green bay packers?	Aaron Rodgers	✓	81.1%
Panda is a national animal of which country?	China	✓	76.8%
Who came up with the theory of relativity?	Albert Einstein	✓	76.4%
When was the first star wars film released?	1977	✓	71.4%
What is the most common blood type in sweden?	A	✗	70.6%
Who is regarded as the founder of psychoanalysis?	Sigmund Freud	✓	69.3%
Who took the first steps on the moon in 1969?	Neil Armstrong	✓	66.8%
Who is the largest supermarket chain in the uk?	Tesco	✓	65.3%
What is the meaning of shalom in english?	peace	✓	64.0%
Who was the author of the art of war?	Sun Tzu	✓	59.6%
Largest state in the us by land mass?	California	✗	59.2%
Green algae is an example of which type of reproduction?	parthenogenesis	✗	56.5%
Vikram samvat calender is official in which country?	India	✓	55.6%
Who is mostly responsible for writing the declaration of independence?	Thomas Jefferson	✓	53.3%

GPT-2 Results

Language Models are Unsupervised Multitask Learners

	LAMBADA (PPL)	LAMBADA (ACC)	CBT-CN (ACC)	CBT-NE (ACC)	WikiText2 (PPL)	PTB (PPL)	enwik8 (BPB)	text8 (BPC)	WikiText103 (PPL)
SOTA	99.8	59.23	85.7	82.3	39.14	46.54	0.99	1.08	18.3
117M	35.13	45.99	87.65	83.4	29.41	65.85	1.16	1.17	37.50
345M	15.60	55.48	92.35	87.1	22.76	47.33	1.01	1.06	26.37
762M	10.87	60.12	93.45	88.0	19.93	40.31	0.97	1.02	22.05
1542M	8.63	63.24	93.30	89.05	18.34	35.76	0.93	0.98	17.48

Table 3. Zero-shot results on many datasets. No training or fine-tuning was performed for any of these results. PTB and results are from (Gong et al., 2018). CBT results are from (Bajgar et al., 2016). LAMBADA accuracy result is from (Hoang and LAMBADA perplexity result is from (Grave et al., 2016). Other results are from (Dai et al., 2019).

GPT-2

- GPT-2 uses a large dataset of new text
 - As it processes text, it generates words
 - Can generate new text
 - Easy to fine-tune
 - **GPT-3** is an even bigger version of GPT-2, but isn't open-source
- Takeaway:**
GPT-2 is astounding for generating realistic-looking new text
- Can fine-tune toward other tasks, too.

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

BERT (a Transformer variant)

BERT is trained on a lot of text data:

- BooksCorpus (800M words)
- English Wikipedia (2.5B words)

Yay, for transfer learning!

BERT-Base model has 12 transformer blocks, 12 attention heads,

110M parameters!

BERT-Large model has 24 transformer blocks, 16 attention heads,

340M parameters!

GPT-2 (a Transformer variant)

GPT-2 is:

- trained on 40GB of text data (8M webpages)!
- 1.5B parameters

GPT-3 is an even bigger version (175B parameters) of GPT-2, but isn't open-source

Yay, for transfer learning!

Concerns

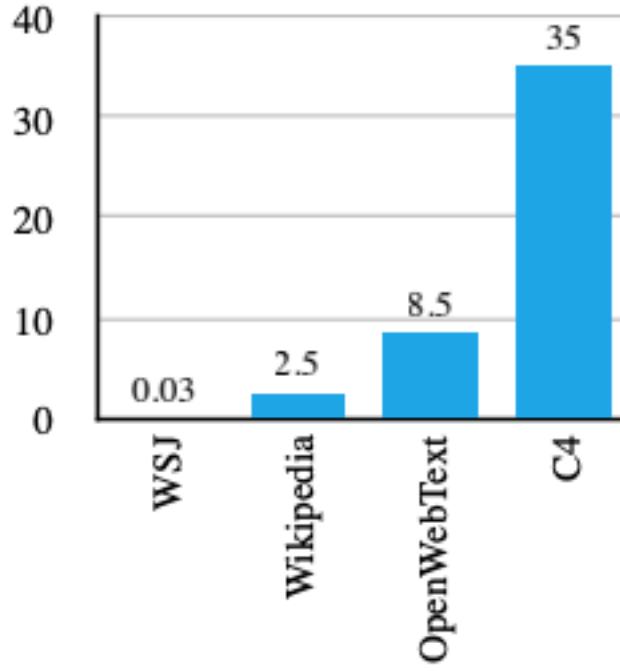
There are several issues to be aware of:

- It is very costly to train these large models. The companies who develop these models easily spend an entire month training one model, which uses **incredible amounts of electricity**.
- **BERT** alone is estimated to cost over **\$1M** for their final models
 - \$2.5k - \$50k (110 million parameter model)
 - \$10k - \$200k (340 million parameter model)
 - \$80k - \$1.6m (1.5 billion parameter model)

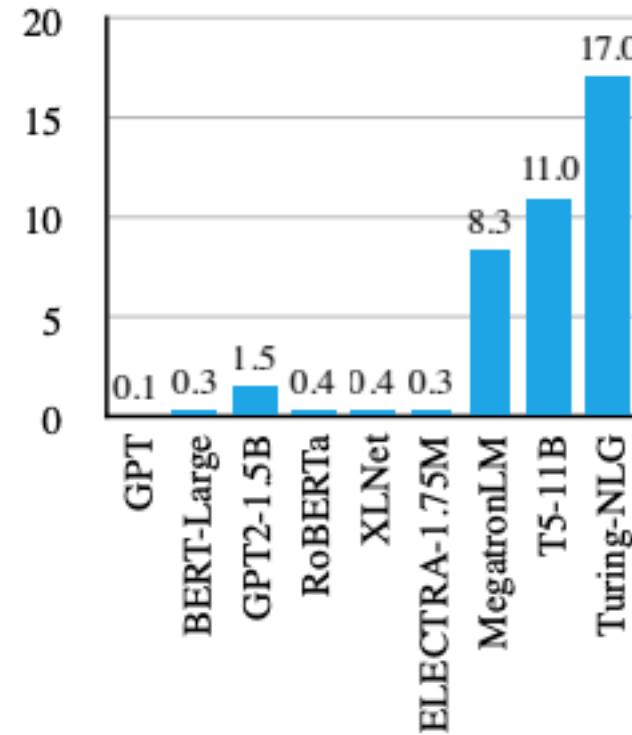
Concerns

It is very costly to train these large models.

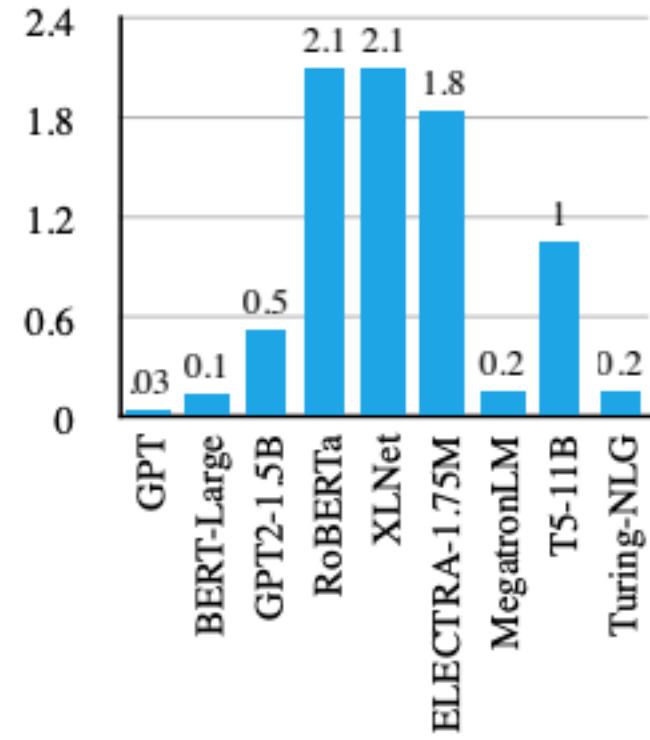
**Data Size
(billion words)**



**Model Size
(billion parameters)**



**Training Volume[†]
(trillion tokens)**



Concerns

- Further, these very large language models have been shown to be **biased** (e.g., in terms of gender, race, sex, etc).
- Converting from one language to another often converts gender neutral pronouns to sexist stereotypes
- Using these powerful LMs comes with **risks of producing** such text and/or evaluating/predicting tasks **based on these biased assumptions.**
- People are researching how to improve this

Concerns

- As computer-generated text starts to become indistinguishable from authentic, human-generated text, consider the ethical impact of fraudulently claiming text to be from a particular author.
- If used maliciously, it can easily contribute toward the problem of Fake News

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Outline



Recap



Transformers



BERT



GPT-2



Concerns



Summary

Summary

- There has been significant NLP progress in the past few years.
- Some of the complex models are incredible, but rely on having a lot of data and computational resources (e.g., Transformers)
- With all **data science** and **machine learning**, it's best to understand your data and task very well, clean your data, and start with a simple model (instead of jumping to the most complex model)

Summary

- NLP is incredibly fun, with infinite number of problems to work on
- I'll teach an **NLP course** next year (most likely Spring 2022).

Some definitions to remember

Models

- N-gram: count statistics; elementary sequence modelling
- FFNN: fixed-length context window; not ideal for sequential modelling
- (Vanilla) RNN: uses context; fair sequence modelling
- LSTM: a variant of an RNN that handles long-range context better
- Seq2Seq: maps 1 sequence to another ($n \rightarrow m$ sequences)
- Attention: determines which elements in **sequence A** pertain to **sequence B**
- Self-Attention: determines which elements to focus on in its own **sequence A**
- Transformers: learns excellent representation, via a **seq2seq** framework with **self-attention** and **attention**
 - BERT: Transformer Encoders that learn great representations and can be fine-tuned on other tasks
 - GPT-2: Transformer Decoders that generate realistic text and can be fine-tuned on other tasks

QUESTIONS?

BACKUP SLIDES

Transformer vs CNN vs RNN

